Machine Learning and the Physical World

Lecture 3 : Gaussian Processes

Carl Henrik Ek - che29@cam.ac.uk

15th of October, 2021

http://carlhenrik.com

Model+Data $\rightarrow$ Predict

$$\text{Data} + \text{Model} \overset{Compute}{\rightarrow} \text{Prediction}$$

Model+Data $\rightarrow$ Predict

$$\text{Data} + \text{Model} \overset{Compute}{\rightarrow} \text{Prediction}$$

- We do not have enough knowledge

$$\text{Model} \overset{Compute}{\not\rightarrow} \text{Prediction}$$

Model+Data $\rightarrow$ Predict

$$\text{Data} + \text{Model} \overset{Compute}{\rightarrow} \text{Prediction}$$

- We do not have enough knowledge

$$\text{Model} \overset{Compute}{\not\rightarrow} \text{Prediction}$$

- We do not have enough data

$$\text{Data} \overset{Compute}{\not\rightarrow} \text{Prediction}$$

Data and Knowledge

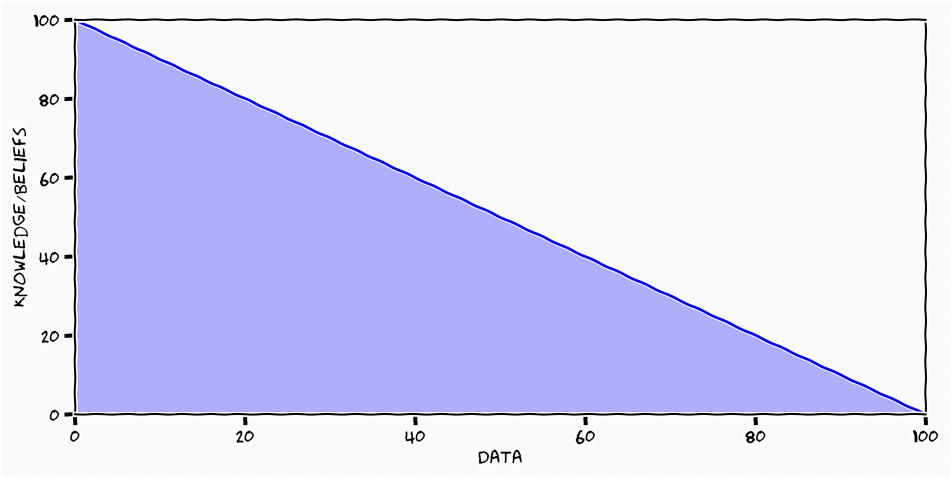

Data and Knowledge

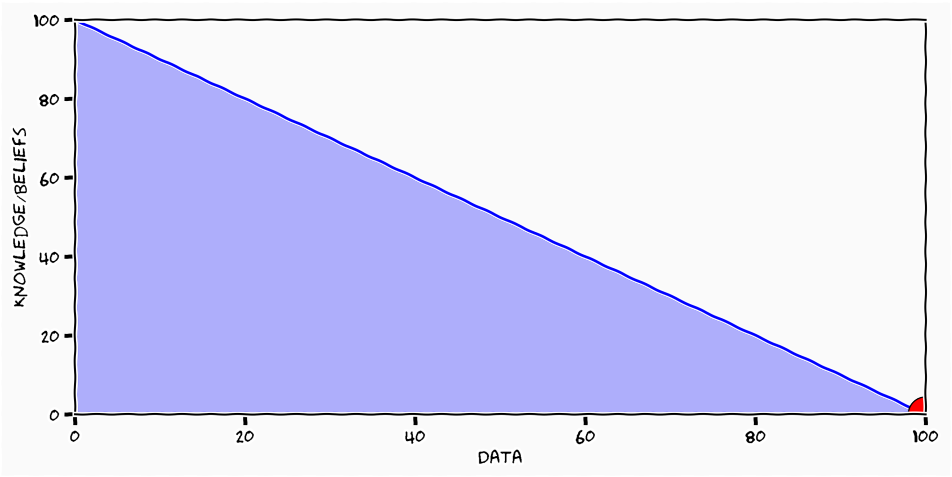

Data and Knowledge

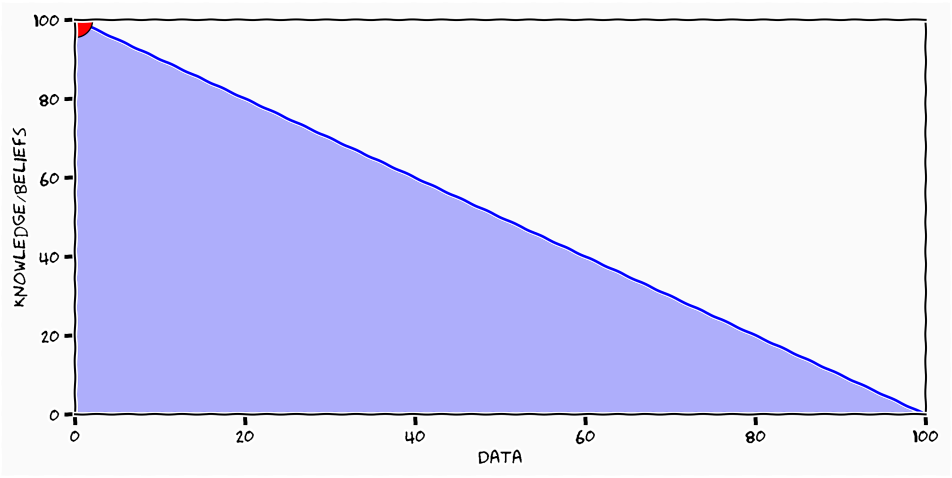

Data and Knowledge

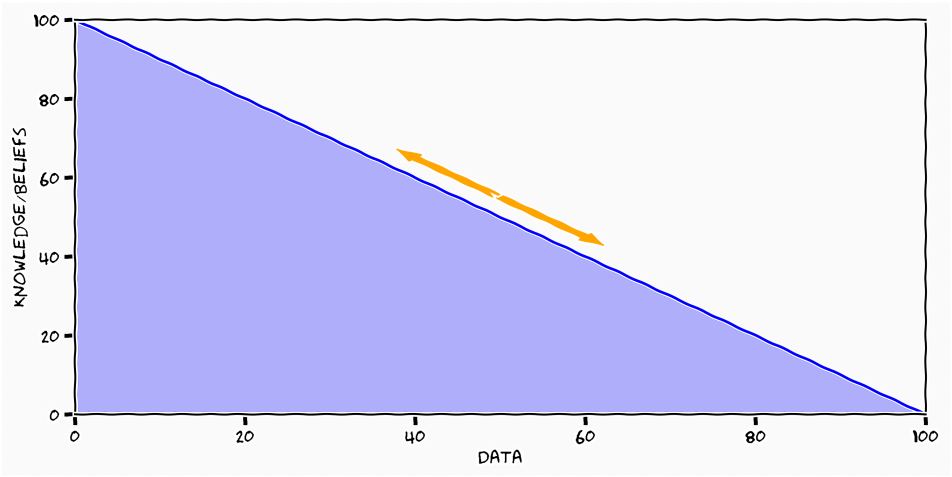

Data and Knowledge

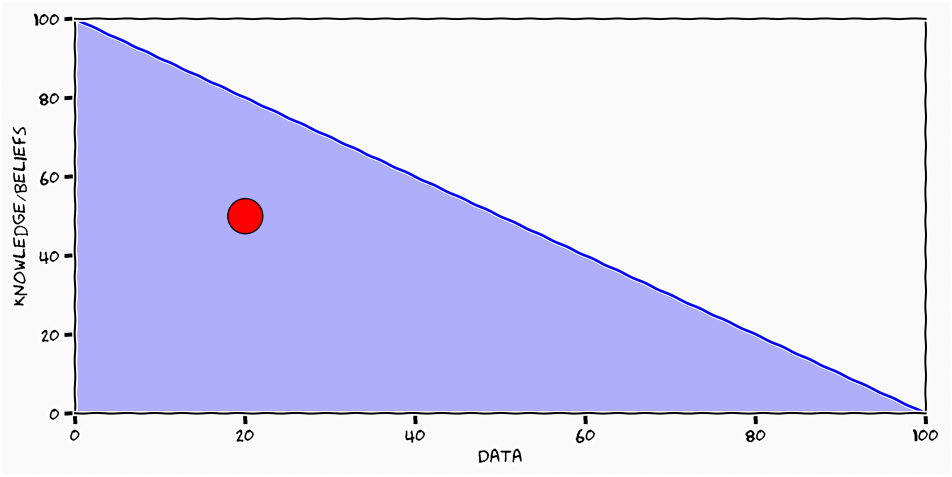

Data and Knowledge

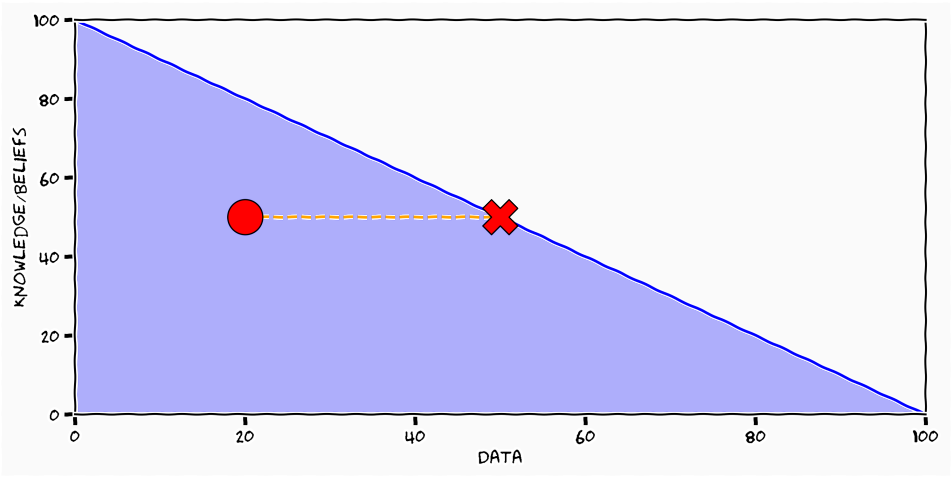

Data and Knowledge

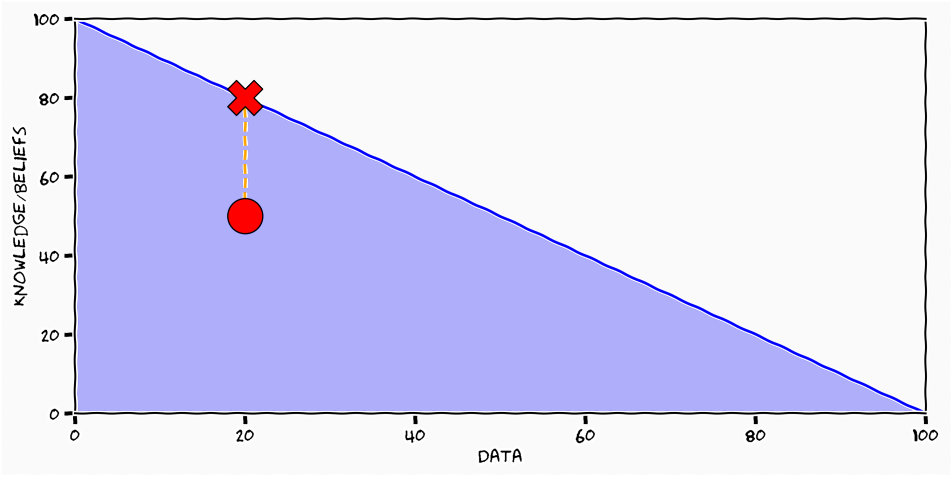

Data and Knowledge

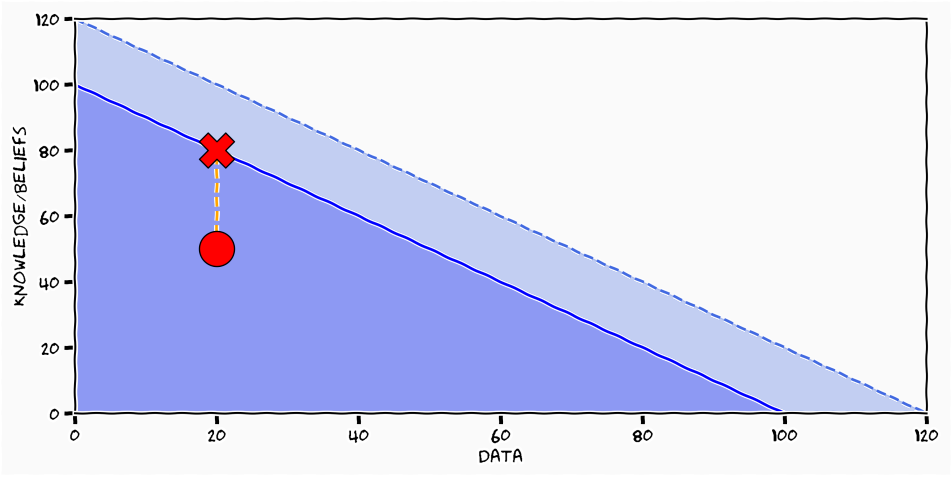

Data and Knowledge

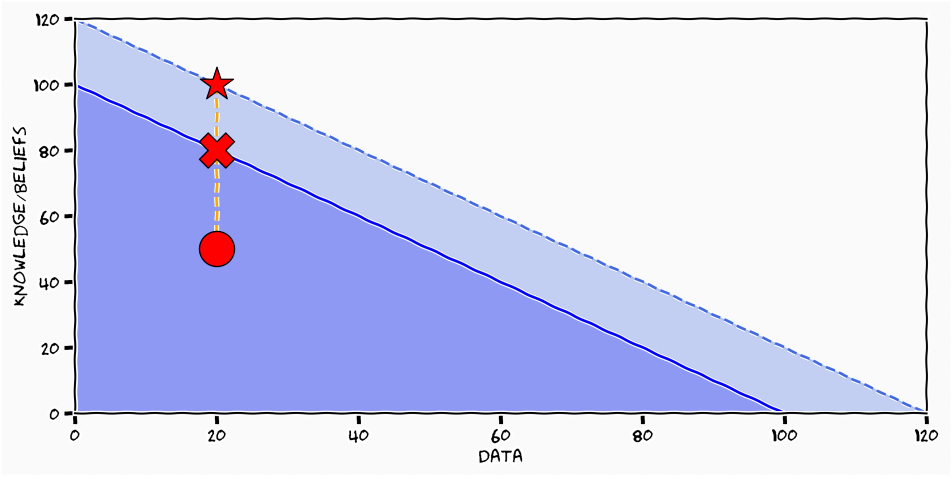

The Scientific Principle

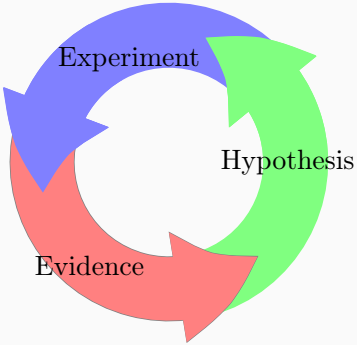

$$p(\theta \mid y) = \frac{p(y \mid \theta)p(\theta)}{\int p(y \mid \theta)p(\theta)\mathrm{d}\theta}$$

Knowledge is Relative

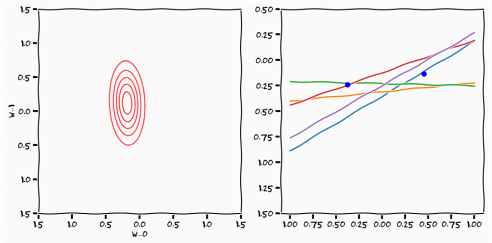

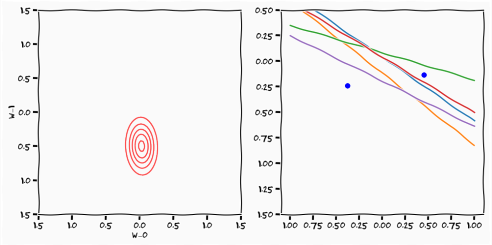

$$\mathcal{A}_{\mathcal{F}}(\mathcal{S})$$

- $\mathcal{F}$ - hypothesis space
- $\mathcal{A}$ - "traverse" mechanism of the hypothesis space
- $\mathcal{S}$ - training data

Statistical Learning Theory in Practice

- Machine Learning as a Framework

$$\text{Data} + \text{Model} \overset{Compute}{\frown\!\!\rightarrow} \text{Prediction}$$

- Machine Learning as a Framework

$$\text{Data} + \text{Model} \overset{Compute}{\rightarrow} \text{Prediction}$$

- Machine Learning as a Science
  - how to construct "handles" to allow us to input knowledge

Norbert Wiener

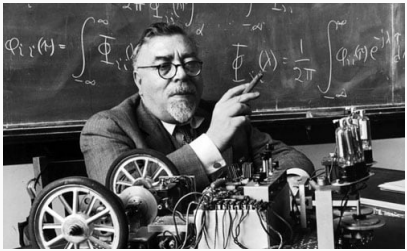

"One of the chief duties of a mathematician in acting as an advisor to scientists is to discourage them from expecting too much of mathematicians."

– Norbert Wiener

Today

- How to specify beliefs over larger classes of hypothesis

Today

- How to specify beliefs over larger classes of hypothesis
- Non-parametrics

Today

- How to specify beliefs over larger classes of hypothesis
- Non-parametrics
- Gaussian processes

Non-parametrics

Hypothesis Classes

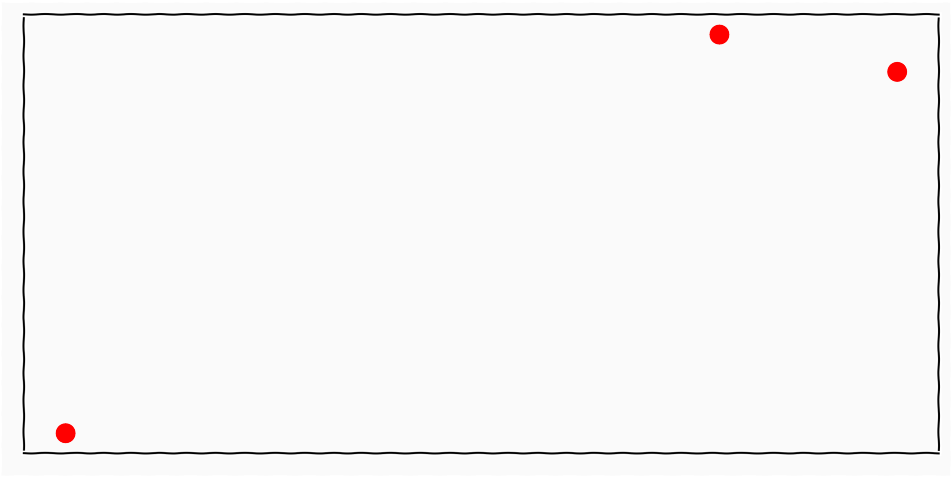

Non-parametrics

Non-parametrics

Non-parametrics

Example

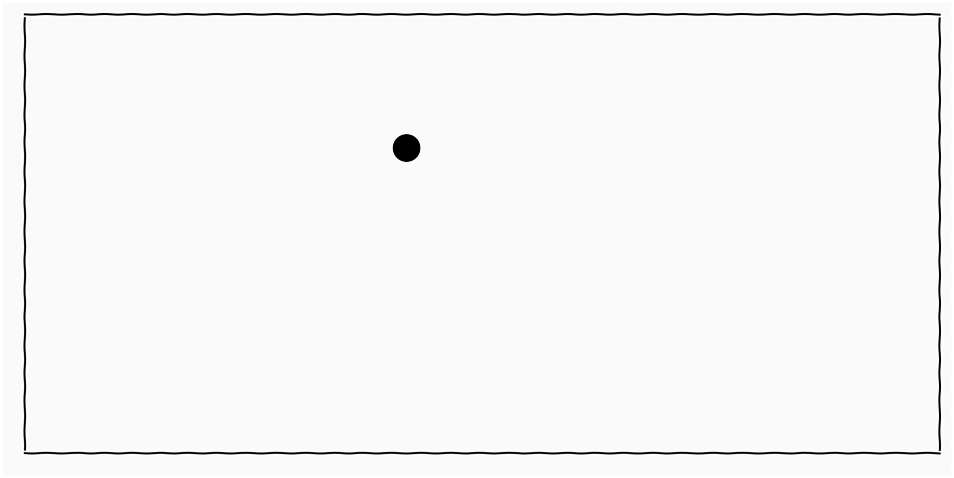

Example

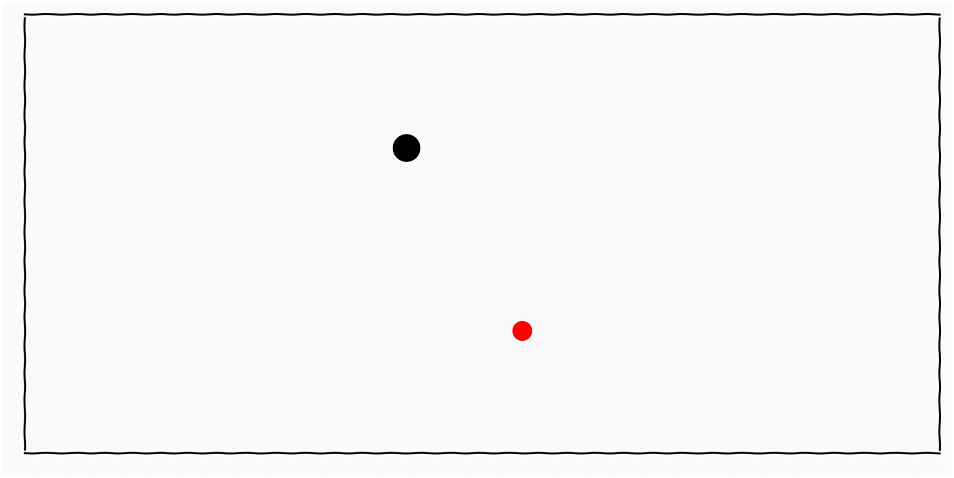

Example

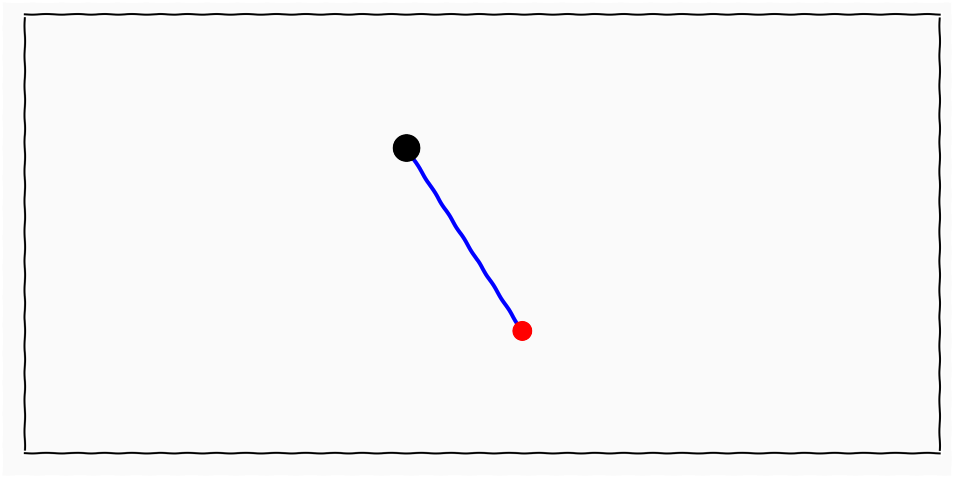

Example

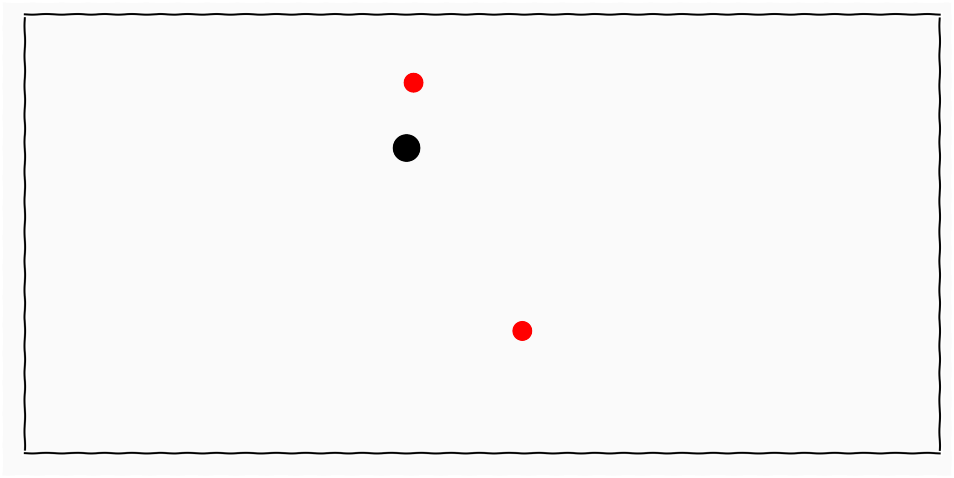

Example

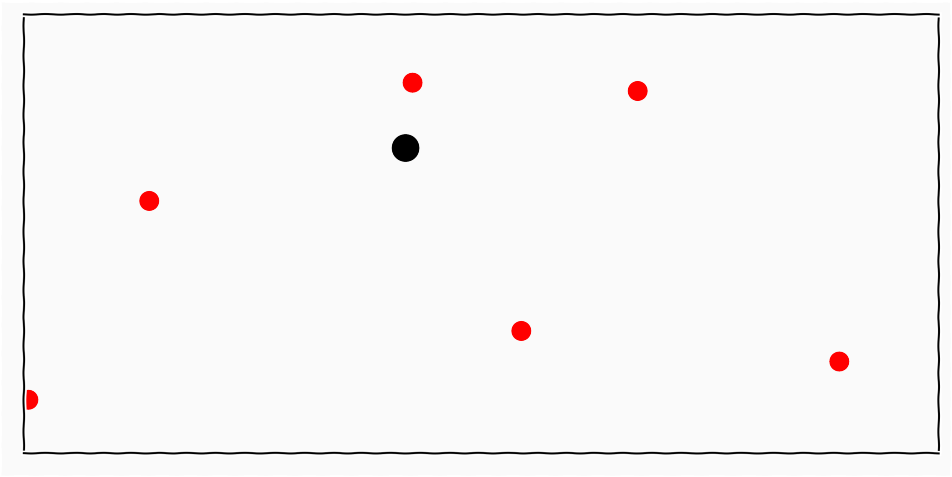

Example

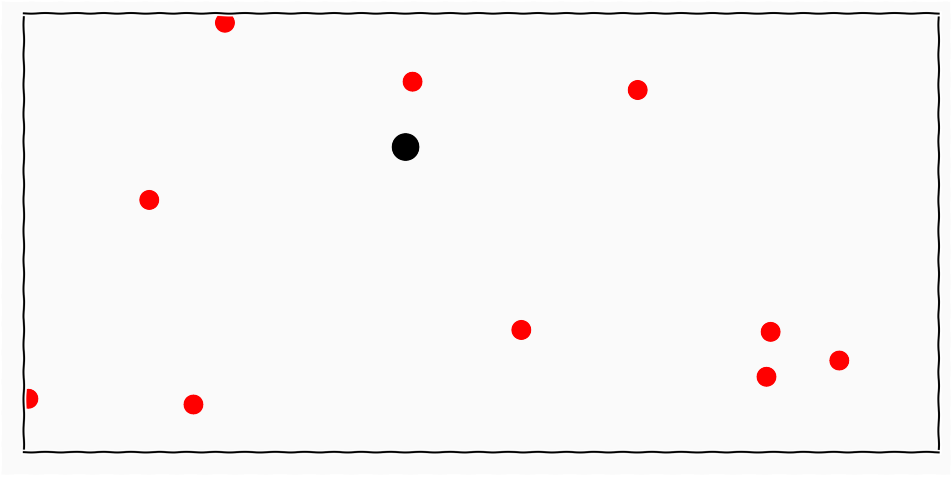

Example

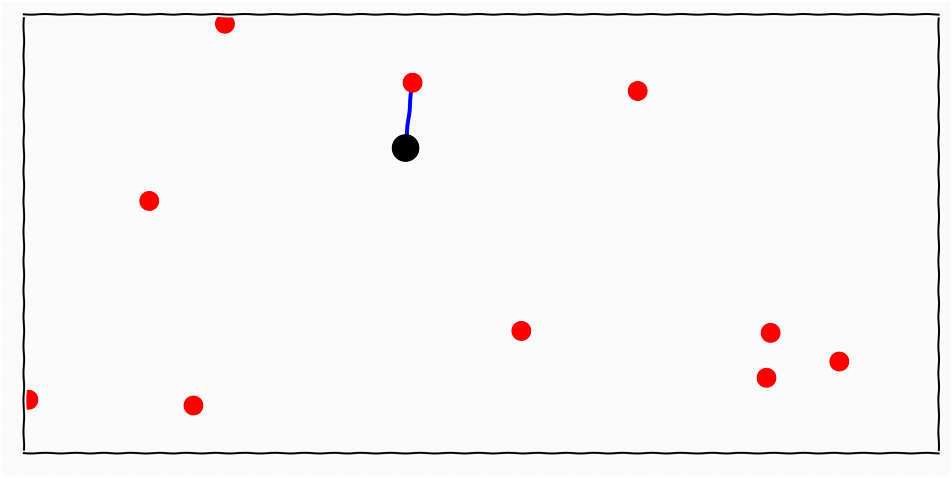

Example

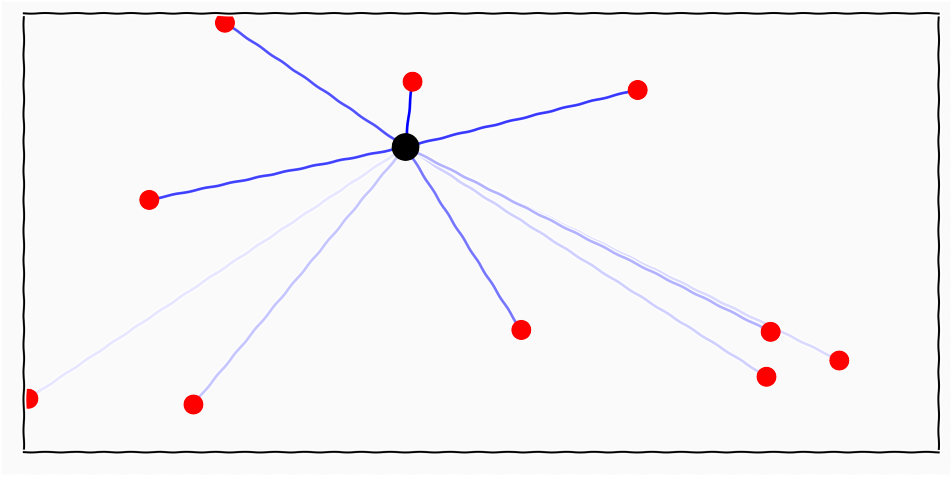

Example

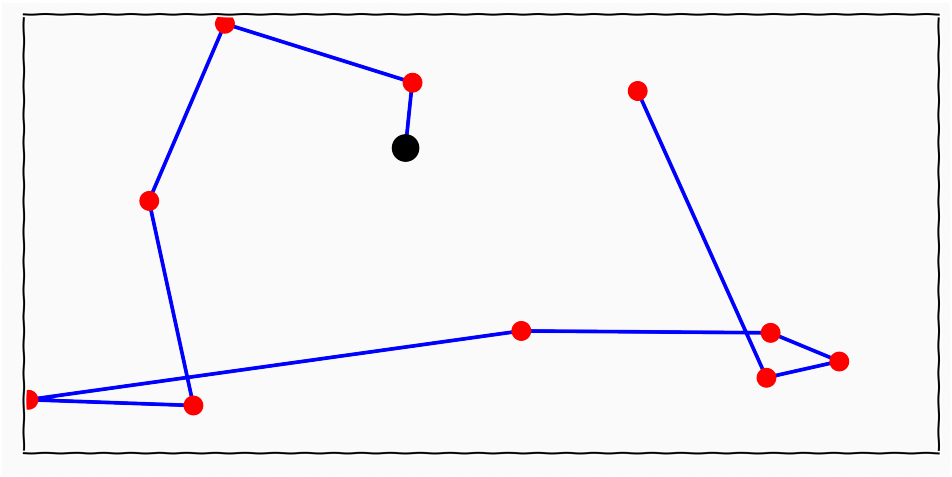

Non-parametric Functions

Lets talk about functions

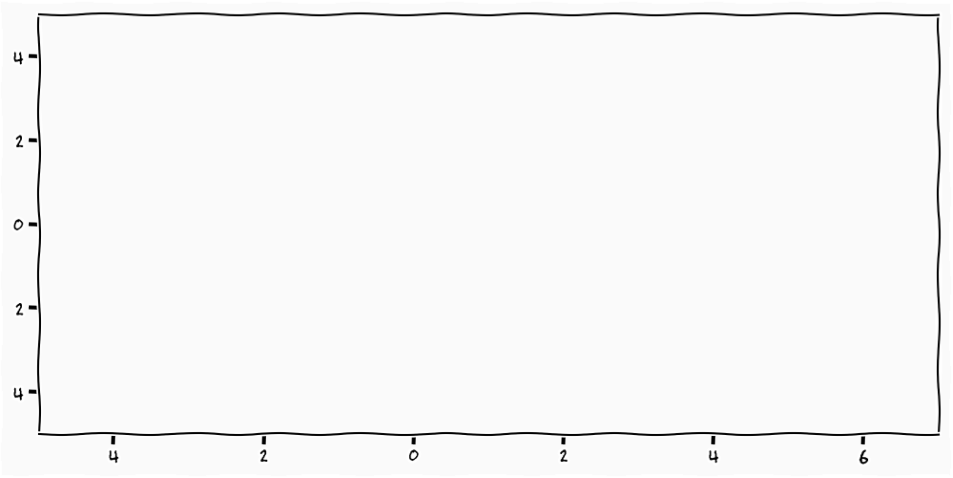

Lets talk about functions

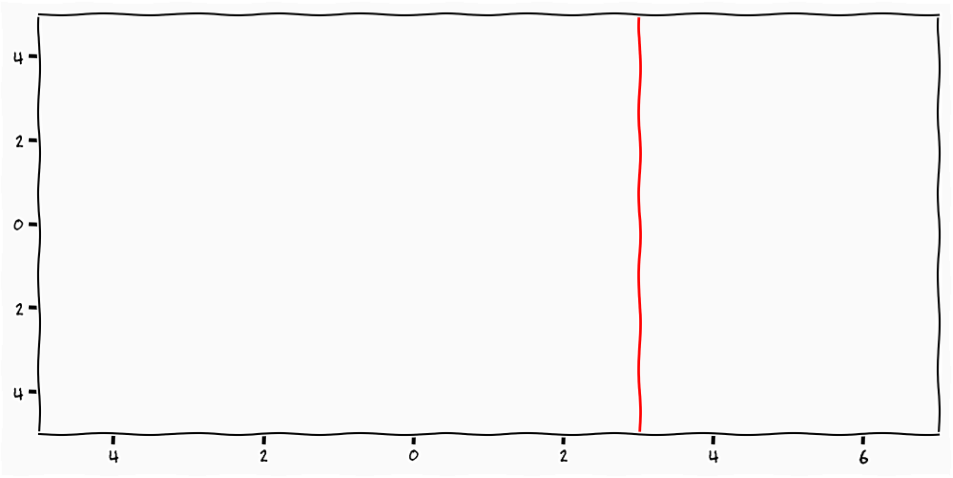

Lets talk about functions

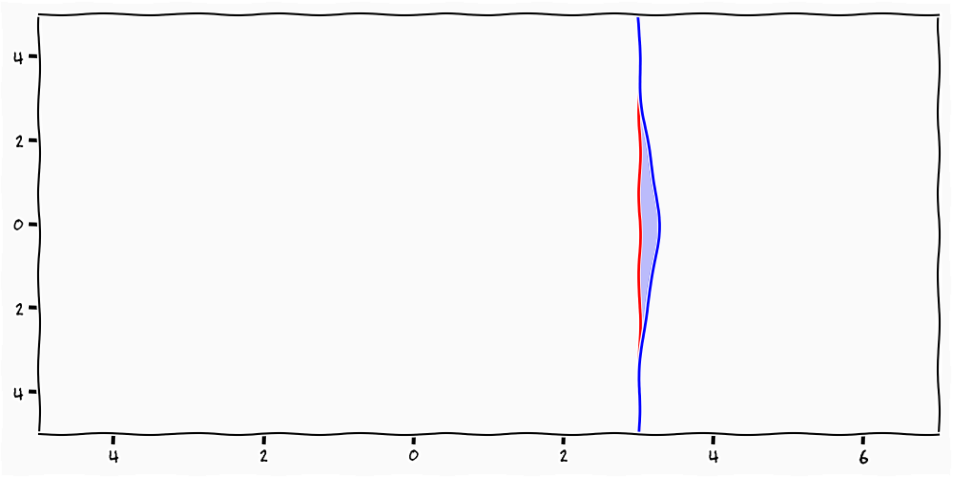

Lets talk about functions

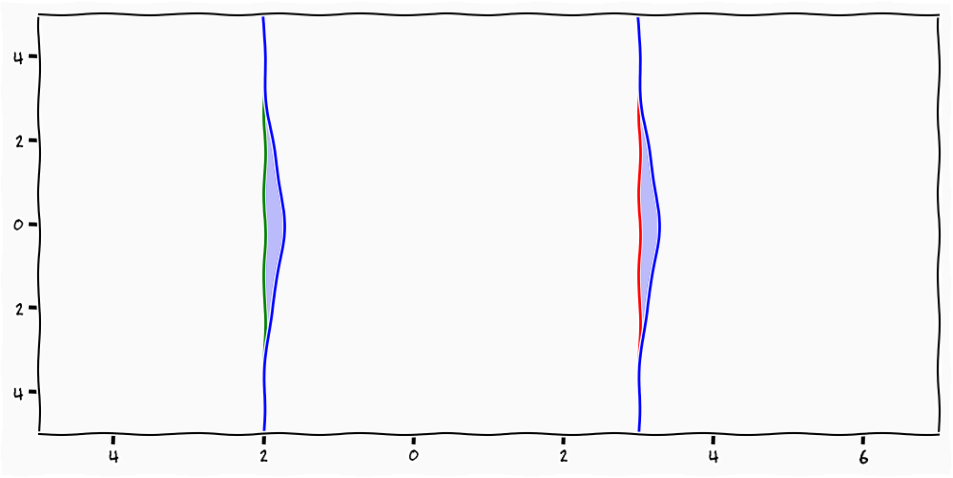

Gaussian function values

$$f_1 = \mathcal{N}(\mu_1, k_1)$$

$$f_1 = \mathcal{N}(\mu_1, k_1)$$

Non-parametric functions

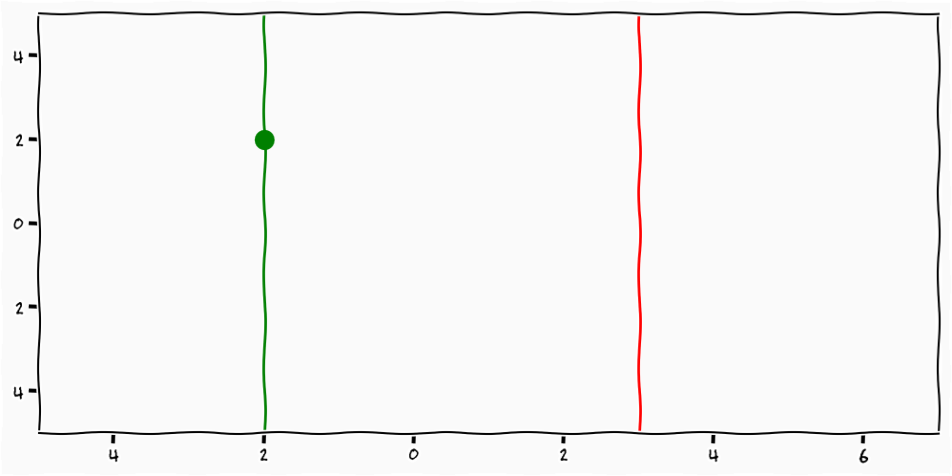

Non-parametric functions

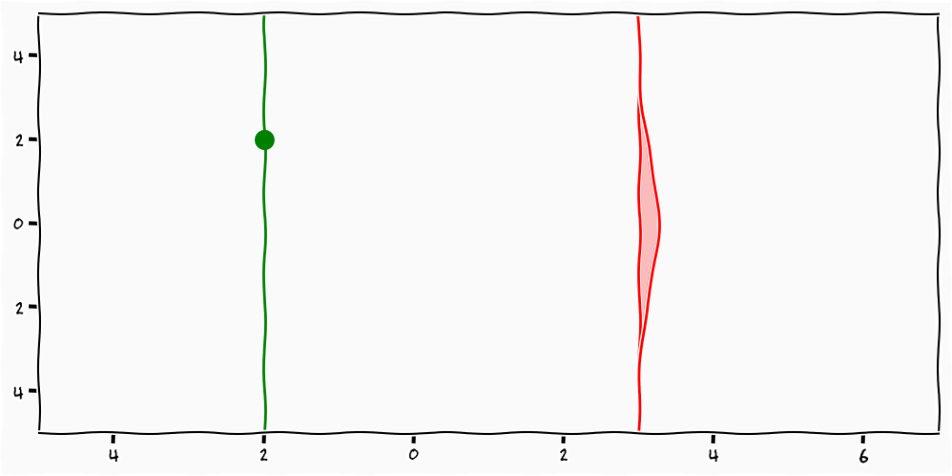

Non-parametric functions

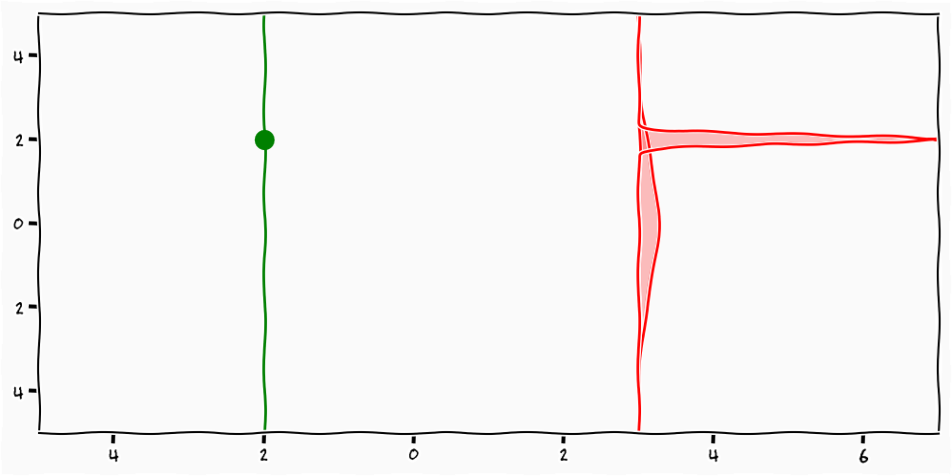

Non-parametric functions

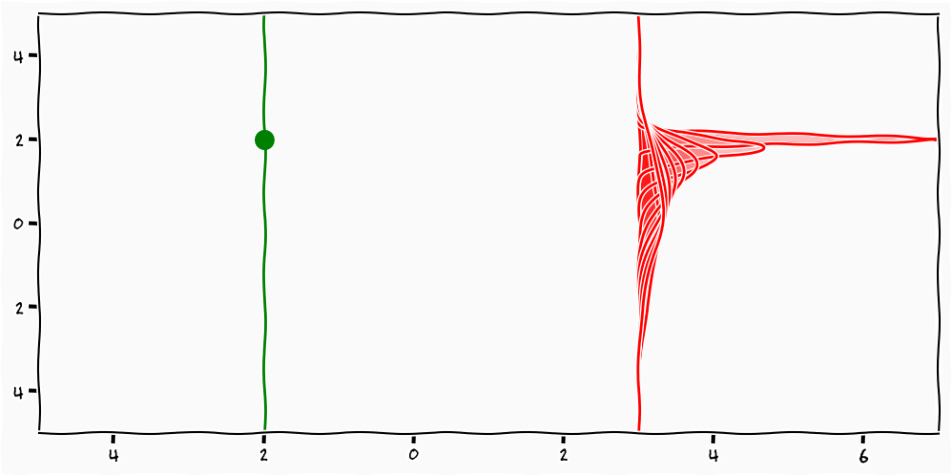

Jointly Gaussian function values

$$\begin{bmatrix} f_1 \\ f_2 \end{bmatrix} = \mathcal{N}\left(\begin{bmatrix} \mu_1 \\ \mu_2 \end{bmatrix}, \begin{bmatrix} k_{11} & ? \\ ? & k_{22} \end{bmatrix} \right)$$

Conditional Gaussians

$$\begin{bmatrix} f_1 \\ f_2 \end{bmatrix} = \mathcal{N}\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 1 & 0.5 \\ 0.5 & 1 \end{bmatrix} \right)$$

Conditional Gaussians

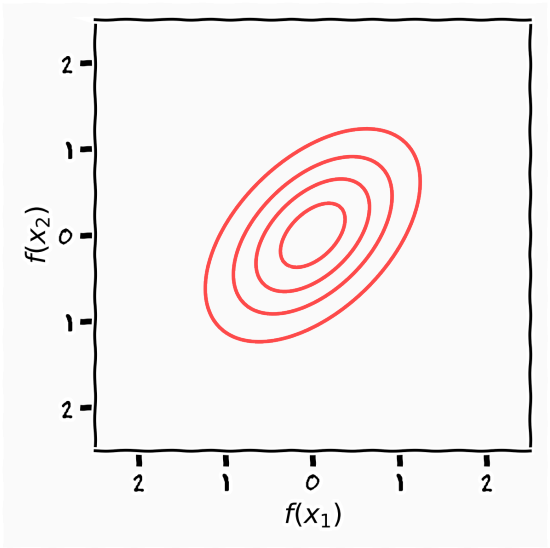

Conditional Gaussians

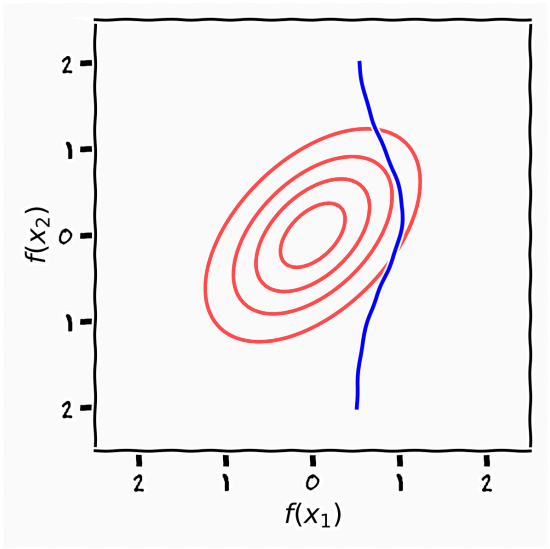

Conditional Gaussians

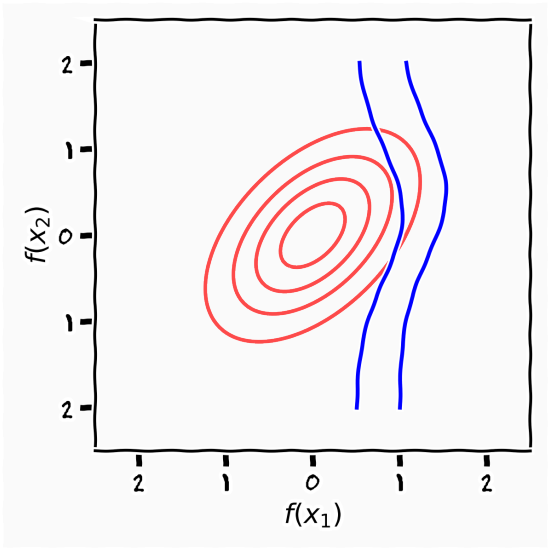

Conditional Gaussians

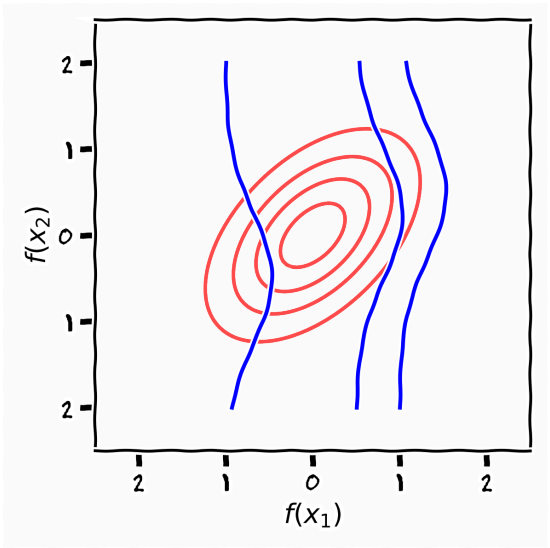

Conditional Gaussians

$$\begin{bmatrix} f_1 \\ f_2 \end{bmatrix} = \mathcal{N}\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 1 & 0.9 \\ 0.9 & 1 \end{bmatrix} \right)$$

Conditional Gaussians

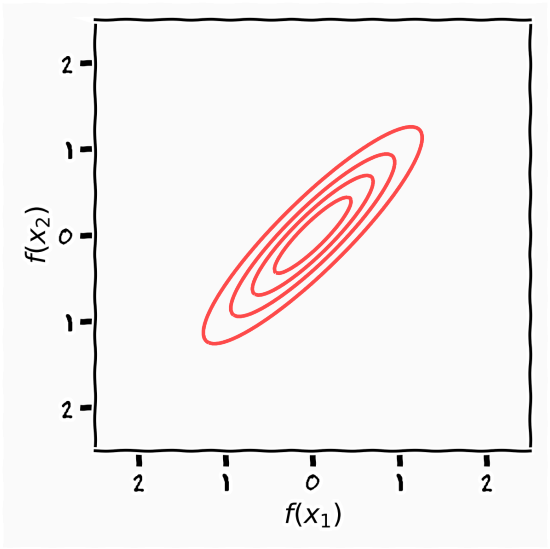

Conditional Gaussians

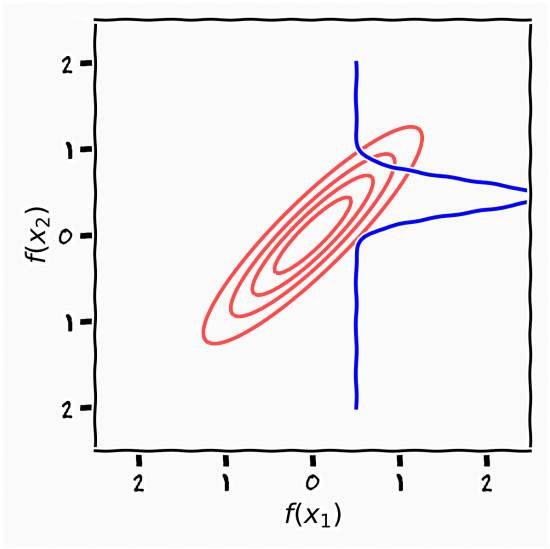

Conditional Gaussians

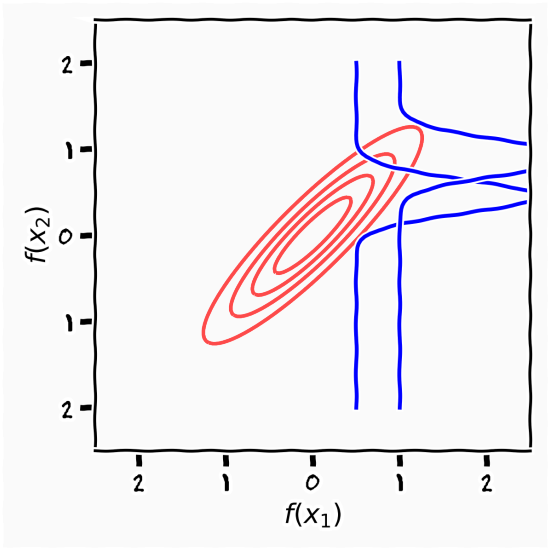

Conditional Gaussians

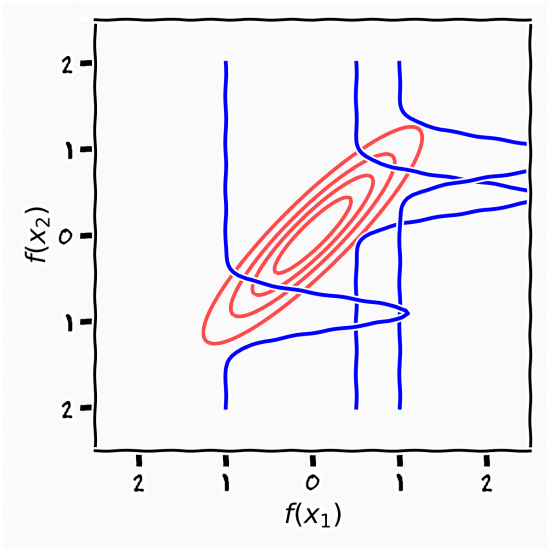

Conditional Gaussians

$$\begin{bmatrix} f_1 \\ f_2 \end{bmatrix} = \mathcal{N}\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \right)$$

Conditional Gaussians

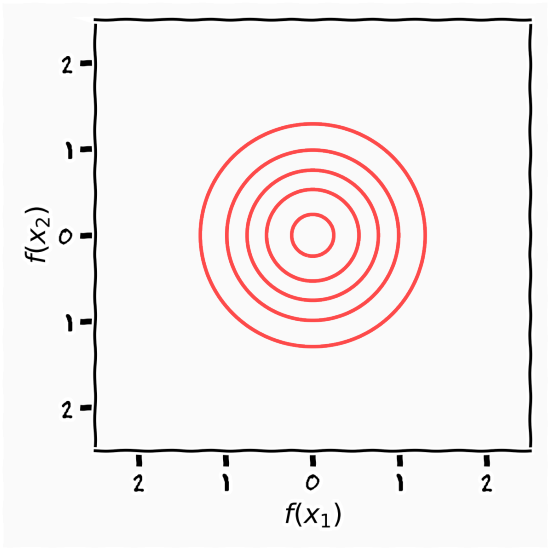

Conditional Gaussians

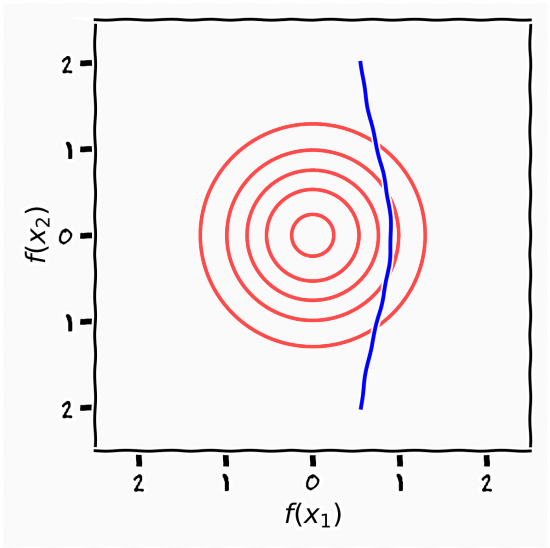

Conditional Gaussians

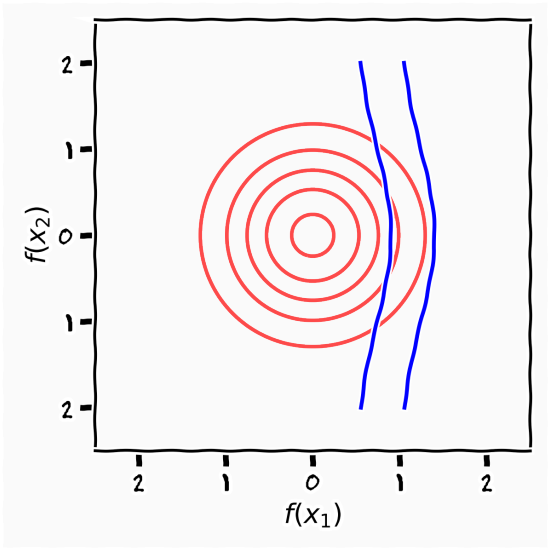

Conditional Gaussians

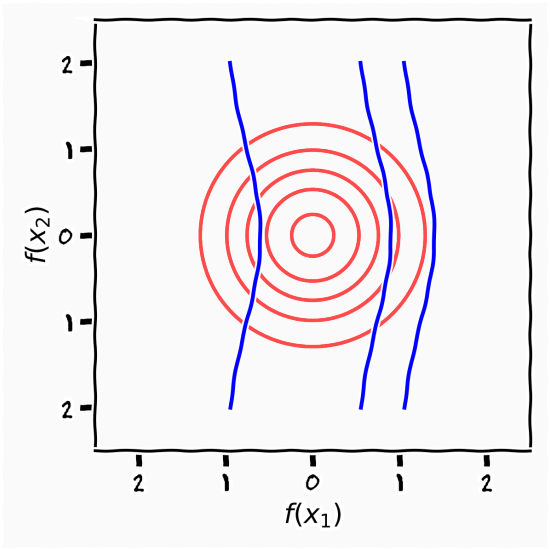

Conditional Gaussians

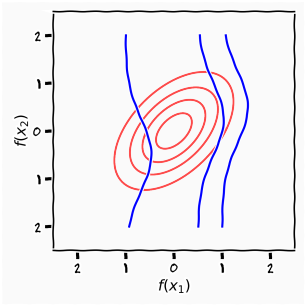

$$\mathsf{N}\left(\begin{bmatrix}0\\0\end{bmatrix}, \begin{bmatrix}1 & 0.5\\0.5 & 1\end{bmatrix}\right)$$

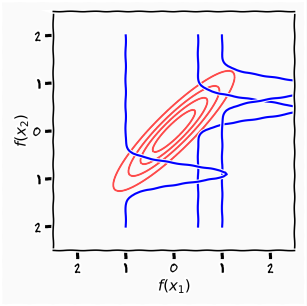

$$\mathsf{N}\left(\begin{bmatrix}0\\0\end{bmatrix}, \begin{bmatrix}1 & 0.9\\0.9 & 1\end{bmatrix}\right)$$

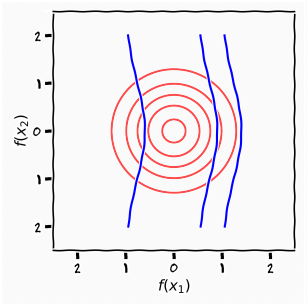

$$\mathsf{N}\left(\begin{bmatrix}0\\0\end{bmatrix}, \begin{bmatrix}1 & 0\\0 & 1\end{bmatrix}\right)$$

Gaussian Samples

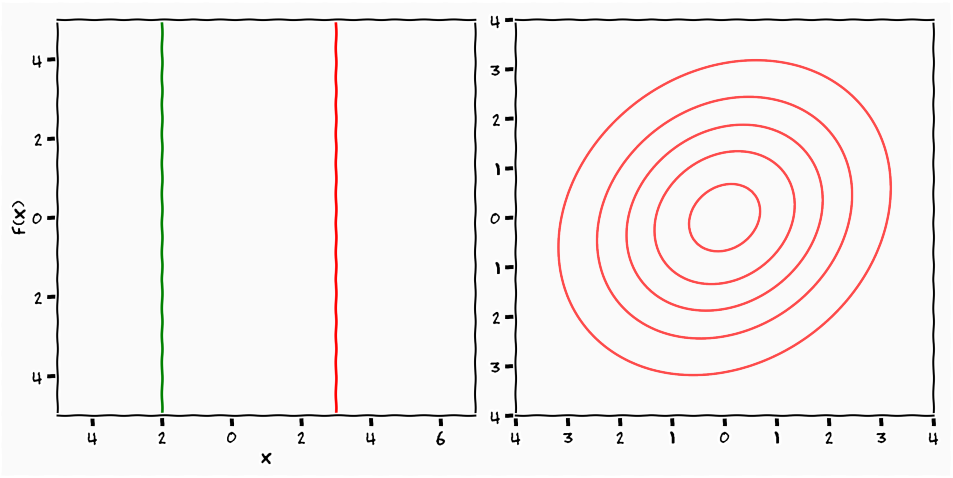

Gaussian Samples

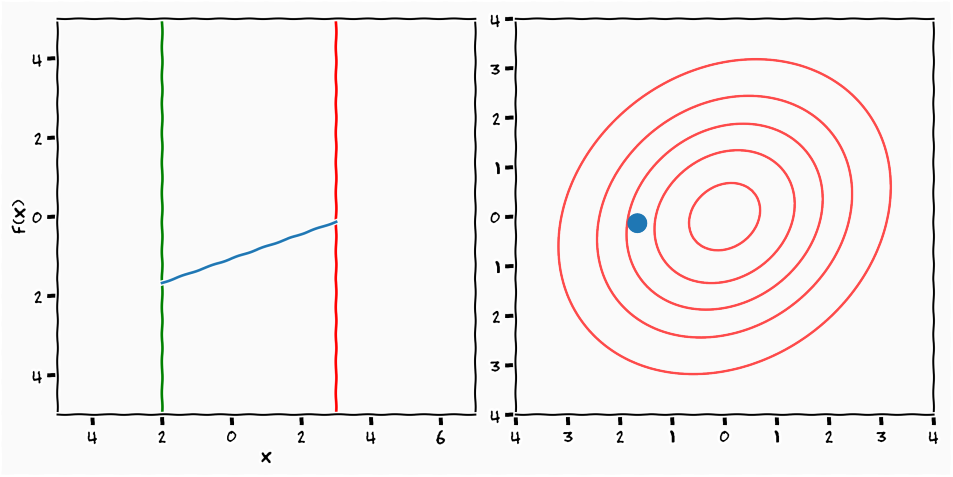

Gaussian Samples

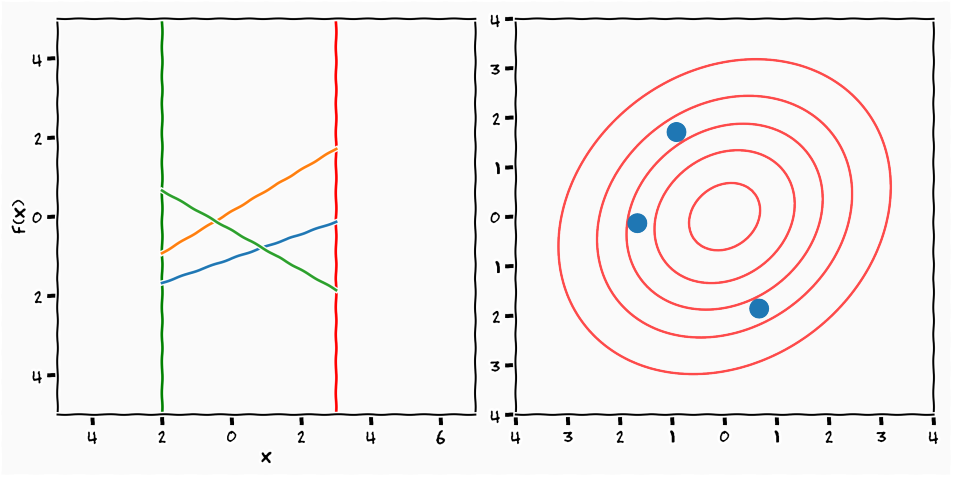

Gaussian Samples

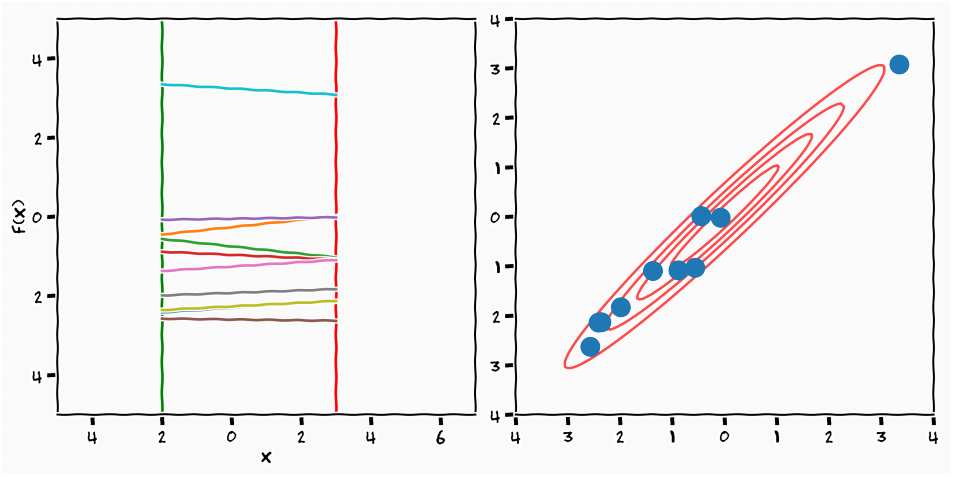

Gaussian Samples

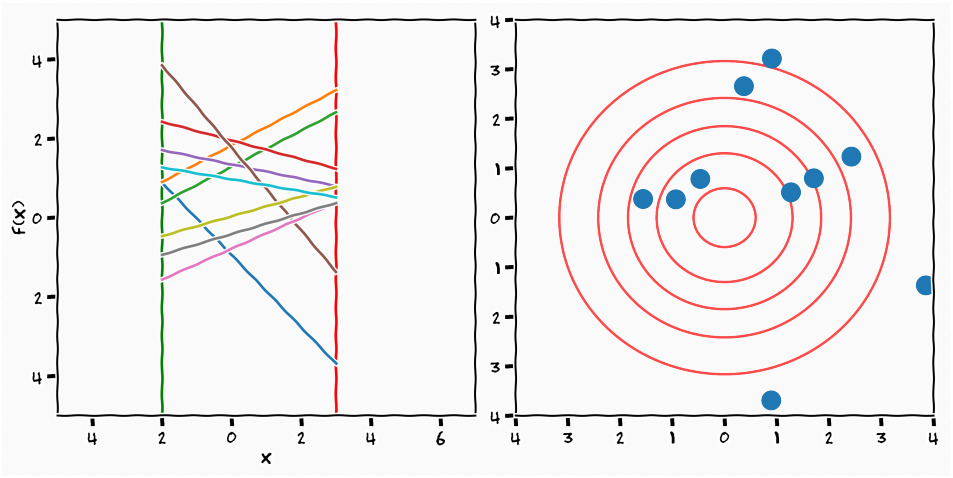

Lets talk about functions

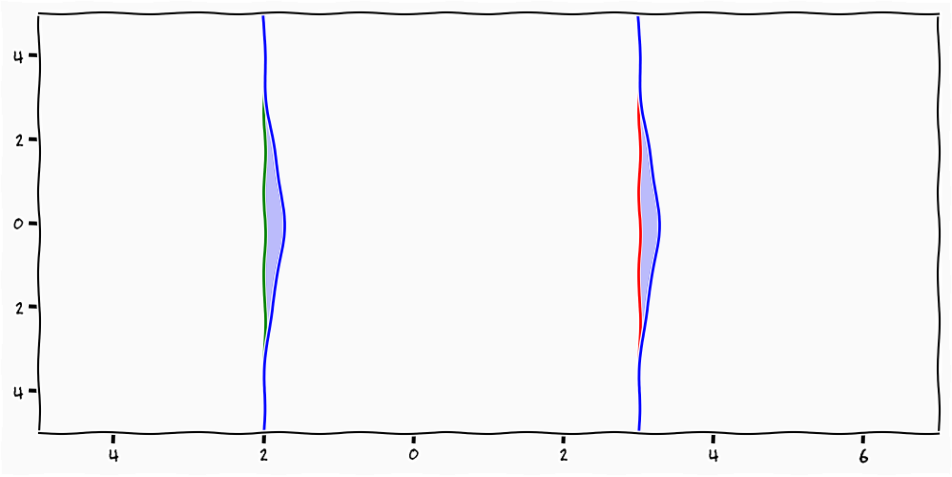

Non-parametric functions

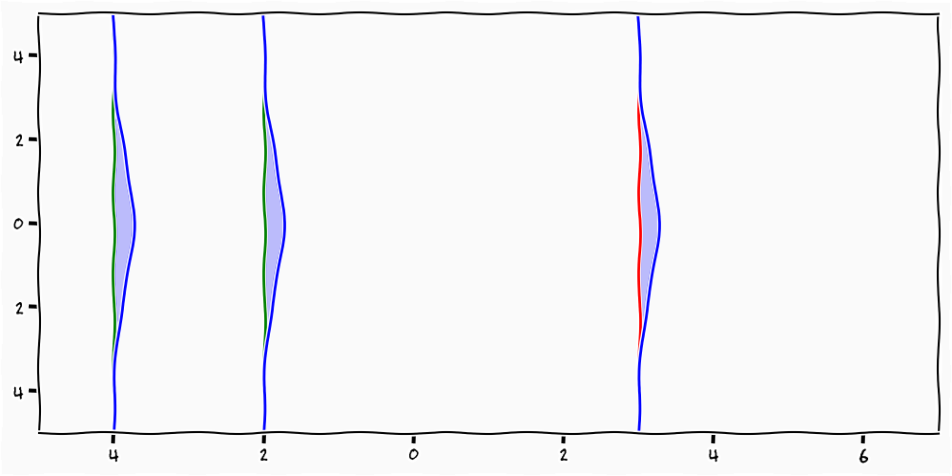

Non-parametric functions

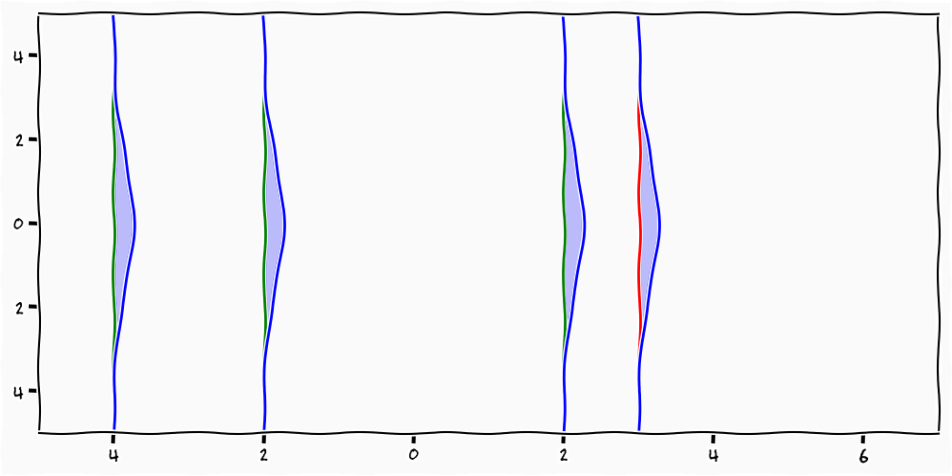

Non-parametric functions

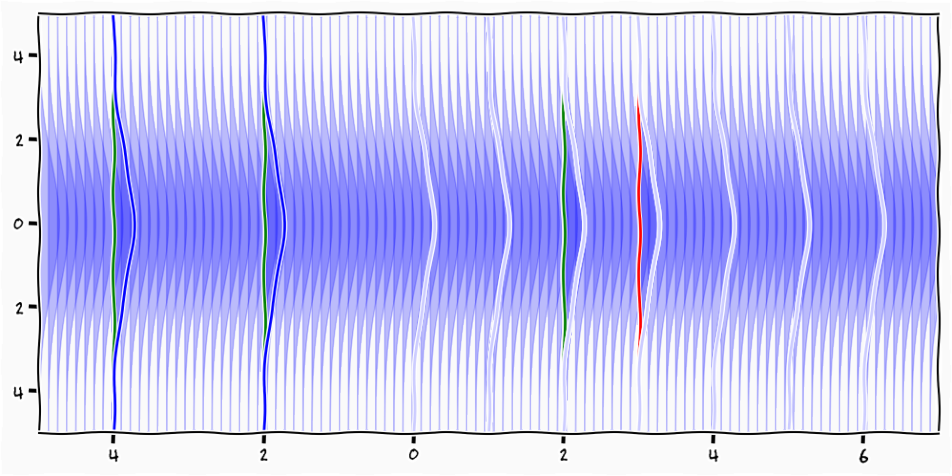

Jointly Gaussian functions II

$$p(\mathbf{f}) = \mathcal{N}\left(\begin{bmatrix} f_1 \\ f_2 \\ \vdots \\ f_N \end{bmatrix} \middle| \begin{bmatrix} \mu_1 \\ \mu_2 \\ \vdots \\ \mu_N \end{bmatrix}, \begin{bmatrix} k_{11} & k_{12} & \dots & k_{1N} \\ k_{21} & k_{22} & \dots & k_{2N} \\ \vdots & \vdots & \ddots & \vdots \\ k_{N1} & k_{N2} & \dots & k_{NN} \end{bmatrix}\right)$$

Gaussian Distribution - Marginal

$$p(x_1, x_2) = \mathcal{N}\left(\begin{array}{c|cc} x_1 & \mu_1 & k_{11} \quad k_{12} \\ x_2 & \mu_2 & k_{21} \quad k_{22} \end{array} \right)$$

$$p(x_1, x_2) = \mathcal{N}\left(\begin{array}{c|cc} x_1 & \mu_1 & k_{11} \quad k_{12} \\ x_2 & \mu_2 & k_{21} \quad k_{22} \end{array} \right)$$

$$\Rightarrow p(x_1) = \int_{x_2} p(x_1, x_2) = \underline{\mathcal{N}\left(x_1 \mid \mu_1, k_{11}\right)}$$

$$p(x_1, x_2) = \mathcal{N}\left(\begin{array}{c|cc} x_1 & \mu_1 & k_{11} \; k_{12} \\ x_2 & \mu_2 & k_{21} \; k_{22} \end{array}\right)$$

$$\Rightarrow p(x_1) = \int_{x_2} p(x_1, x_2) = \underline{\mathcal{N}\left(x_1 \mid \mu_1, k_{11}\right)}$$

$$p(x_1, x_2, \ldots, x_N) = \mathcal{N}\left(\begin{array}{c|ccccc} x_1 & \mu_1 & k_{11} & k_{12} & \cdots & k_{1N} \\ x_2 & \mu_2 & k_{21} & k_{22} & \cdots & k_{2N} \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ x_N & \mu_N & k_{N1} & k_{N2} & \cdots & k_{NN} \end{array}\right)$$

$$p(x_1, x_2) = \mathcal{N}\left(\begin{array}{c|cc} x_1 & \mu_1 & k_{11} \quad k_{12} \\ x_2 & \mu_2 & k_{21} \quad k_{22} \end{array}\right)$$

$$\Rightarrow p(x_1) = \int_{x_2} p(x_1, x_2) = \underline{\mathcal{N}\left(x_1 \mid \mu_1, k_{11}\right)}$$

$$p(x_1, x_2, \ldots, x_N) = \mathcal{N}\left(\begin{array}{c|ccccc} x_1 & \mu_1 & k_{11} & k_{12} & \cdots & k_{1N} \\ x_2 & \mu_2 & k_{21} & k_{22} & \cdots & k_{2N} \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ x_N & \mu_N & k_{N1} & k_{N2} & \cdots & k_{NN} \end{array}\right)$$

$$\Rightarrow \quad p(x_1) = \int_{x_2,\ldots,x_N} p(x_1, x_2, \ldots, x_N) = \underline{\mathcal{N}\left(x_1 \mid \mu_1, k_{11}\right)}$$

Gaussian Distribution - Marginal

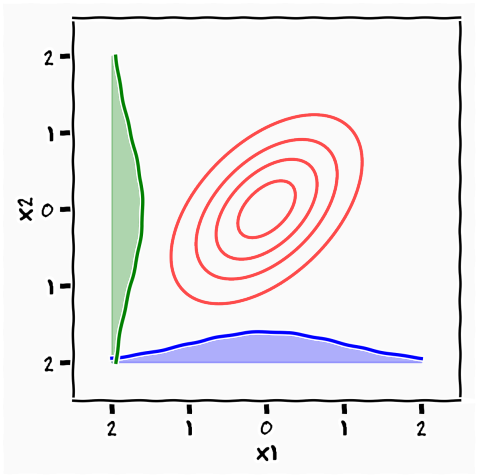

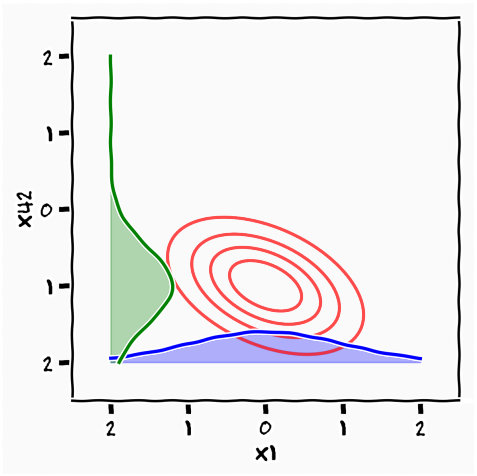

Marginal Property (Consistency)

For all measurable sets $F_i \subseteq \mathbb{R}^n$ and probability measure $\mathcal{N}$

$$\mathcal{N}_{t_1 \cdot t_k}\left(F_1 \times \cdot \times F_k\right) = \mathcal{N}_{t_1 \cdots t_k, t_{k+1} \cdot t_{k+m}}\left(F_1 \times \cdot \times F_k \times \mathbb{R}^n \times \cdot \times \mathbb{R}^n\right)$$

Gaussian Samples

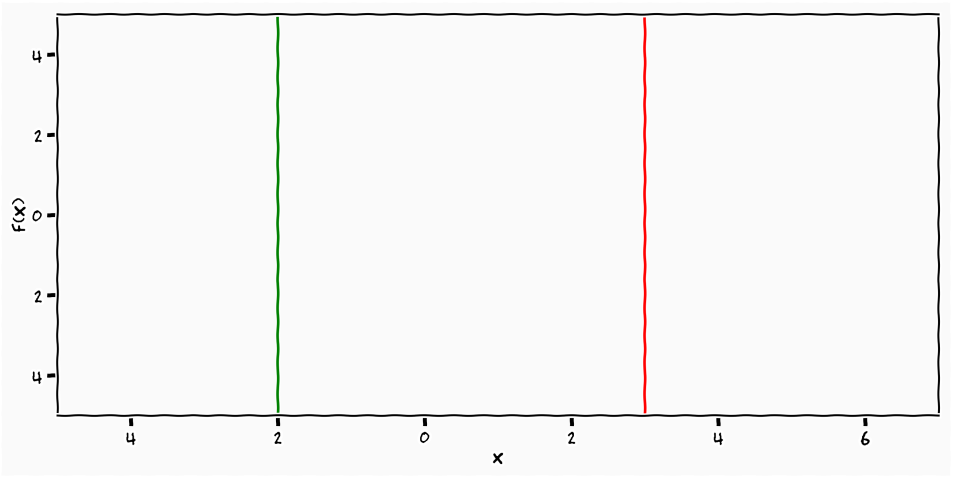

Gaussian Samples

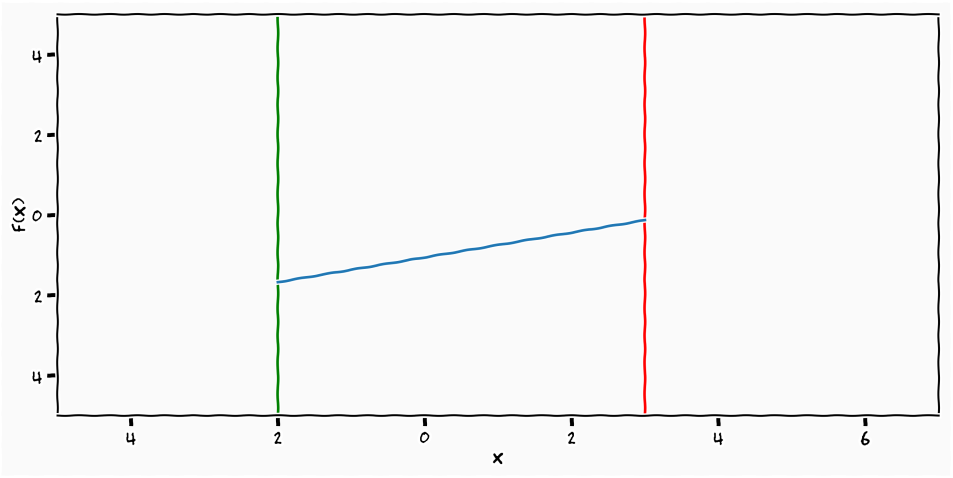

Gaussian Samples

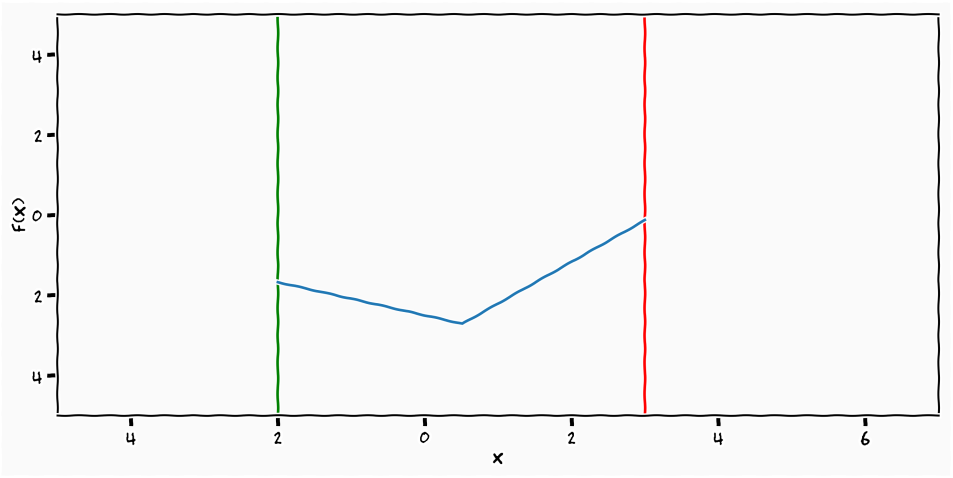

Gaussian Samples

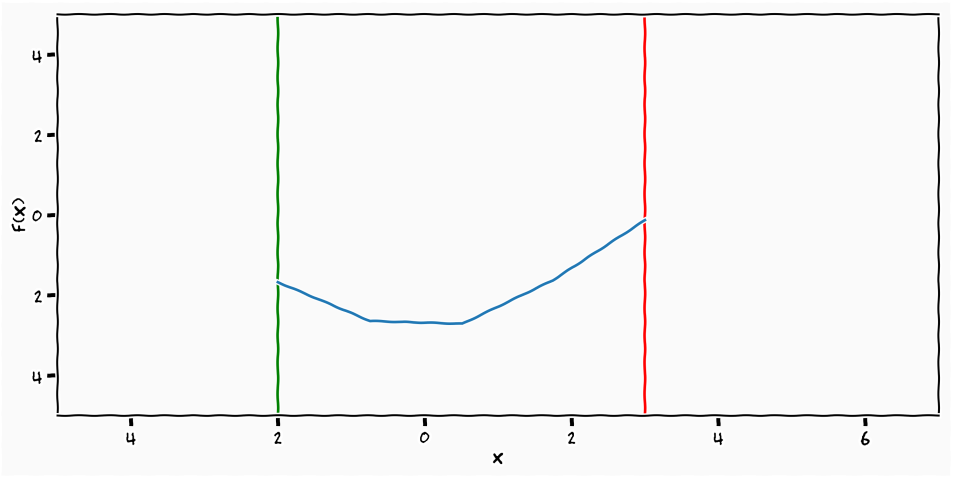

Gaussian Samples

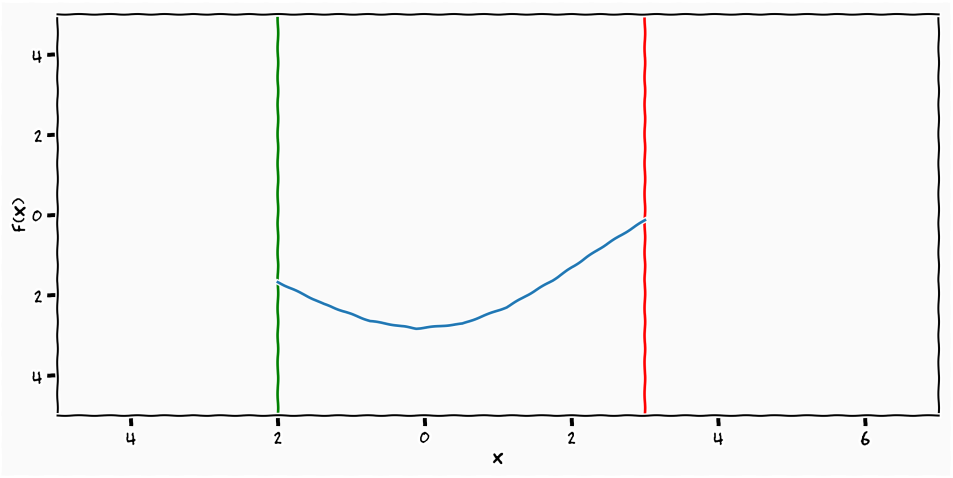

Gaussian Samples

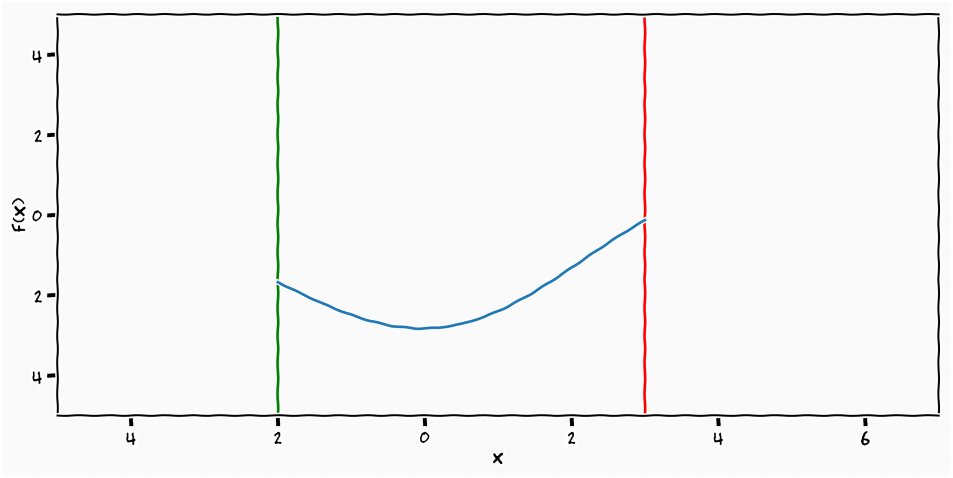

Gaussian Samples

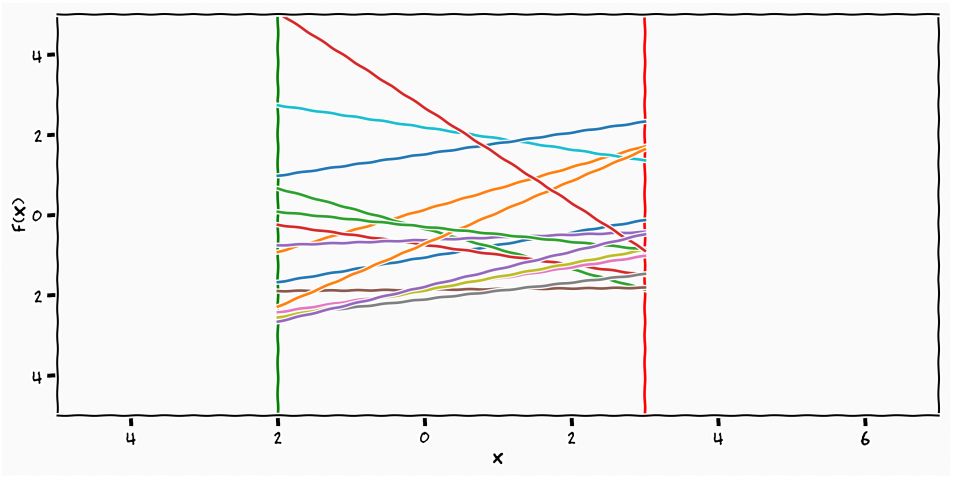

Gaussian Samples

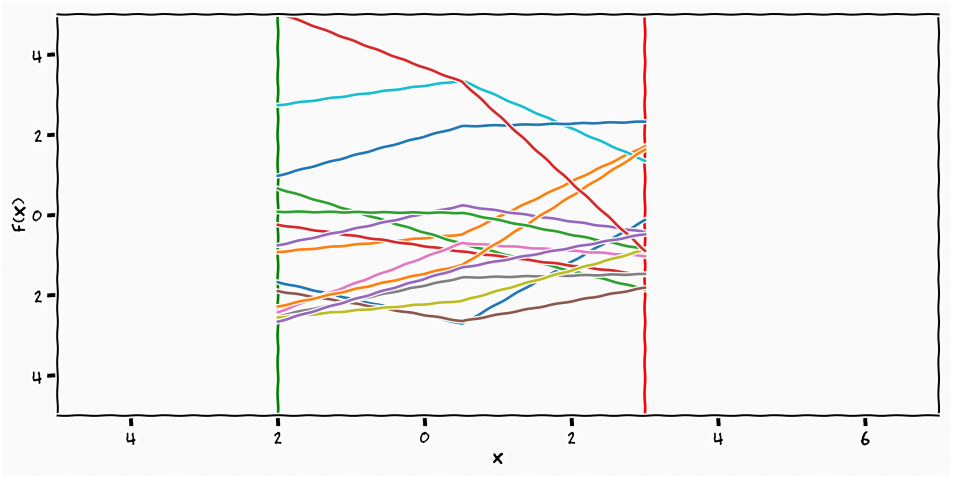

Gaussian Samples

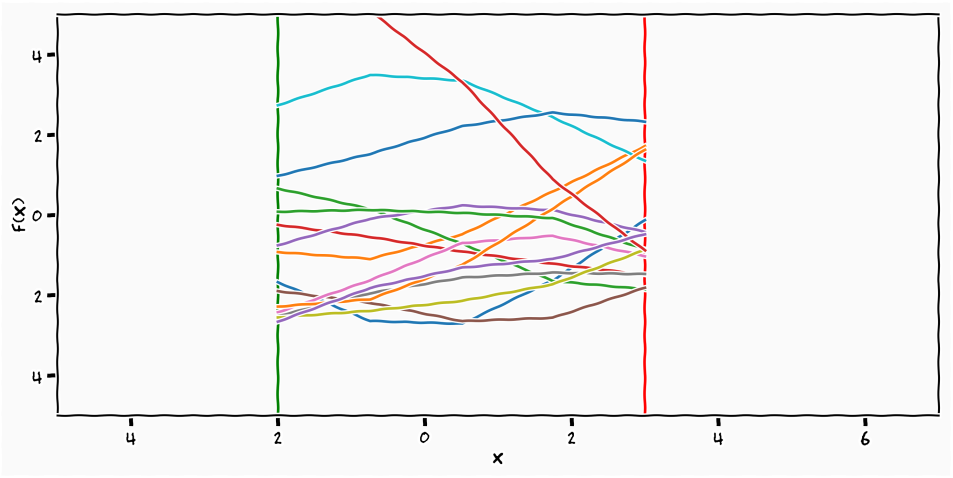

Gaussian Samples

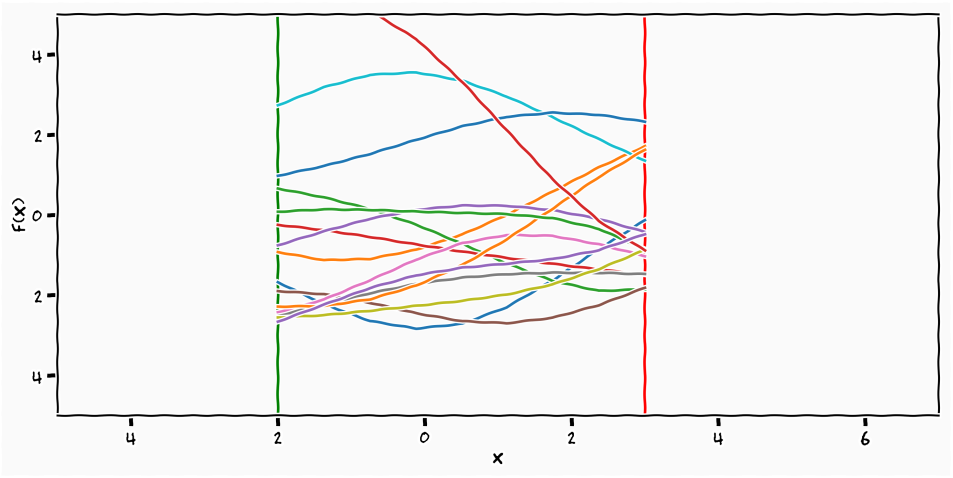

Gaussian Samples

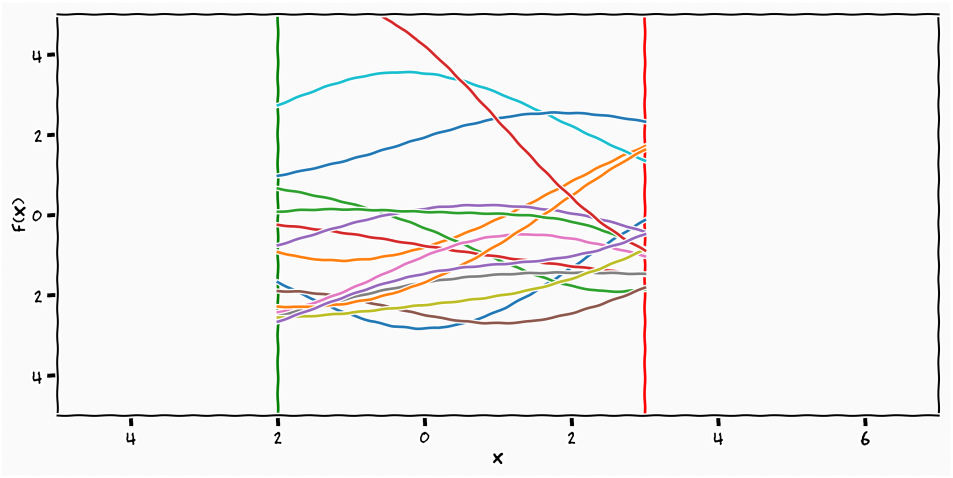

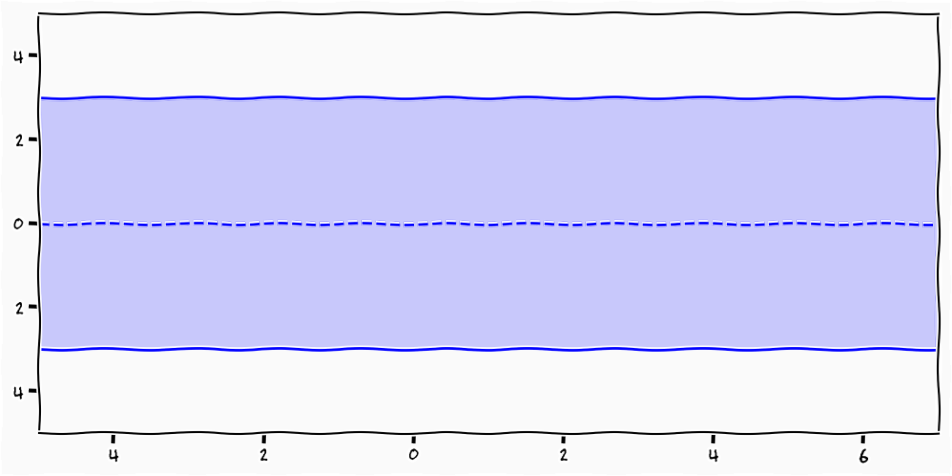
4
2
0
2
4
4
2
0
2
4
6

$$p(\mathbf{f}) = \mathcal{N}\left(\begin{bmatrix} f_1 \\ f_2 \\ \vdots \\ f_N \\ \vdots \end{bmatrix} \middle\| \begin{bmatrix} \mu_1 \\ \mu_2 \\ \vdots \\ \mu_N \\ \vdots \end{bmatrix}, \begin{bmatrix} k_{11} & k_{12} & \dots & k_{1N} & \dots \\ k_{21} & k_{22} & \dots & k_{2N} & \dots \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ k_{N1} & k_{N2} & \dots & k_{NN} & \dots \\ \vdots & \vdots & \dots & \vdots & \ddots \end{bmatrix}\right)$$

$$\mathcal{GP}(\cdot,\cdot) \qquad M \in \mathbb{R}^{\infty \times N} \qquad \mathcal{N}(\cdot,\cdot)$$

$$\rightarrow$$

$$\infty \qquad\qquad N$$

The Gaussian distribution is the projection of the infinite Gaussian process

Definition (Gaussian Process)
A Gaussian process is a collection of random variables who are jointly Gaussian distributed index by a infinite index set

$$p(\mathbf{f}) = \mathcal{N}\left(\begin{bmatrix} f_1 \\ f_2 \\ \vdots \\ f_N \\ \vdots \end{bmatrix} \middle| \begin{bmatrix} \mu(x_1) \\ \mu(x_2) \\ \vdots \\ \mu(x_N) \\ \vdots \end{bmatrix}, \begin{bmatrix} k(x_1,x_1) & k(x_1,x_2) & \dots & k(x_1,x_N) & \dots \\ k(x_2,x_1) & k(x_2,x_2) & \dots & k(x_2,x_N) & \dots \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ k(x_N,x_1) & k(x_N,x_2) & \dots & k(x_N,x_N) & \dots \\ \vdots & \vdots & \dots & \vdots & \ddots \end{bmatrix}\right)$$

"Parametrisation"

$$k_{ij} = k(x_i, x_j)$$

- We parameterise the covariance as a function of the input

"Parametrisation"

$$k_{ij} = k(x_i, x_j)$$

- We parameterise the covariance as a function of the input
- the index set of the measure is the uncountable infinity

"Parametrisation"

$$k_{ij} = k(x_i, x_j)$$

- We parameterise the covariance as a function of the input
- the index set of the measure is the uncountable infinity
- Your "handle" to input your knowledge into a GP is the covariance function

"Parametrisation"

$$k_{ij} = k(x_i, x_j)$$

- We parameterise the covariance as a function of the input
- the index set of the measure is the uncountable infinity
- Your "handle" to input your knowledge into a GP is the covariance function
  - *you specify the degree of covariance between data-points*

"Parametrisation"

$$k_{ij} = k(x_i, x_j)$$

- We parameterise the covariance as a function of the input
- the index set of the measure is the uncountable infinity
- Your "handle" to input your knowledge into a GP is the covariance function
  - *you specify the degree of covariance between data-points*
- If this "parametrisation" aligns well with your knowledge a GP is the way forward!

Gaussian Processes

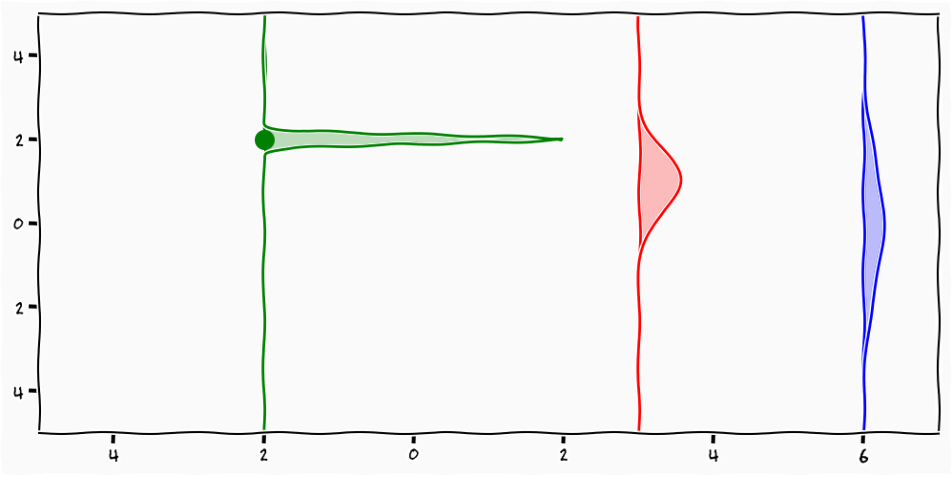

Gaussian Processes Samples

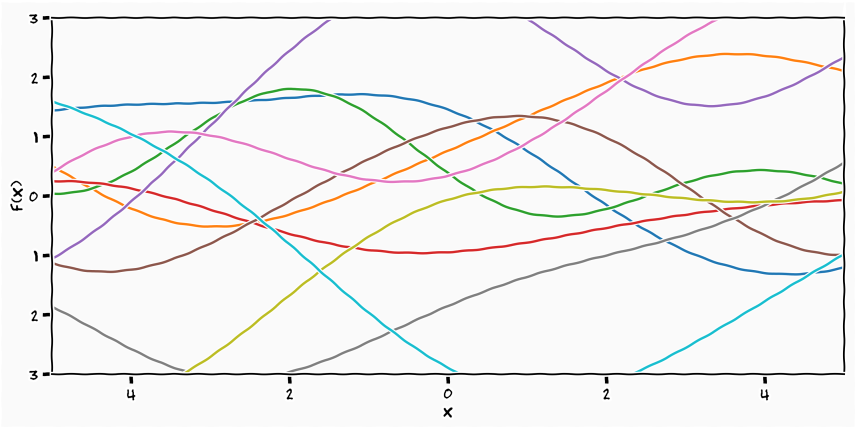

$$k(x_i, x_j) = 3 \cdot e^{-\frac{(x_i - x_j)^2}{15}}$$

Gaussian Processes Samples

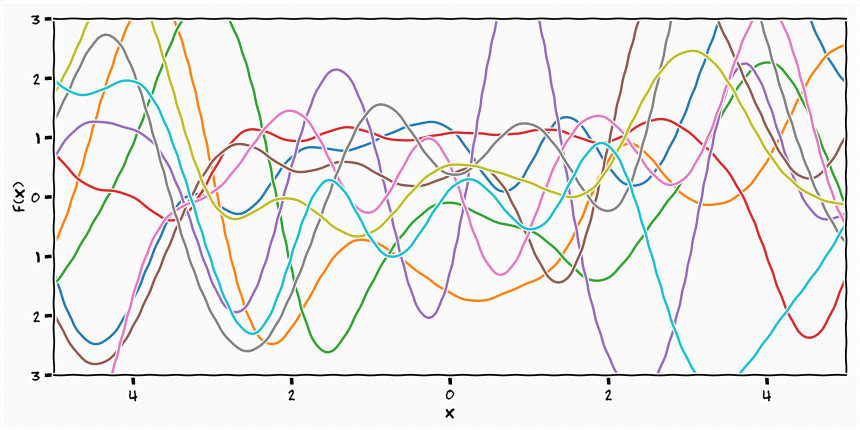

$$k(x_i, x_j) = 3 \cdot e^{-\frac{(x_i - x_j)^2}{1}}$$

Gaussian Processes Samples

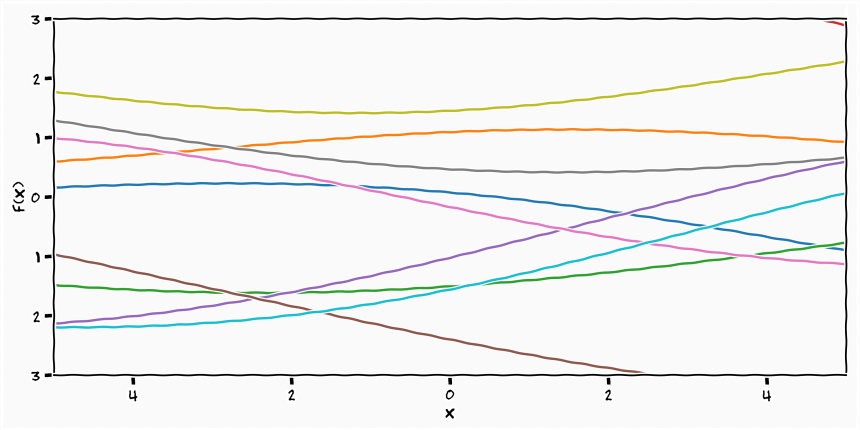

$$k(x_i, x_j) = 3 \cdot e^{-\frac{(x_i - x_j)^2}{150}}$$

Gaussian Processes

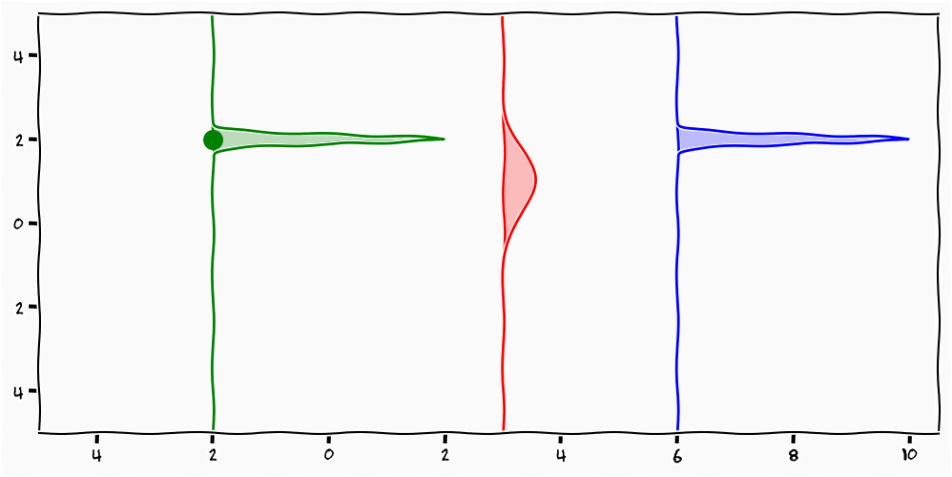

Bayesian Inference

Bayes' Rule

$$p(\mathbf{f}_* \mid \mathbf{f}) = \frac{p(\mathbf{f}, \mathbf{f}_*)}{p(\mathbf{f})} = \frac{p(\mathbf{f}, \mathbf{f}_*)}{\int p(\mathbf{f}, \mathbf{f}_*) \mathrm{d}\mathbf{f}_*}$$

Marginal Likelihood

$$\int p(\mathbf{f}, \mathbf{f}_*) \mathrm{d}\mathbf{f}_* = \int p(\mathbf{f} \mid \mathbf{f}_*) p(\mathbf{f}_*) \mathrm{d}\mathbf{f}_*$$

- Take every possible function value/marginal $\mathbf{f}_*$ at location $\mathbf{x}_*$ according to their probability

$$\int p(\mathbf{f}, \mathbf{f}_*) \mathrm{d}\mathbf{f}_* = \int p(\mathbf{f} \mid \mathbf{f}_*) p(\mathbf{f}_*) \mathrm{d}\mathbf{f}_*$$

- Take every possible function value/marginal $\mathbf{f}_*$ at location $\mathbf{x}_*$ according to their probability
- Check if these marginals are consistent with the marginals we observe $\mathbf{f}$ at location $\mathbf{x}$

Gaussian Processes: Posterior Samples

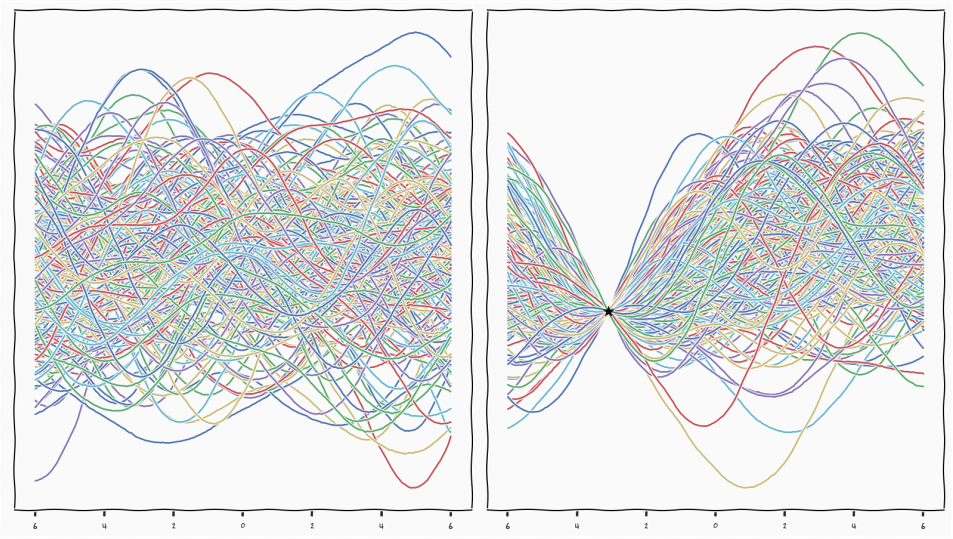

Gaussian Process: "Predictive Posterior"

$$p(\mathbf{f}, \mathbf{f}_*) = p(\mathbf{f}_* \mid \mathbf{f})p(\mathbf{f})$$

- We have defined $p(\mathbf{f}, \mathbf{f}_*)$

Gaussian Process: "Predictive Posterior"

$$p(\mathbf{f}, \mathbf{f}_*) = p(\mathbf{f}_* \mid \mathbf{f})p(\mathbf{f})$$

- We have defined $p(\mathbf{f}, \mathbf{f}_*)$
- We know through the marginal property of the Gaussian that $p(\mathbf{f})$ is consistent

$$p(\mathbf{f}, \mathbf{f}_*) = p(\mathbf{f}_* \mid \mathbf{f})p(\mathbf{f})$$

- We have defined $p(\mathbf{f}, \mathbf{f}_*)$
- We know through the marginal property of the Gaussian that $p(\mathbf{f})$ is consistent
- We know that $p(\mathbf{f}_* \mid \mathbf{f})$ is Gaussian

Gaussian Process: "Predictive Posterior"

$$p(\mathbf{f}, \mathbf{f}_*) = p(\mathbf{f}_* \mid \mathbf{f})p(\mathbf{f})$$

- We have defined $p(\mathbf{f}, \mathbf{f}_*)$
- We know through the marginal property of the Gaussian that $p(\mathbf{f})$ is consistent
- We know that $p(\mathbf{f}_* \mid \mathbf{f})$ is Gaussian
- $\Rightarrow$ *We can just solve for* $p(\mathbf{f}_* \mid \mathbf{f})$

- All instantiations are jointly Gaussian

$$\begin{bmatrix} \mathbf{f} \\ \mathbf{f}_* \end{bmatrix} \sim \mathcal{N}\left(\begin{bmatrix} \mathbf{0} \\ \mathbf{0} \end{bmatrix}, \begin{bmatrix} k(\mathbf{x}, \mathbf{x}) & k(\mathbf{x}, \mathbf{x}_*) \\ k(\mathbf{x}_*, \mathbf{x}) & k(\mathbf{x}_*, \mathbf{x}_*) \end{bmatrix}\right)$$

- Conditional Gaussian

$$p(f_*|\mathbf{x}_*, \mathbf{x}, \mathbf{f}) = \mathcal{N}(k(\mathbf{x}_*, \mathbf{x})^{\mathrm{T}} k(\mathbf{x}, \mathbf{x})^{-1}\mathbf{f}, \\ k(\mathbf{x}_*, \mathbf{x}_*) - k(\mathbf{x}_*, \mathbf{x})^{\mathrm{T}} k(\mathbf{x}, \mathbf{x})^{-1} k(\mathbf{x}, \mathbf{x}_*)$$

Conditional Gaussians

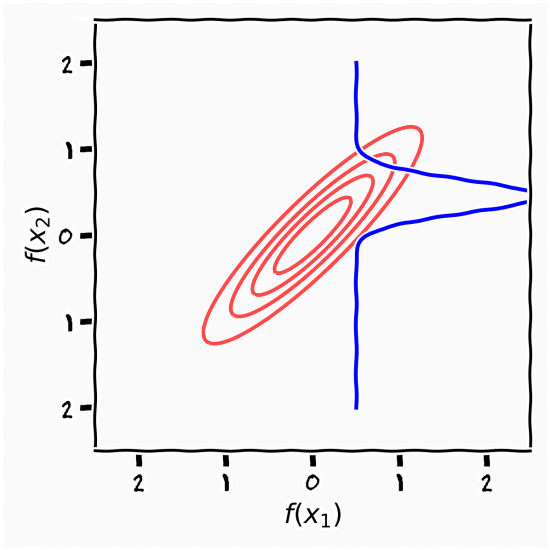

Gaussian Processes: "Predictive Posterior Samples"

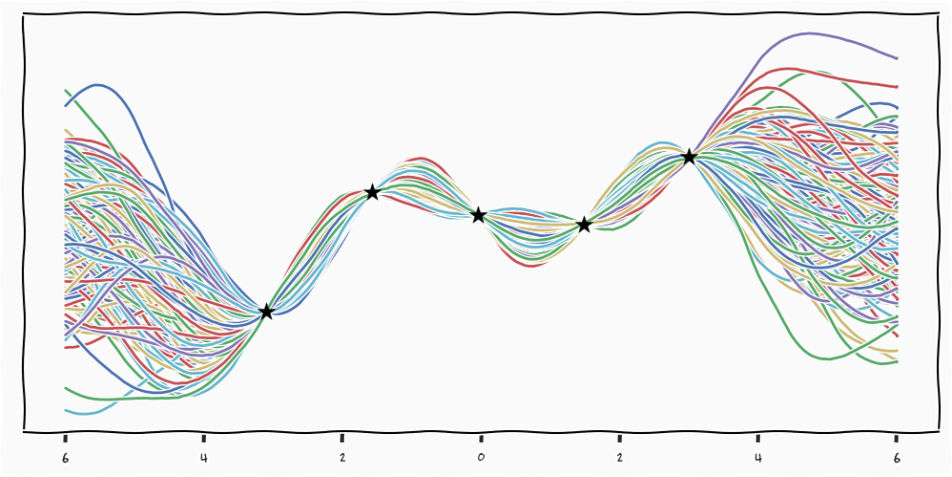

Gaussian Processes: "Predictive Posterior Process"

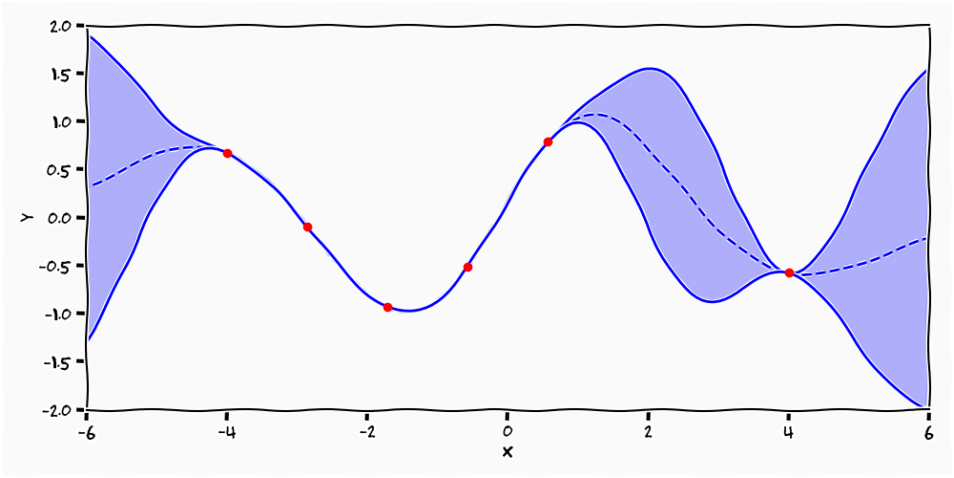

Posterior Processes

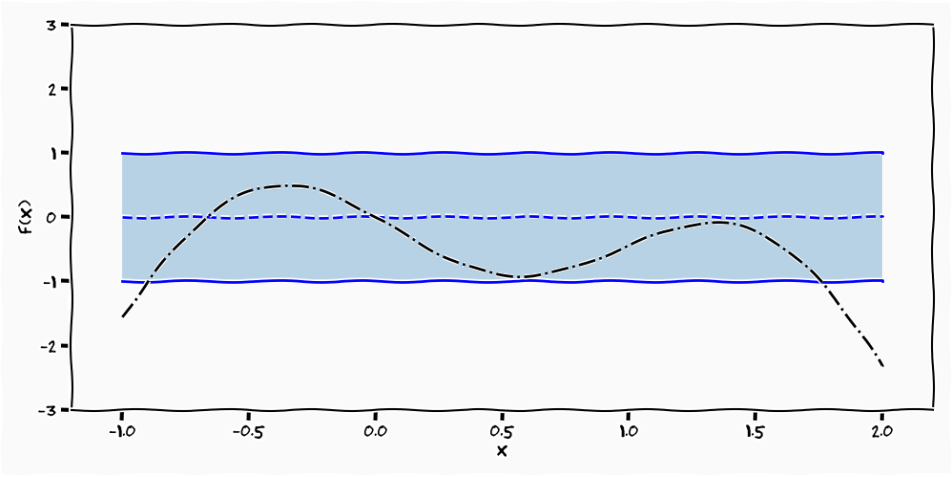

Posterior Processes

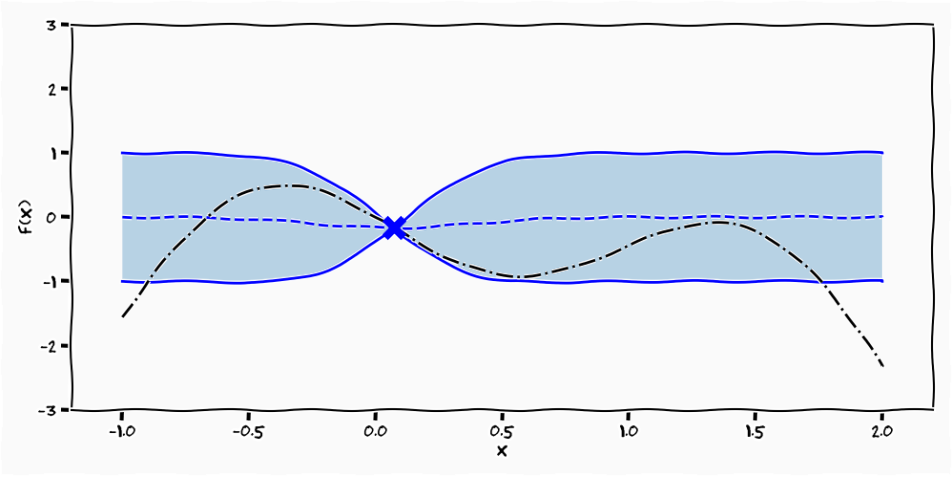

Posterior Processes

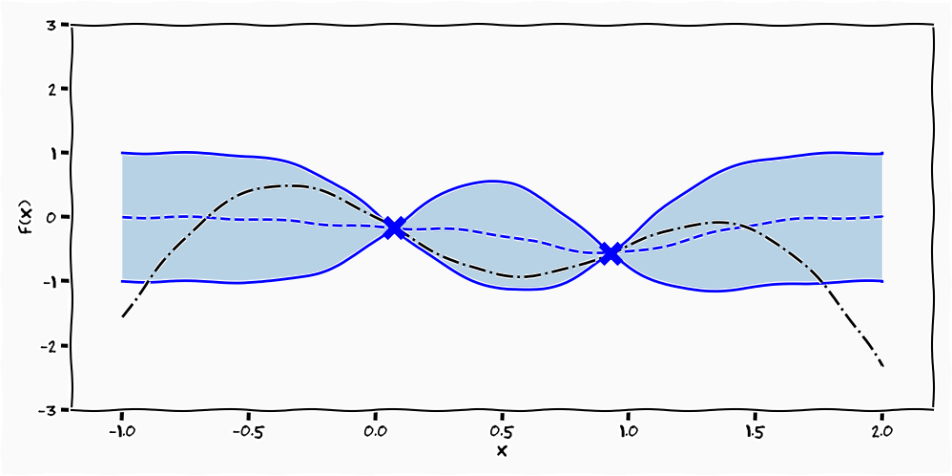

Posterior Processes

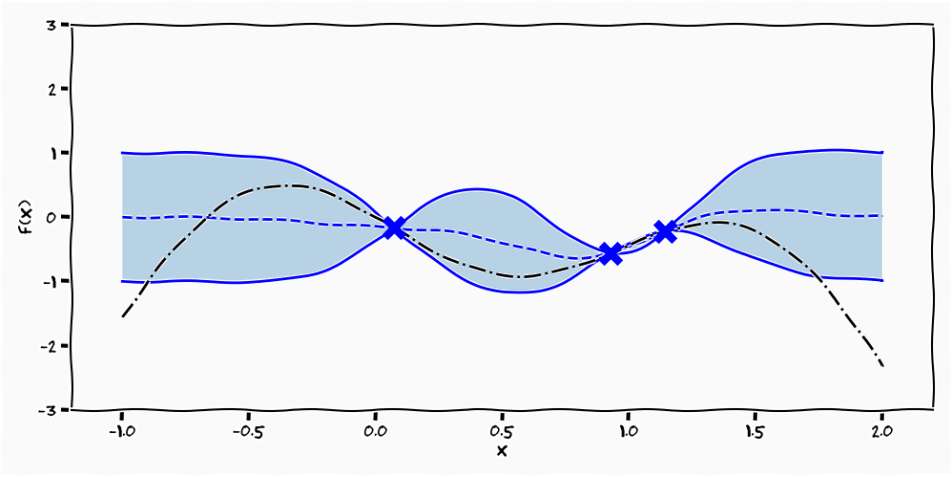

Posterior Processes

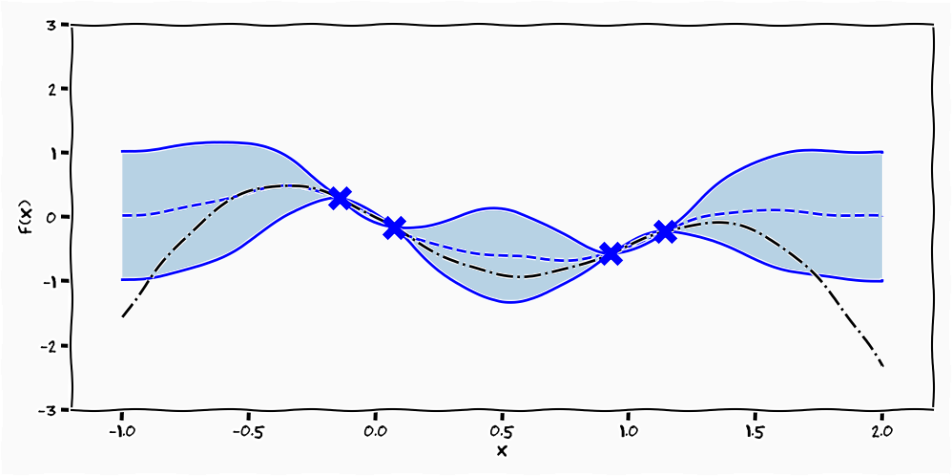

Posterior Processes

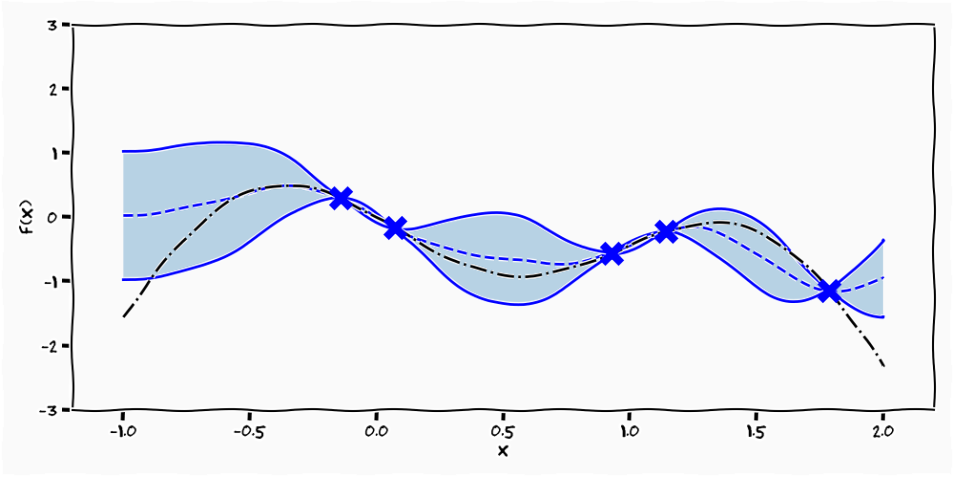

Posterior Processes

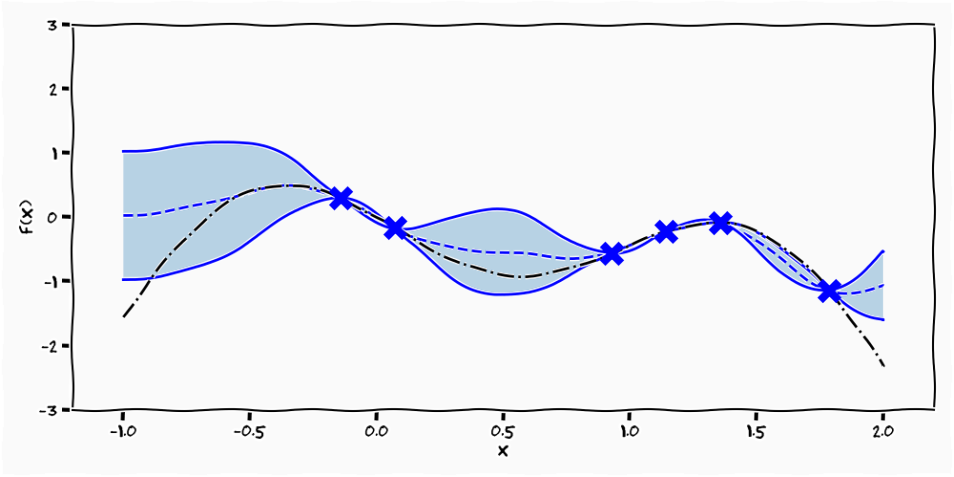

Posterior Processes

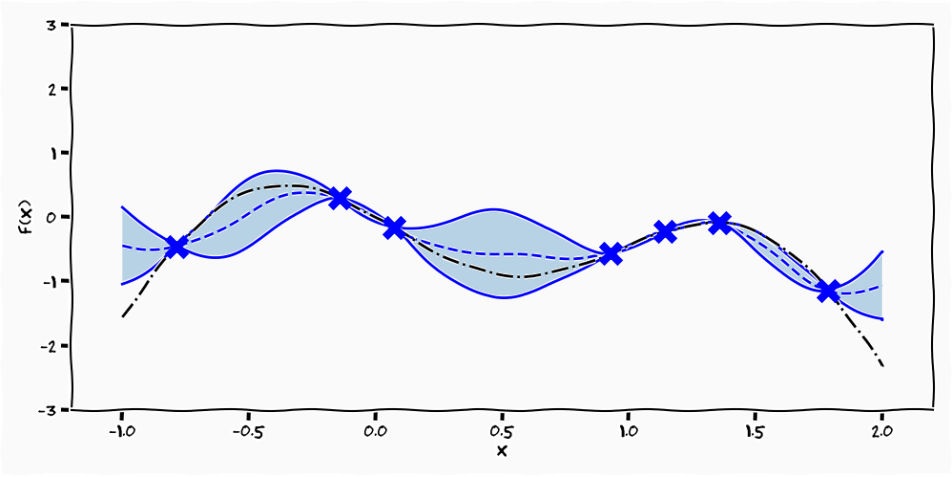

Posterior Processes

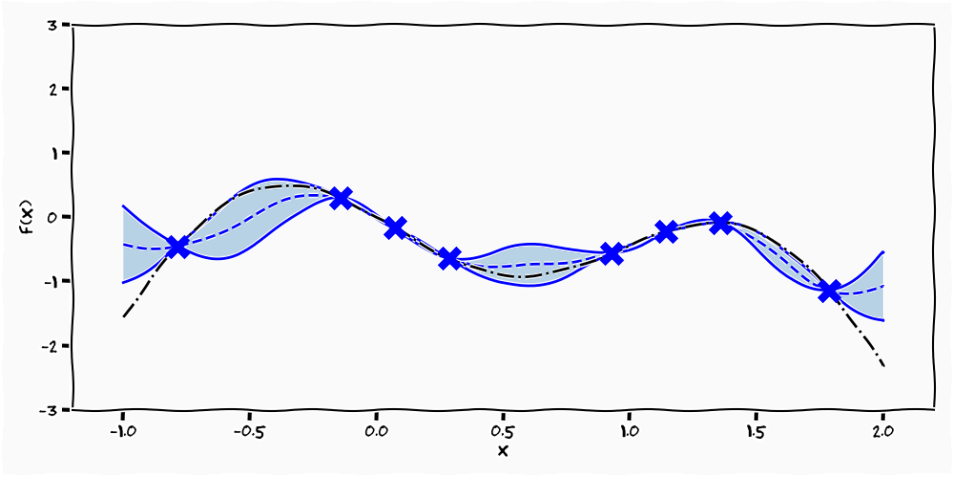

Posterior Processes

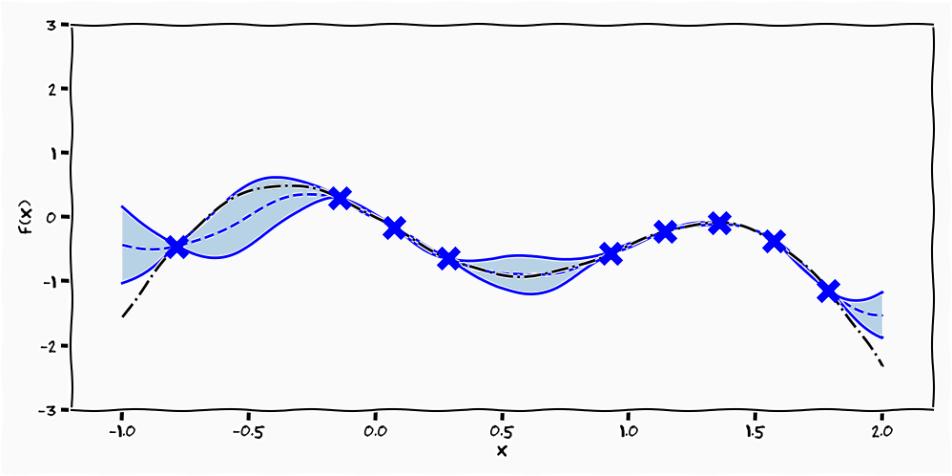

Posterior Processes

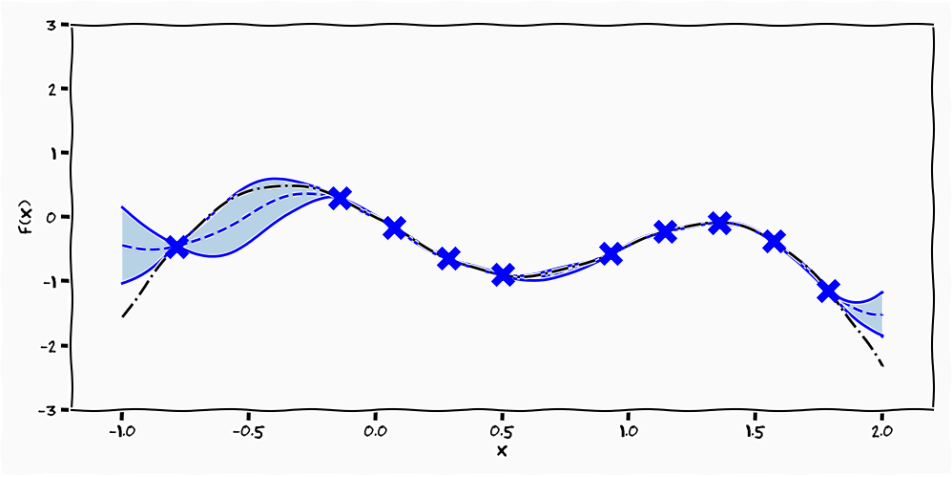

Posterior Processes

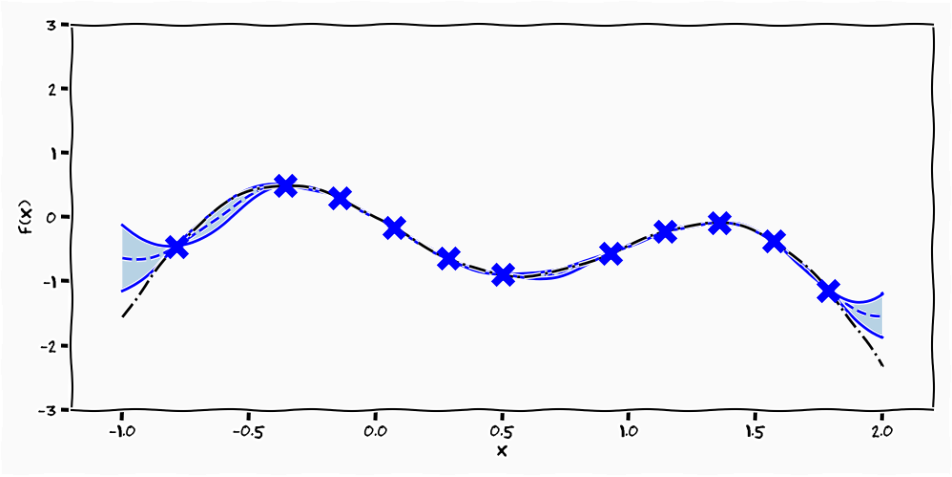

Posterior Processes

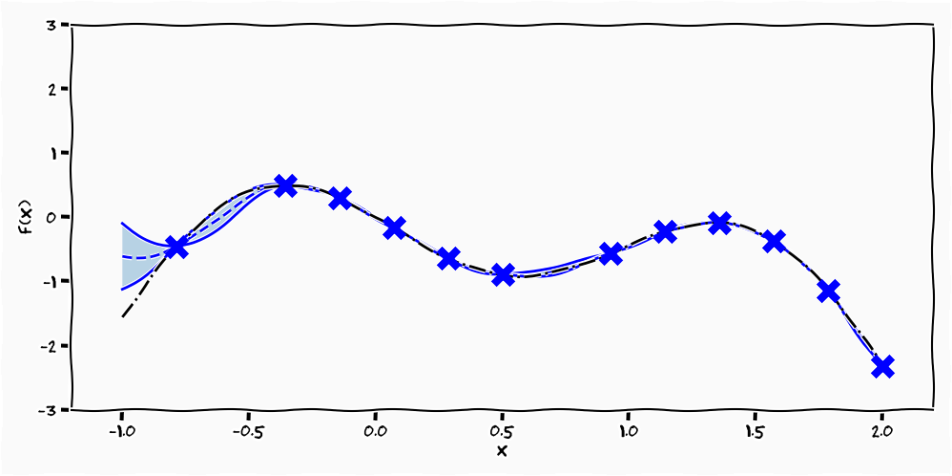

Posterior Processes

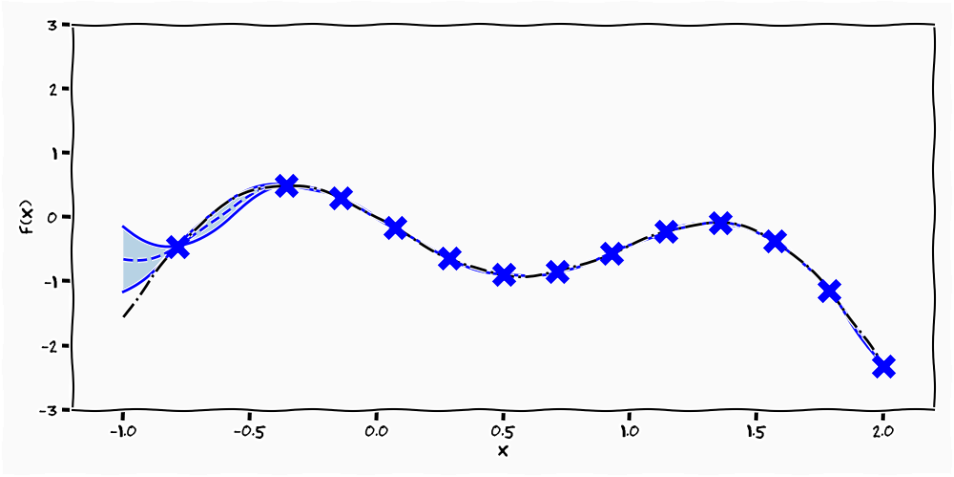

Posterior Processes

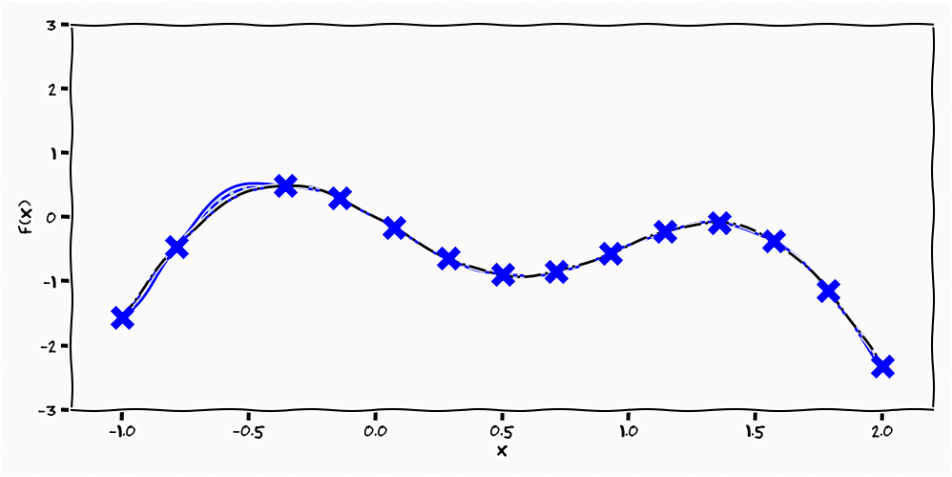

Posterior Processes

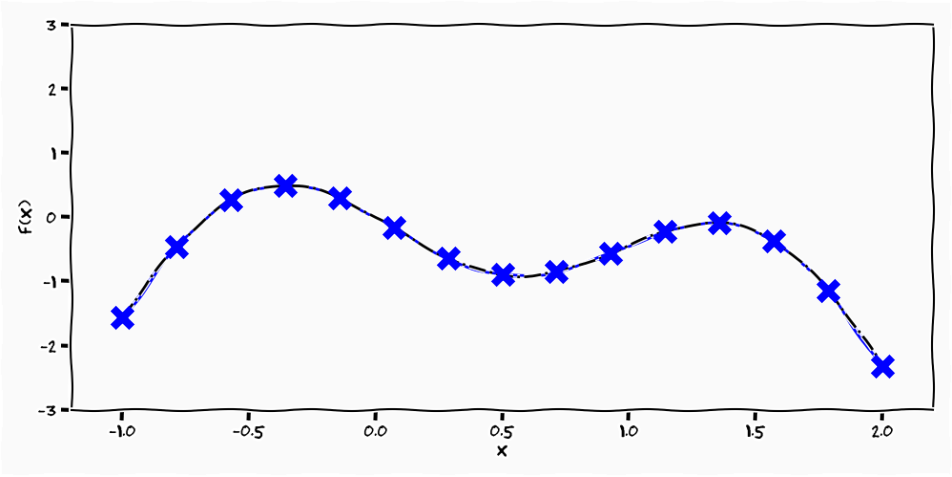

Model

- So far we have only looked at the prior
  - the same as when we sampled from $p(\mathbf{w})$ in the previous lecture
  - the "predictive posterior" has really been the "conditional prior"
- Now lets introduce a likelihood

$$p(\mathbf{y}, \mathbf{f}) = p(\mathbf{y} \mid \mathbf{f})p(\mathbf{f}) = p(\mathbf{f}) \prod_{i=1} p(y_i \mid f_i)$$

$$p(\mathbf{y} \mid \mathbf{f}) = \mathcal{N}(\mathbf{y} \mid \mathbf{f}, \sigma^2 \mathbf{I})$$

- Same motivation as with linear regression
- we want to explain away our ignorance from the data using σ

$$\begin{bmatrix} \mathbf{y} \\ \mathbf{f}_* \end{bmatrix} \sim \mathcal{N}\left(\begin{bmatrix} \mathbf{0} \\ \mathbf{0} \end{bmatrix}, \begin{bmatrix} k(\mathbf{x},\mathbf{x}) + \sigma^2\mathbf{I} & k(\mathbf{x},\mathbf{x}_*) \\ k(\mathbf{x}_*,\mathbf{x}) & k(\mathbf{x}_*,\mathbf{x}_*) \end{bmatrix}\right)$$

$$p(f_*|\mathbf{x}_*,\mathbf{x},\mathbf{y},\boldsymbol{\theta}) = \mathcal{N}(k(\mathbf{x}_*,\mathbf{x})^{\mathrm{T}}(K(\mathbf{x},\mathbf{x}) + \sigma^2\mathbf{I}))^{-1}\mathbf{y},$$
$$k(\mathbf{x}_*,\mathbf{x}_*) - k(\mathbf{x}_*,\mathbf{x})^{\mathrm{T}}(K(\mathbf{x},\mathbf{x}) + \sigma^2\mathbf{I})^{-1}K(\mathbf{x},\mathbf{x}_*))$$

Posterior Samples

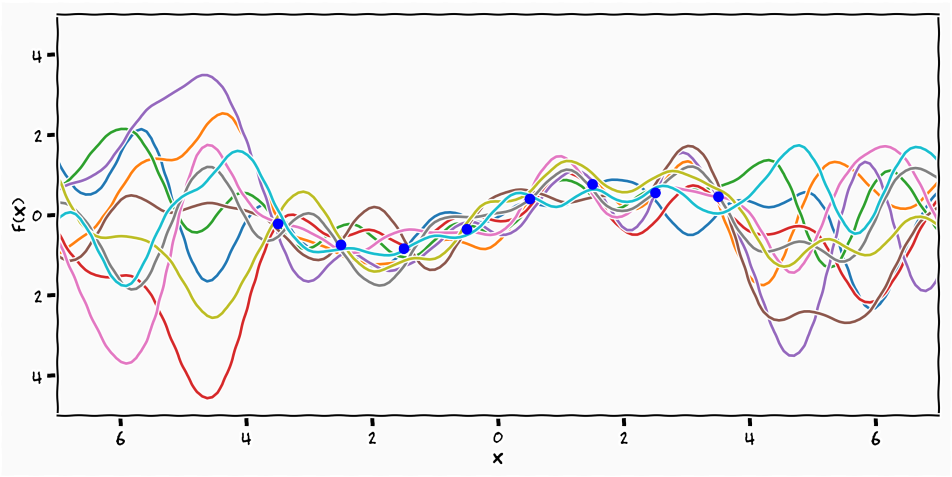

Posterior Samples

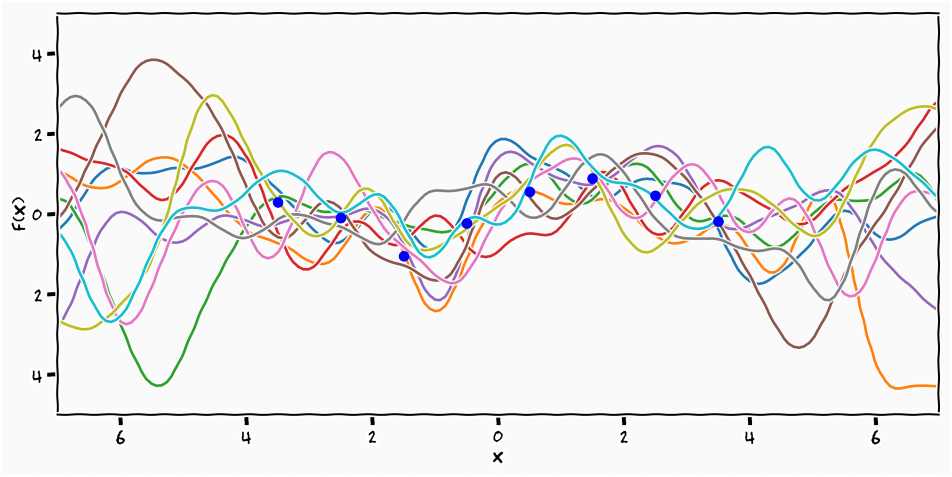

Gaussian Processes: Predictive Posterior Process

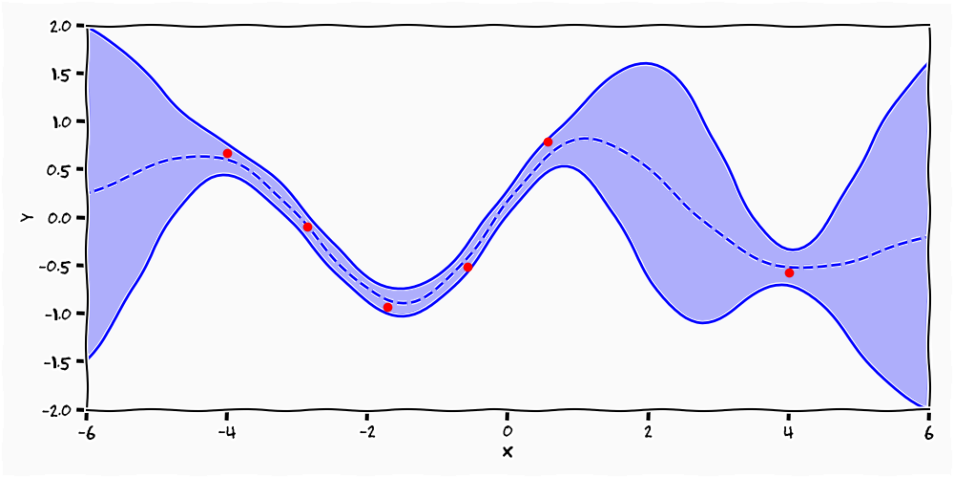

Gaussian Processes: Predictive Posterior Process

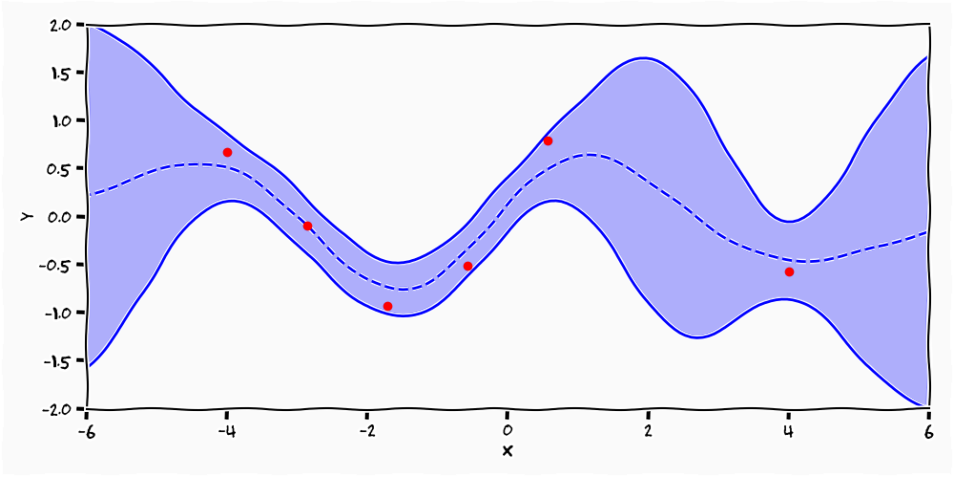

Gaussian Processes: Predictive Posterior Process

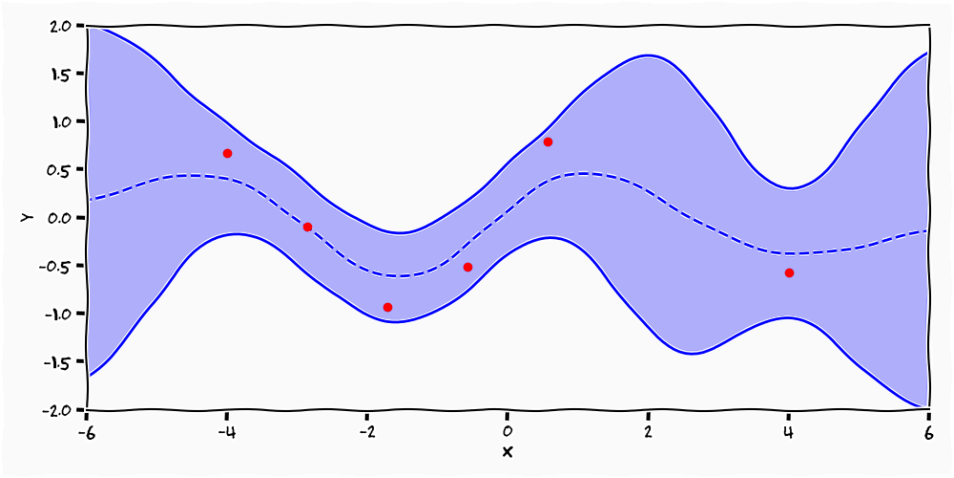

Gaussian Processes: Predictive Posterior Process

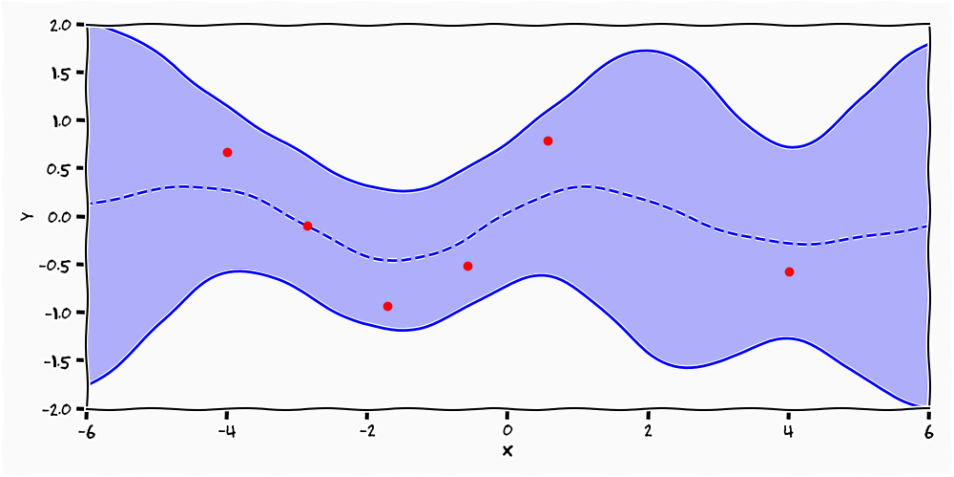

Gaussian Processes: Predictive Posterior Process

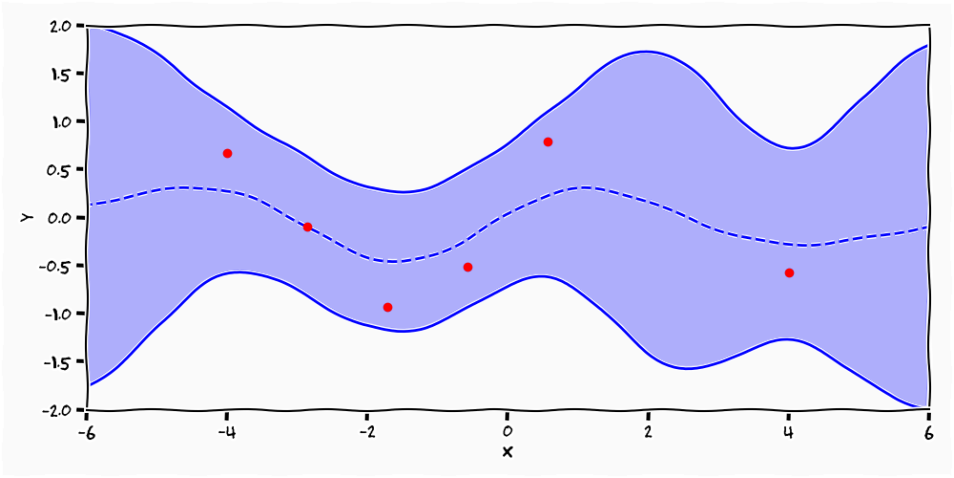

Two views

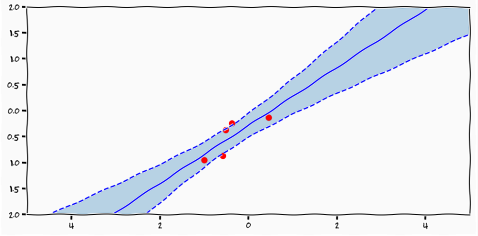
2.0
1.5
1.0
0.5
0.0
0.5
1.0
1.5
2.0
4
2
0
2
4

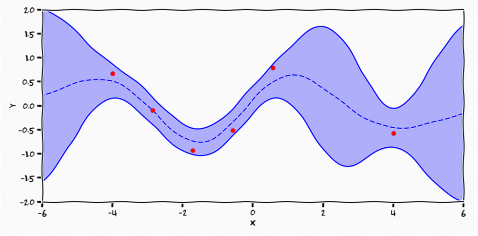
2.0
1.5
1.0
0.5
0.0
−0.5
−1.0
−1.5
−2.0
y
−6
−4
−2
0
2
4
6
x

- Gaussian Process - *Non-parametric formulation*

$$p(y_* \mid \mathbf{y}, \mathbf{x}_*, \mathbf{X}, \theta) = \mathcal{N}(y_* \mid \mu_*, K_*)$$
$$\mu_* = k(\mathbf{x}_*, \mathbf{x}) \left(\sigma^2 \mathbf{I} + k(\mathbf{x}, \mathbf{x})\right)^{-1} \mathbf{y}$$
$$K_* = k(\mathbf{x}_*, \mathbf{x}_*) - k(\mathbf{x}_*, \mathbf{x}) \left(k(\mathbf{x}, \mathbf{x}) + \sigma^2 \mathbf{I})\right)^{-1} k(\mathbf{x}, \mathbf{x}_*)$$

- Gaussian Process - *Non-parametric formulation*

$$p(y_* \mid \mathbf{y}, \mathbf{x}_*, \mathbf{X}, \theta) = \mathcal{N}(y_* \mid \mu_*, K_*)$$
$$\mu_* = k(\mathbf{x}_*, \mathbf{x}) \left(\sigma^2 \mathbf{I} + k(\mathbf{x}, \mathbf{x})\right)^{-1} \mathbf{y}$$
$$K_* = k(\mathbf{x}_*, \mathbf{x}_*) - k(\mathbf{x}_*, \mathbf{x}) \left(k(\mathbf{x}, \mathbf{x}) + \sigma^2 \mathbf{I})\right)^{-1} k(\mathbf{x}, \mathbf{x}_*)$$

- Linear Regression - *Parametric formulation*

$$p(y_* | \mathbf{y}, \mathbf{x}_*, \mathbf{X}, \alpha, \beta) = \int p(y_* | \mathbf{x}_*, \mathbf{w}, \beta) p(\mathbf{w} | \mathbf{y}, \mathbf{X}, \alpha, \beta) \mathrm{d}\mathbf{w}$$
$$= \mathcal{N}\left(y_* \mid \mu_*(x_*), \Sigma_*(x_*)\right)$$
$$\mu_* = (\beta(\alpha \mathbf{I} + \beta \mathbf{x}^{\mathrm{T}} \mathbf{x})^{-1} \mathbf{x} \mathbf{y}^{\mathrm{T}} \mathbf{x}_* = \mathbf{x}_*^{\mathrm{T}} \mathbf{x} \left(\beta(\alpha \mathbf{I} + \beta \mathbf{x}^{\mathrm{T}} \mathbf{x})\right)^{-1} \mathbf{y}$$
$$\Sigma_* = \frac{1}{\beta} + \mathbf{x}_*^{\mathrm{T}} (\alpha \mathbf{I} + \beta \mathbf{x}^{\mathrm{T}} \mathbf{x})^{-1} \mathbf{x}$$

$$\mu_*^{\mathsf{GP}} = k(\mathbf{x}_*, \mathbf{x}) \left(\sigma^2 \mathbf{I} + k(\mathbf{x}, \mathbf{x})\right)^{-1} \mathbf{y} \qquad \mu_*^{\mathsf{Lin}} = \mathbf{x}_*^{\mathrm{T}} \mathbf{x} \left(\beta(\alpha \mathbf{I} + \beta \mathbf{x}^{\mathrm{T}} \mathbf{x})\right)^{-1} \mathbf{y}$$

- both means are linear with respect to the output data

[1] or use a kernel $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^{\mathrm{T}} \mathbf{x}$

$$\mu_*^{\mathsf{GP}} = k(\mathbf{x}_*, \mathbf{x}) \left(\sigma^2 \mathbf{I} + k(\mathbf{x}, \mathbf{x})\right)^{-1} \mathbf{y} \qquad \mu_*^{\mathsf{Lin}} = \mathbf{x}_*^{\mathrm{T}} \mathbf{x} \left(\beta(\alpha \mathbf{I} + \beta \mathbf{x}^{\mathrm{T}} \mathbf{x})\right)^{-1} \mathbf{y}$$

- both means are linear with respect to the output data
- how about if we take a basis function approach such that $\mathbf{x}_*^{\mathrm{T}} \mathbf{x} = \Phi(\mathbf{x}_*)^{\mathrm{T}} \Phi(\mathbf{x})$

[1] or use a kernel $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^{\mathrm{T}} \mathbf{x}$

$$\mu_*^{\text{GP}} = k(\mathbf{x}_*, \mathbf{x}) \left(\sigma^2 \mathbf{I} + k(\mathbf{x}, \mathbf{x})\right)^{-1} \mathbf{y} \qquad \mu_*^{\text{Lin}} = \mathbf{x}_*^{\mathrm{T}} \mathbf{x} \left(\beta(\alpha \mathbf{I} + \beta \mathbf{x}^{\mathrm{T}} \mathbf{x})\right)^{-1} \mathbf{y}$$

- both means are linear with respect to the output data
- how about if we take a basis function approach such that $\mathbf{x}_*^{\mathrm{T}} \mathbf{x} = \Phi(\mathbf{x}_*)^{\mathrm{T}} \Phi(\mathbf{x})$
- if we think of $\mathbf{x}$ as the center of each basis function[1]

[1] or use a kernel $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^{\mathrm{T}} \mathbf{x}$

$$\mu_*^{\text{GP}} = k(\mathbf{x}_*, \mathbf{x}) \left(\sigma^2 \mathbf{I} + k(\mathbf{x}, \mathbf{x})\right)^{-1} \mathbf{y} \qquad \mu_*^{\text{Lin}} = \mathbf{x}_*^{\mathrm{T}} \mathbf{x} \left(\beta(\alpha \mathbf{I} + \beta \mathbf{x}^{\mathrm{T}} \mathbf{x})\right)^{-1} \mathbf{y}$$

- both means are linear with respect to the output data
- how about if we take a basis function approach such that $\mathbf{x}_*^{\mathrm{T}} \mathbf{x} = \Phi(\mathbf{x}_*)^{\mathrm{T}} \Phi(\mathbf{x})$
- if we think of $\mathbf{x}$ as the center of each basis function[1]
  - a basis function per data point

[1] or use a kernel $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^{\mathrm{T}} \mathbf{x}$

$$\mu_*^{\text{GP}} = k(\mathbf{x}_*, \mathbf{x}) \left(\sigma^2 \mathbf{I} + k(\mathbf{x}, \mathbf{x})\right)^{-1} \mathbf{y} \qquad \mu_*^{\text{Lin}} = \mathbf{x}_*^{\mathrm{T}} \mathbf{x} \left(\beta(\alpha \mathbf{I} + \beta \mathbf{x}^{\mathrm{T}} \mathbf{x})\right)^{-1} \mathbf{y}$$

- both means are linear with respect to the output data
- how about if we take a basis function approach such that $\mathbf{x}_*^{\mathrm{T}} \mathbf{x} = \Phi(\mathbf{x}_*)^{\mathrm{T}} \Phi(\mathbf{x})$
- if we think of $\mathbf{x}$ as the center of each basis function[1]
  - a basis function per data point
- if we could parametrise $\Phi(\mathbf{x}_*)^{\mathrm{T}} \Phi(\mathbf{x}) = k(\mathbf{x}_*, \mathbf{x})$ as a function

[1] or use a kernel $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^{\mathrm{T}} \mathbf{x}$

$$\mu_*^{\text{GP}} = k(\mathbf{x}_*, \mathbf{x}) \left(\sigma^2 \mathbf{I} + k(\mathbf{x}, \mathbf{x})\right)^{-1} \mathbf{y} \qquad \mu_*^{\text{Lin}} = \mathbf{x}_*^{\mathrm{T}} \mathbf{x} \left(\beta(\alpha \mathbf{I} + \beta \mathbf{x}^{\mathrm{T}} \mathbf{x})\right)^{-1} \mathbf{y}$$

- both means are linear with respect to the output data
- how about if we take a basis function approach such that $\mathbf{x}_*^{\mathrm{T}} \mathbf{x} = \Phi(\mathbf{x}_*)^{\mathrm{T}} \Phi(\mathbf{x})$
- if we think of $\mathbf{x}$ as the center of each basis function[1]
  - a basis function per data point
- if we could parametrise $\Phi(\mathbf{x}_*)^{\mathrm{T}} \Phi(\mathbf{x}) = k(\mathbf{x}_*, \mathbf{x})$ as a function
- this leads to the interpretation of GPs as infinite basis functions

[1] or use a kernel $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^{\mathrm{T}} \mathbf{x}$

Summary

Summary

- Non-parametrics
  - parametrise relationship between data

Summary

- Non-parametrics
  - parametrise relationship between data
- Gaussian processes
  - Implementation: "just a big big Gaussian"
  - Theory: projection of an infinite stochastic process

For all permutations π, measurable sets $F_i \subseteq \mathbb{R}^n$ and probability measure ν

1. Exchangeable

$$\nu_{t_{\pi(1)} \cdots t_{\pi(k)}} \left(F_{\pi(1)} \times \cdot \times F_{\pi(k)} \right) = \nu_{t_1 \cdots t_k} \left(F_1 \times \cdots \times F_k \right)$$

2. Marginal

$$\nu_{t_1 \cdot t_k} \left(F_1 \times \cdot \times F_k \right) = \nu_{t_1 \cdots t_k, t_{k+1} \cdot t_{k+m}} \left(F_1 \times \cdot \times F_k \times \mathbb{R}^n \times \cdot \times \mathbb{R}^n \right)$$

In this case the finite dimensional probability measure is a realisation of an underlying stochastic process

eof