Machine Learning and the Physical World

Lecture 7 : Probabilistic Numerics

Carl Henrik Ek - che29@cam.ac.uk

31st of October, 2023

http://carlhenrik.com


Data + Model
Compute
Prediction

The role of Uncertainty/Ignorance

$$p(y) = \int p(y \mid f) p(f) \mathrm{d}f$$

What do we do with uncertainty?

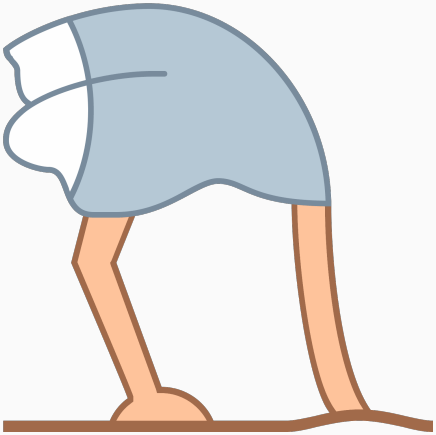

Bayesian Optimisation

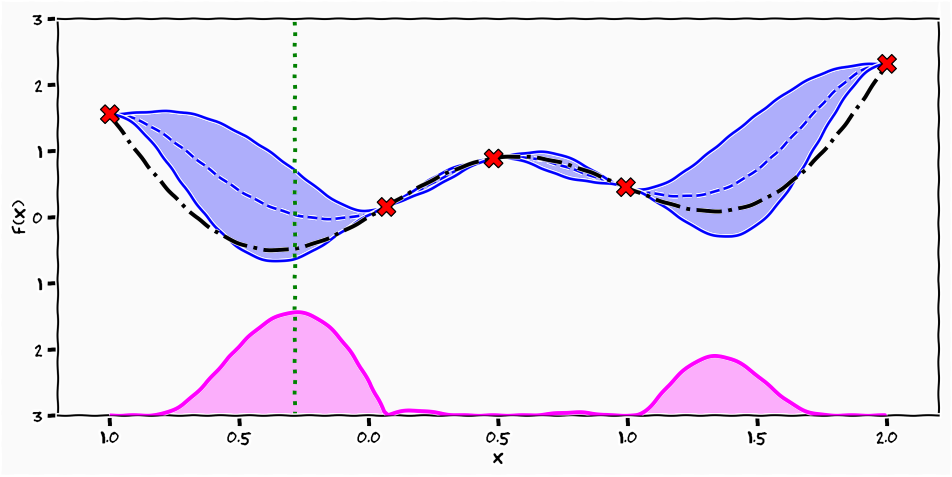

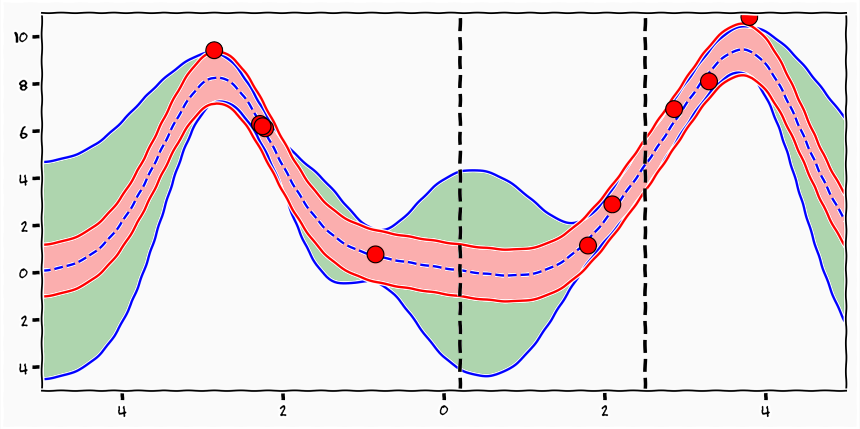

$$y_i = f_i + \epsilon$$

Aleatoric/Stochastic "Randomness" inherent in system, or noise in our measurement of system

Uncertainty Quantification

Aleatoric/Stochastic "Randomness" inherent in system, or noise in our measurement of system

Epistemic Uncertainty related to our ignorance of a the underlying system

Uncertainty Quantification

Aleatoric/Stochastic "Randomness" inherent in system, or noise in our measurement of system

Epistemic Uncertainty related to our ignorance of a the underlying system

Computational *Uncertainty related to finite computation, or intractable computations*

"The need for probability only arises out of uncertainty: It has no place if we are certain that we know all aspects of a problem. But our lack of knowledge also must not be complete, otherwise we would have nothing to evaluate. There is thus a spectrum of degrees of uncertainty. While the probability for the sixth decimal digit of a number in a table of logarithms to equal 6 is 1/10 a priori, in reality, all aspects of the corresponding problem are well determined, and, if we wanted to make the effort, we could find out its exact value. The same holds for interpolation, for the integration methods of Cotes or Gauss, etc"

– Henri Poincare, 1896

Computational Decisions

- Computation is expensive, how much knowledge will I gain from computing more?

Computational Decisions

- Computation is expensive, how much knowledge will I gain from computing more?
- What should I compute in order to reduce my uncertainty as much as possible?

Computational Decisions

- Computation is expensive, how much knowledge will I gain from computing more?
- What should I compute in order to reduce my uncertainty as much as possible?
- How much should I trust the computation I have done?

Computational Decisions

- Computation is expensive, how much knowledge will I gain from computing more?
- What should I compute in order to reduce my uncertainty as much as possible?
- How much should I trust the computation I have done?
- How precise should I do down-stream tasks based on the information from a specific computation?

Why Probabilistic Numerics?

"[round-off errors] are strictly very complicated but uniquely defined number theoretical functions [of the inputs], yet our ignorance of their true nature is such that we best treat them as random variables."

– Neumann et al., 1947

I believe in...

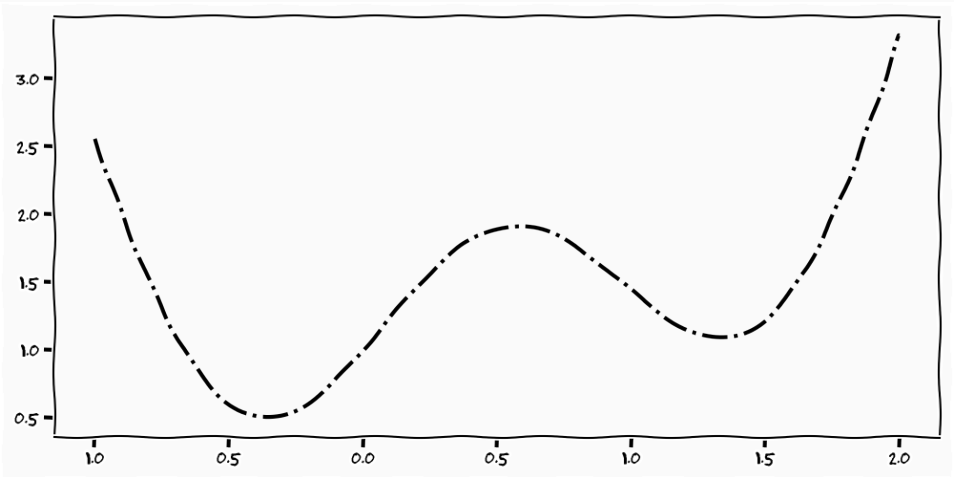
3.0
2.5
2.0
1.5
1.0
0.5
1.0
0.5
0.0
0.5
1.0
1.5
2.0

Formalisation [Cockayne et al., 2017]

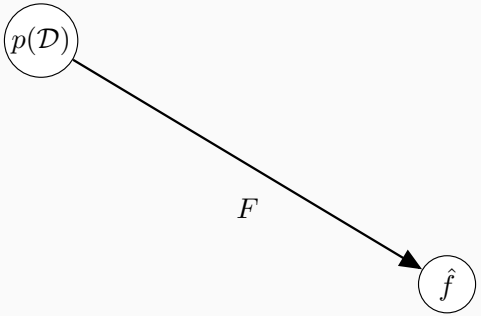

$$F : p(\mathcal{D}) \to p(\mathcal{Y}|\mathcal{X})$$

Formalisation [Cockayne et al., 2017]

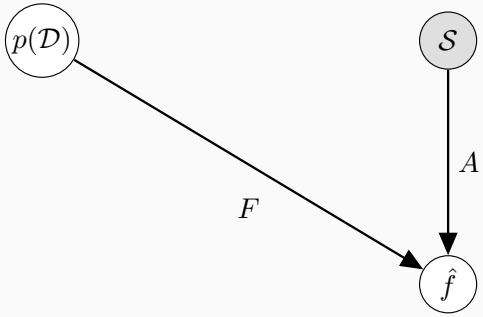

$$A \circ \mathcal{S} \approx F \circ p(\mathcal{D})$$

Formalisation

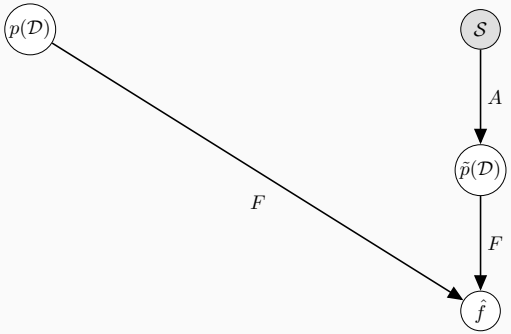

$$A \circ S \circ p(\mathcal{D}) \approx p(\mathcal{D})$$

Formalisation [Cockayne et al., 2017]

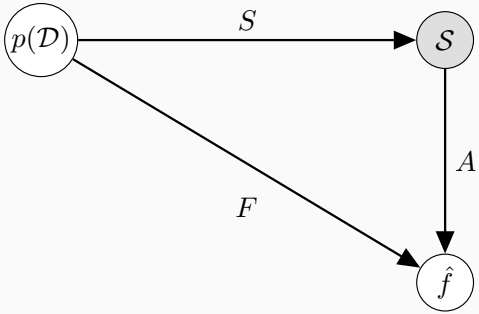

$$A \circ S \circ p(\mathcal{D}) \approx F \circ p(\mathcal{D})$$

Bayesian Optimisation

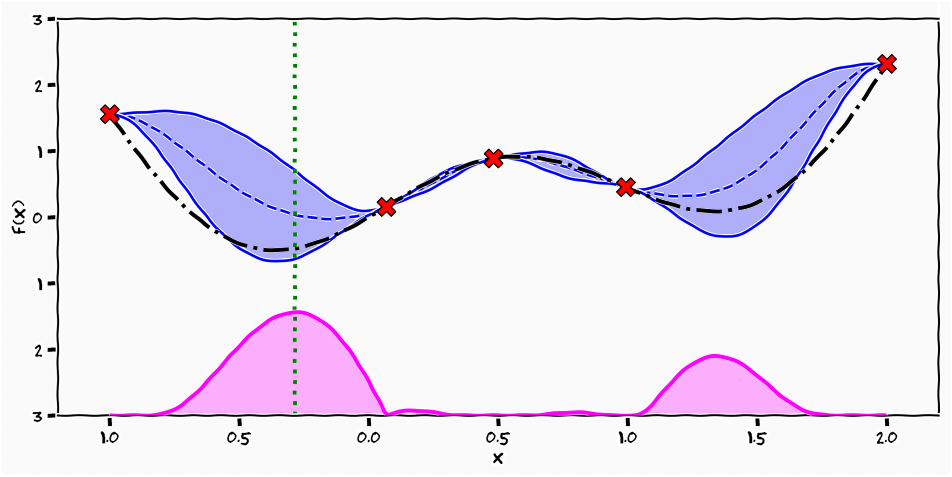

Bayesian Optimisation

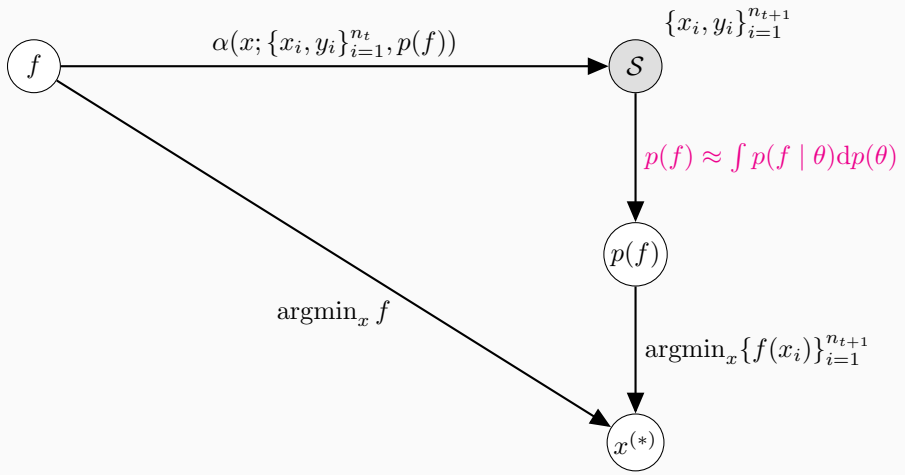

Formalisation [Cockayne et al., 2017]

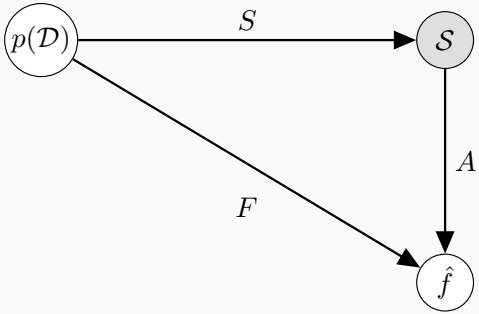

$$A \circ S \circ p(\mathcal{D}) \approx F \circ p(\mathcal{D})$$

Numerical Computations[1]

Linear Algebra given $As = y$ estimate x s.t. $Ax = b$

Optimisation given $\nabla f(x_i)$ estimate x s.t. $\nabla f(x) = 0$

Analysis given $f(x, t)$ estimate $x(t)$ s.t. $\mathrm{d}x = f(x, t)$

Quadrature given $f(x_i)$ estimate $\int_a^b f(x)\mathrm{d}x$

[1] https://www.cs.toronto.edu/~duvenaud/talks/odes_runge_kutta_nips.pdf

Quantity of Interest

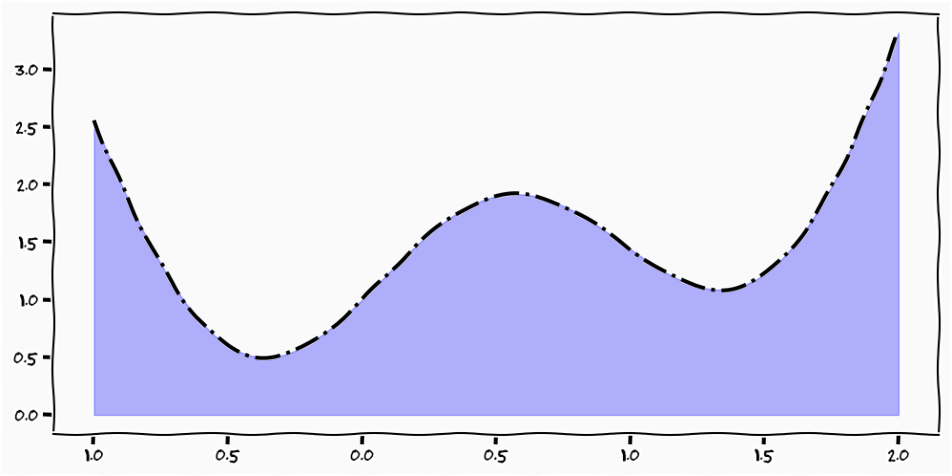

Integration is a significant numerical problem in many fields of science and engineering. It is a key step in inference, where it is encountered when averaging over the many states of the world consistent with observed data. Indeed, a provocative Bayesian view is that integration is the single challenge separating us from systems that fully automate statistics. More speculatively still, such systems may even exhibit artificial intelligence (ai).
– Hennig, Osbourne, Kersting

Integration

$$F := \int f(x) \mathrm{d}\nu(x)$$

- $\nu(x)$ is the measure that we are integrating over

$$\underbrace{p(\mathcal{D})}_{F} = \int \underbrace{p(\mathcal{D} \mid \theta)}_{f(\theta)} \underbrace{p(\theta)\mathrm{d}\theta}_{\mathrm{d}\nu(\theta)}$$

- marginalisation[2] is integration over the prior probability measure on the parameter

[2] think of computing the evidence

Quadrature [Cockayne et al., 2017]

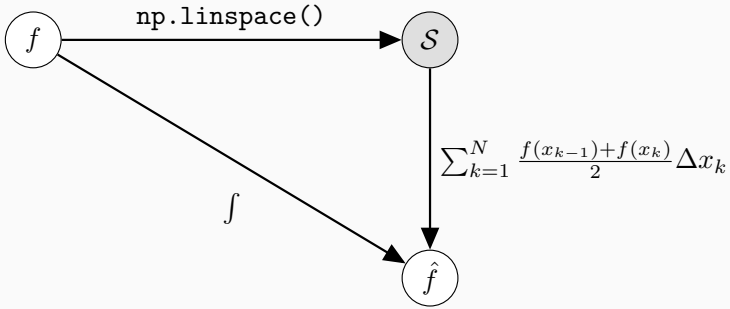

$$A \circ \mathcal{S} \approx \int f(x) \mathrm{d}x$$

Quadrature

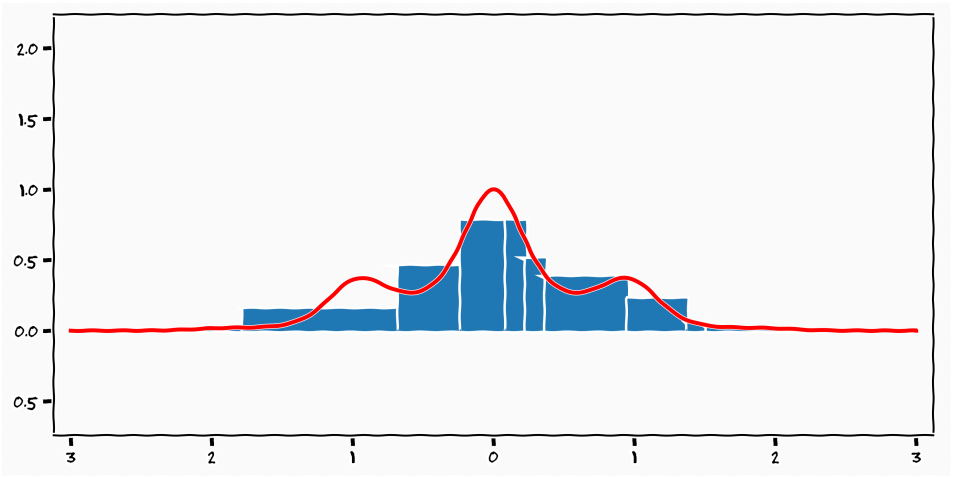

Quadrature

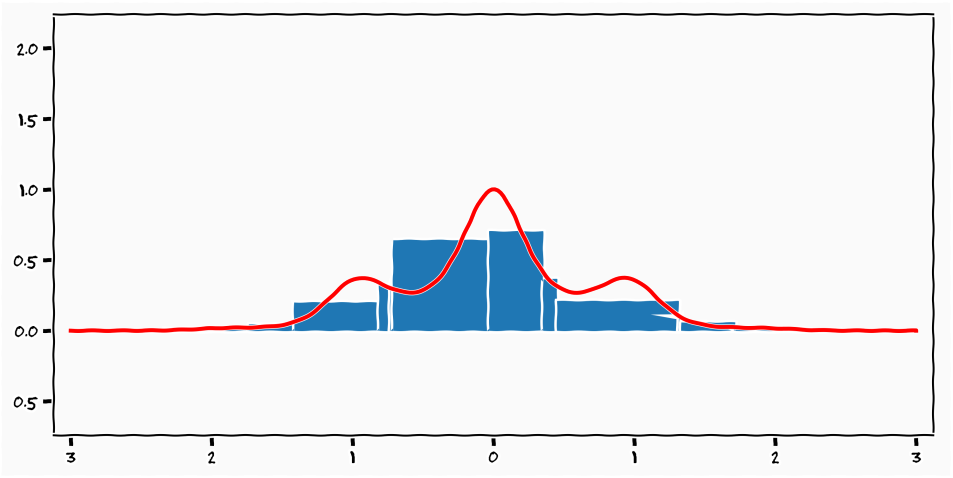

Numerical Computation

A numerical method estimates a function's latent property given the result of computations.

Numerical Computation

A numerical method estimates a function's latent property given the result of computations.

$\frac{\text{Numerical algorithms}}{\text{Statistical inference}}$ takes data in the form of $\frac{\text{evaluations of computations}}{\text{measurements of observed variables}}$ and aims to return predictions of the quantity of interest.

> *A numerical method estimates a function's latent property given the result of computations.*

$\frac{\text{Numerical algorithms}}{\text{Statistical inference}}$ takes data in the form of $\frac{\text{evaluations of computations}}{\text{measurements of observed variables}}$ and aims to return predictions of the quantity of interest.

Should we think about computation as inference?

$$p(\hat{f} \mid S)$$

Decision which algorithm to use when

Use of Computational Uncertainty

$$p(\hat{f} \mid S)$$

Decision which algorithm to use when

Decision efficient use of expensive algorithms

Use of Computational Uncertainty

$$p(\hat{f} \mid S)$$

Decision which algorithm to use when

Decision efficient use of expensive algorithms

Decision when to stop computation

Use of Computational Uncertainty

$$p(\hat{f} \mid S)$$

Decision which algorithm to use when

Decision efficient use of expensive algorithms

Decision when to stop computation

Decision effect on downstream tasks

Albert Valentionvic Suldin (1924-1996) worked on error minimising estimators for numerical algorithms, how to design algorithms from a statistical perspective

When computation was expensive

Albert Valentionvic Suldin (1924-1996) worked on error minimising estimators for numerical algorithms, how to design algorithms from a statistical perspective

Frederick Michael Larkin (1936-1982) incorporating the notion of prior knowledge into numerical algorithms to make robust calculations

Bayesian Quadrature

Quantity of interest

$$F := \int_{-3}^{3} \underbrace{e^{-(sin(3x))^2 - x^2}}_{f(x)} \mathrm{d}x$$

- $f(x)$ fully specified and deterministic
- F is deterministic
- F cannot be computed analytically

Integration

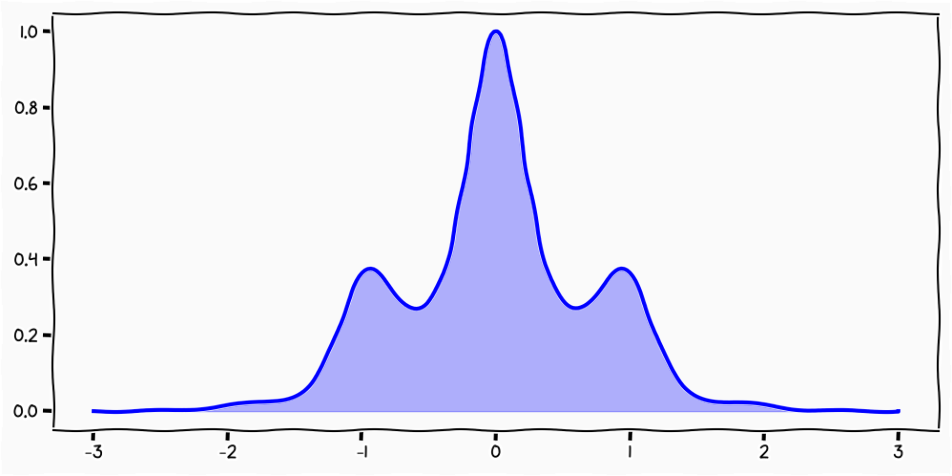

What we would like

$$p(F \mid Y)$$

- given that I have seen data Y what is my belief about the integral

What we would like

$$p(F \mid Y)$$

- given that I have seen data Y what is my belief about the integral
- allows for "active learning"

What we would like

$$p(F \mid Y)$$

- given that I have seen data Y what is my belief about the integral
- allows for "active learning"
- exploration/exploitation etc.

Emulation

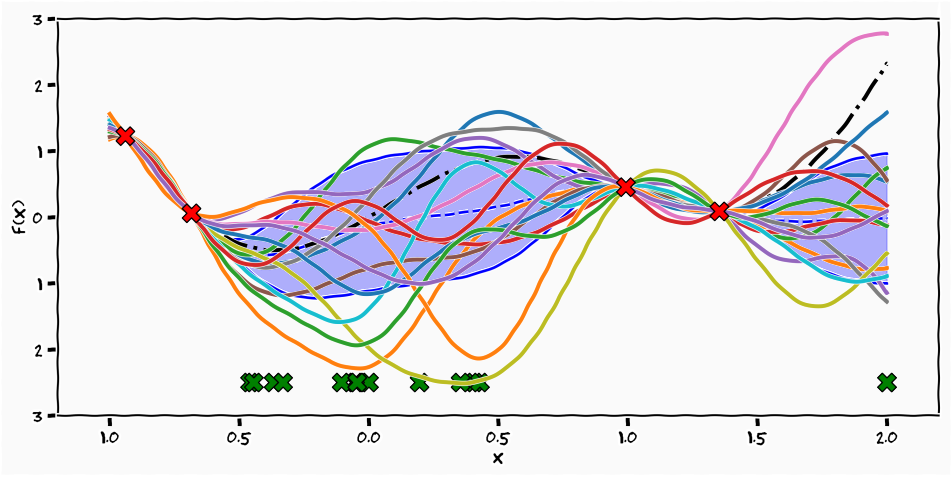

Quadrature

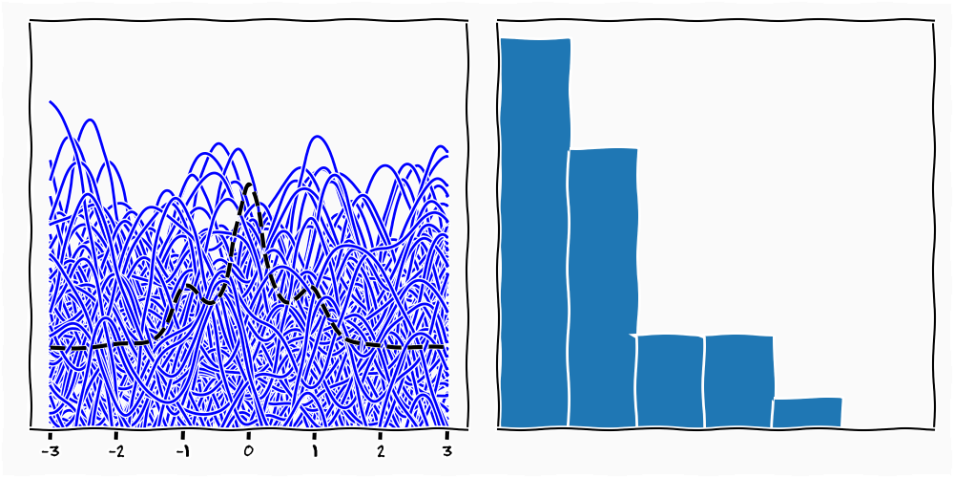

Quadrature

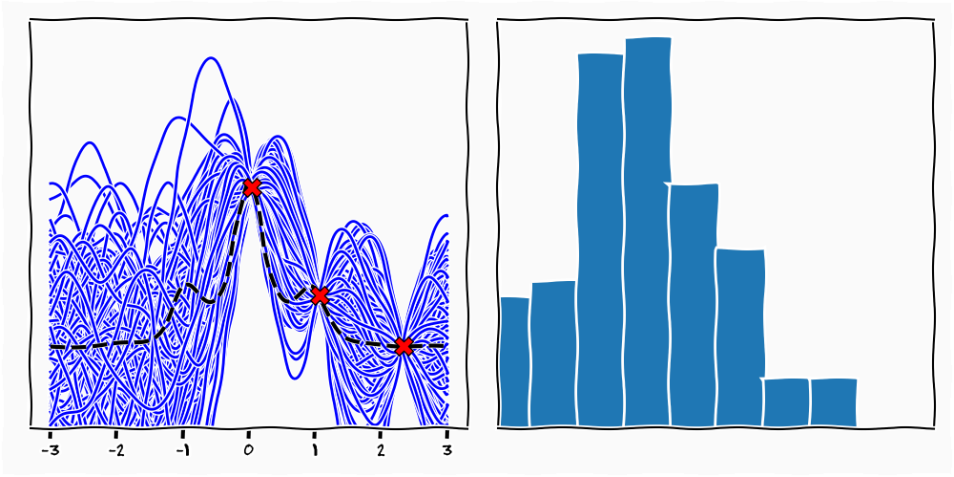

Quadrature

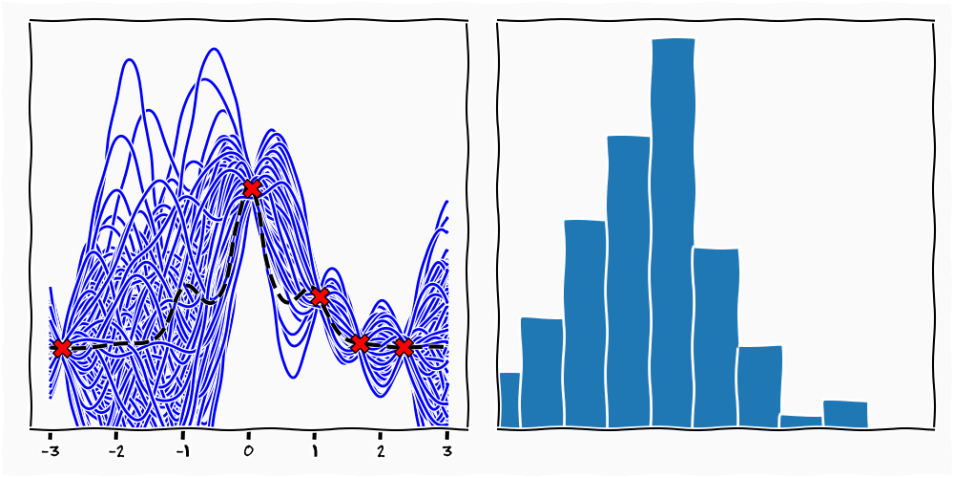

Quadrature

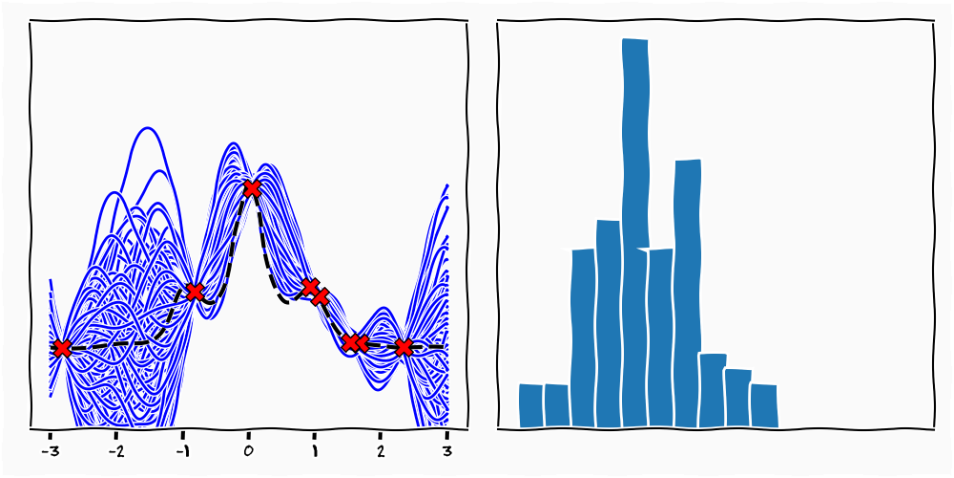

Quadrature

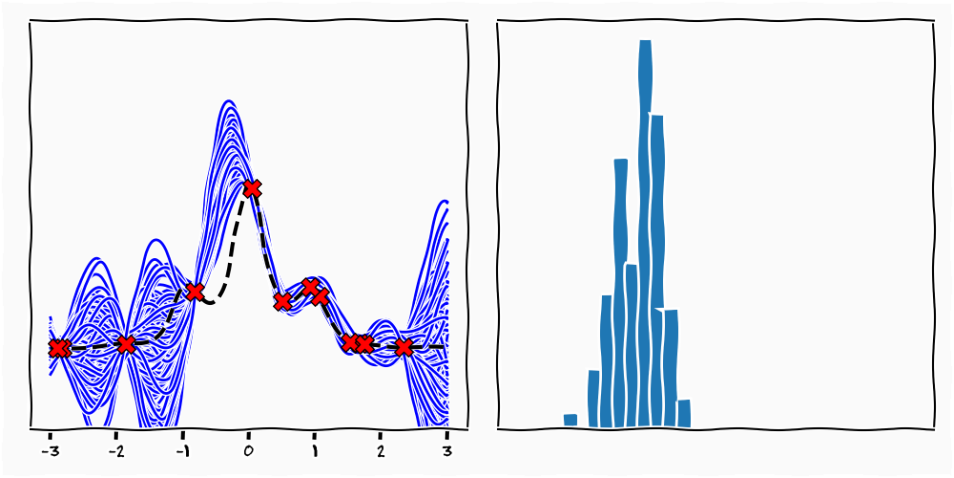

Quadrature

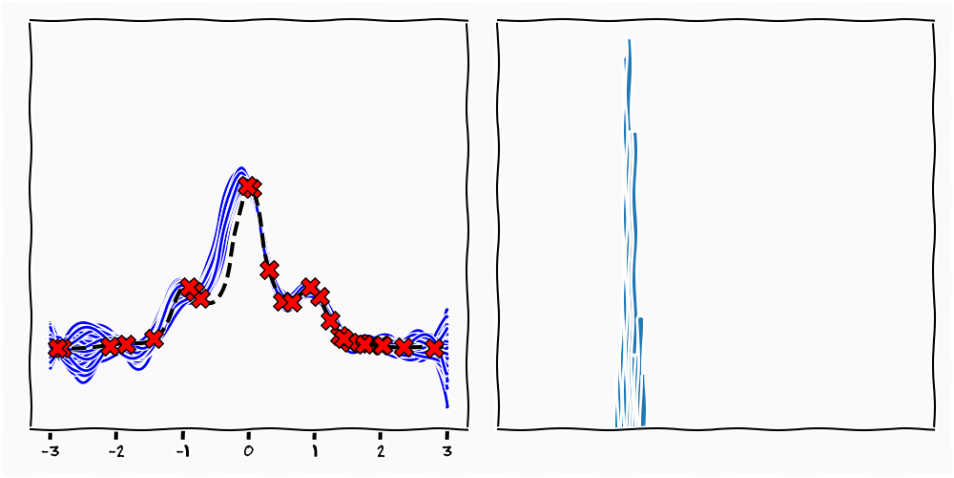

Quadrature

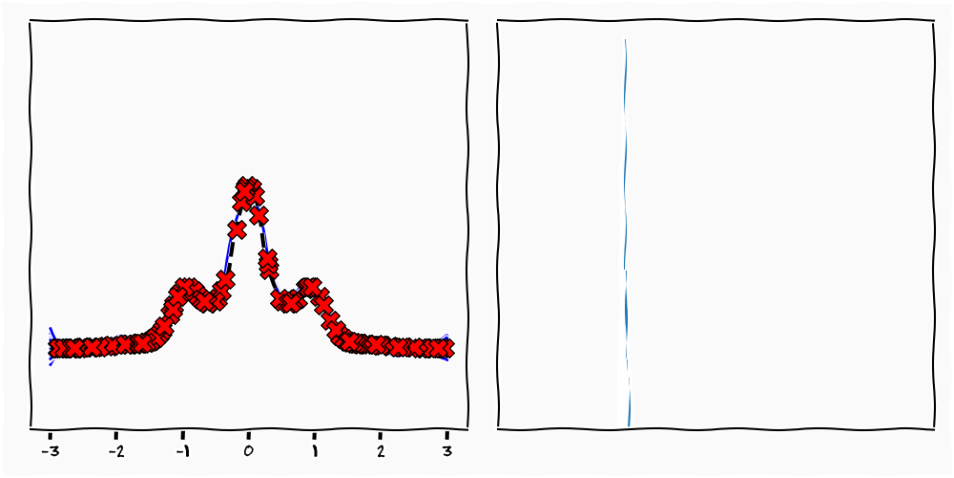

$$F := \int_{-3}^{3} \underbrace{e^{-(sin(3x))^2 - x^2}}_{f(x)} \mathrm{d}x$$

Knowledge

- $f(x)$ strictly positive $\Rightarrow F > 0$
- bounded above by,

$$f(x) \leq e^{-x^2}$$

- Therefore,

$$0 < F < \int_{-\infty}^{\infty} e^{-x^2} \mathrm{d}x = \sqrt{\pi}$$

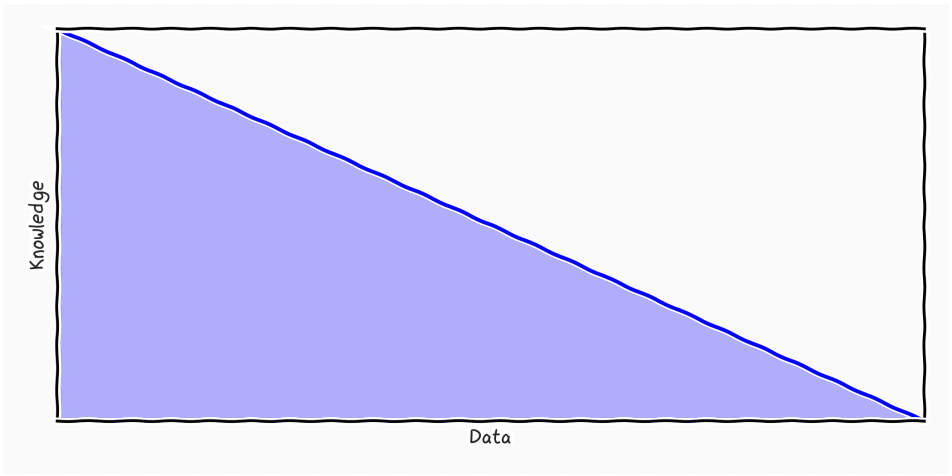
Knowledge
Data

$$F := \int f(x) \mathrm{d}\nu(x)$$

- $\nu(x)$ is the measure that we are integrating over

Bayesian Quadrature [O'Hagan, 1991]

$$p(F, Y) = \int p(F \mid f) p(Y \mid f) p(f) \mathrm{d}f$$

Bayesian Quadrature [O'Hagan, 1991]

$$
\begin{aligned}
p(F, Y) &= \int p(F \mid f) p(Y \mid f) p(f) \mathrm{d}f \\
&= \int \delta \left(F - \int_{\mathcal{X}} f \mathrm{d}x \right) \prod_{i}^{N} \delta(y_i - f(x_i)) p(f) \mathrm{d}f
\end{aligned}
$$

Bayesian Quadrature [O'Hagan, 1991]

$$p\begin{pmatrix} Y \\ F \end{pmatrix} = \mathcal{N}\left(\begin{bmatrix} \mathbf{m}_X \\ \int m_X(x)\mathrm{d}x \end{bmatrix}, \begin{bmatrix} k(X,X) & \int k(X,x)\mathrm{d}x \\ \int k(x,X)\mathrm{d}x & \int\int k(x,x')\mathrm{d}x\mathrm{d}x' \end{bmatrix}\right)$$

- We can derive $p(F \mid Y)$ through our normal conditioning procedure

Bayesian Quadrature [O'Hagan, 1991]

$$p\begin{pmatrix} Y \\ F \end{pmatrix} = \mathcal{N}\left(\begin{bmatrix} \mathbf{m}_X \\ \int m_X(x)\mathrm{d}x \end{bmatrix}, \begin{bmatrix} k(X,X) & \int k(X,x)\mathrm{d}x \\ \int k(x,X)\mathrm{d}x & \int\int k(x,x')\mathrm{d}x\mathrm{d}x' \end{bmatrix}\right)$$

- We can derive $p(F \mid Y)$ through our normal conditioning procedure
- $p(F \mid Y) = \mathcal{N}(\mu_F, k_F)$ is a uni-variate Gaussian

Statistical Inference

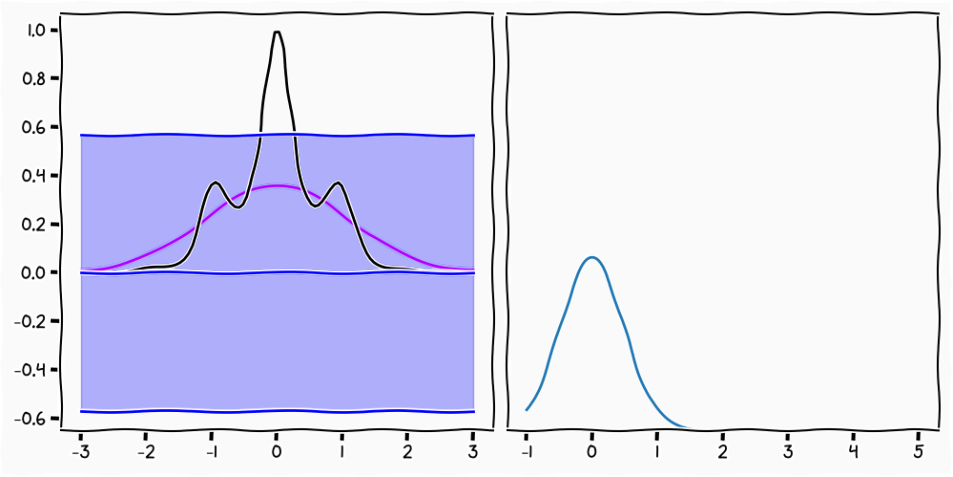

Statistical Inference

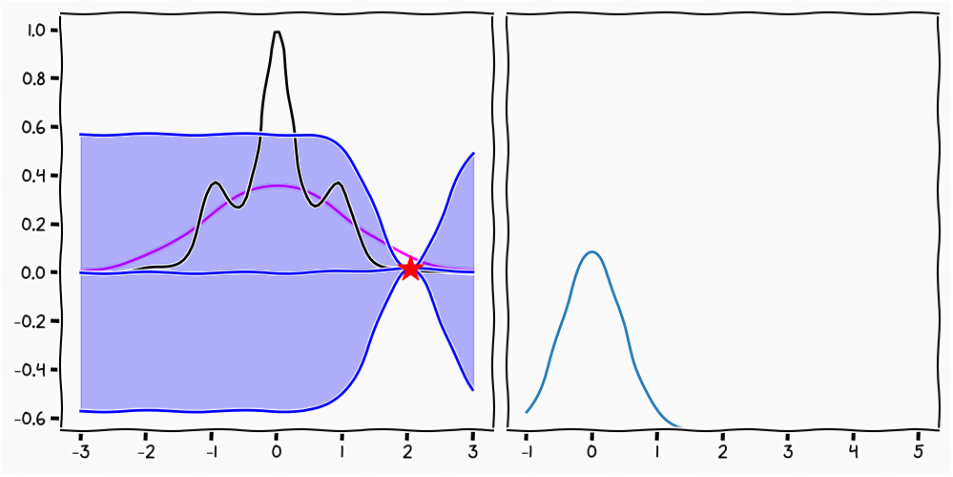

Statistical Inference

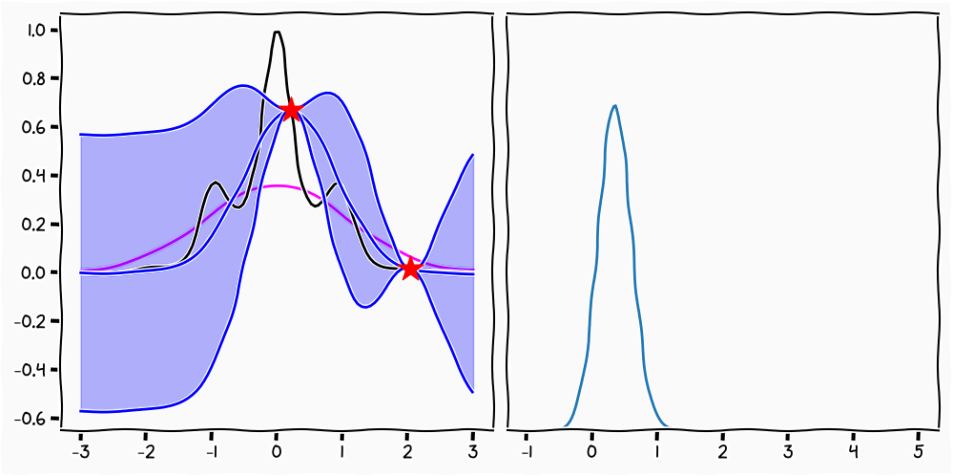

Statistical Inference

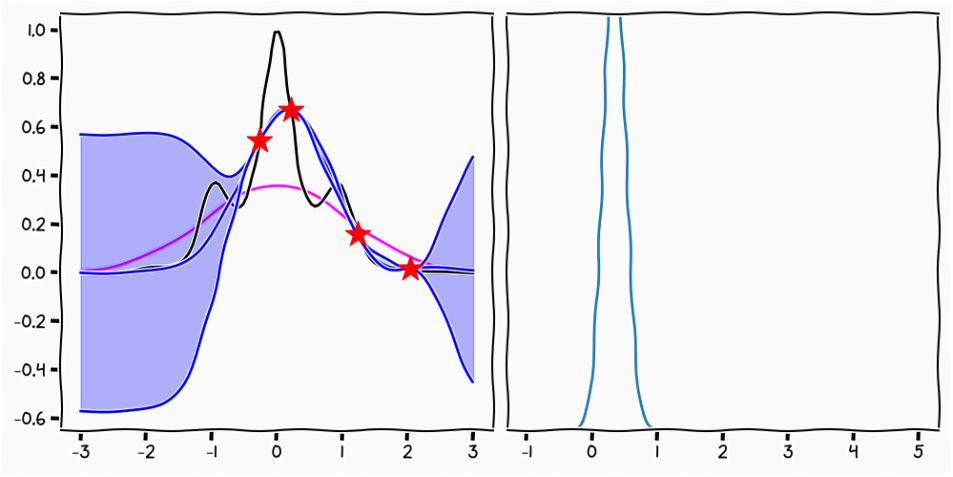

$$p\left(\begin{array}{c} Y \\ F \end{array}\right) = \mathcal{N}\left(\left[\begin{array}{c} \mathbf{m}_X \\ \int m_X(x)\mathrm{d}x \end{array}\right], \left[\begin{array}{cc} k(X,X) & \int k(X,x)\mathrm{d}x \\ \int k(x,X)\mathrm{d}x & \int\int k(x,x')\mathrm{d}x\mathrm{d}x' \end{array}\right]\right)$$

- We can derive $p(F \mid Y)$ through our normal conditioning procedure
- $p(F \mid Y) = \mathcal{N}(\mu_F, k_F)$ is a uni-variate Gaussian
- $p(Y \mid F) = \mathcal{N}(\mu_Y, k_Y)$ is a Gaussian process

Integral Constrained Samples

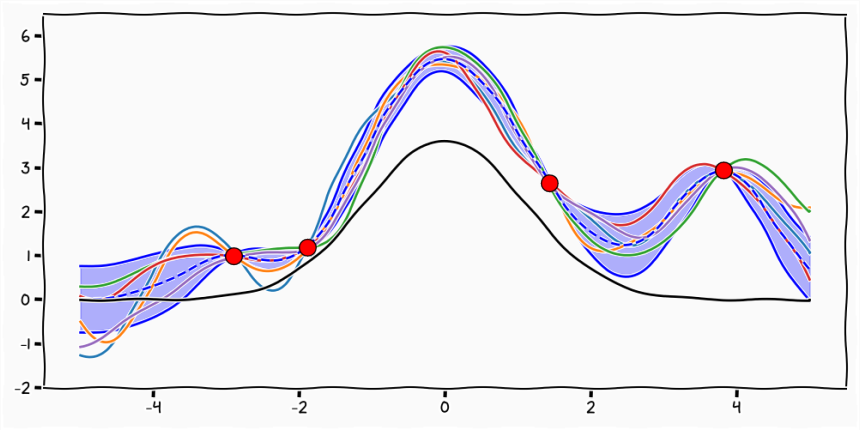

$$F = 4.0$$

Integral Constrained Samples

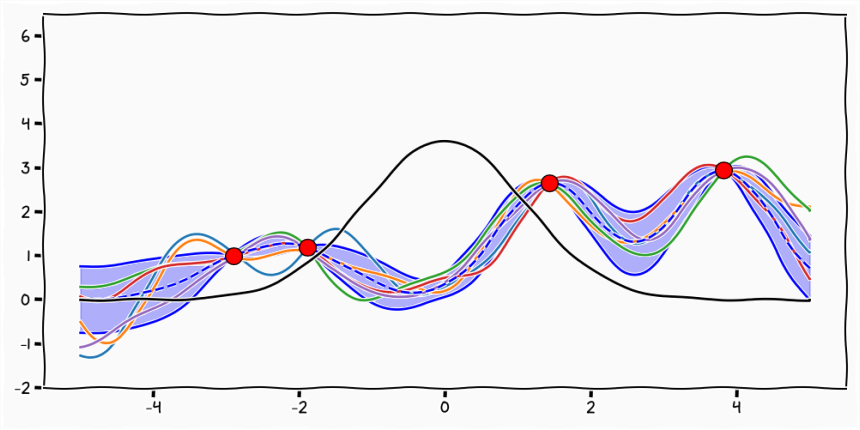

$$F = 1.0$$

Information Operator[3]

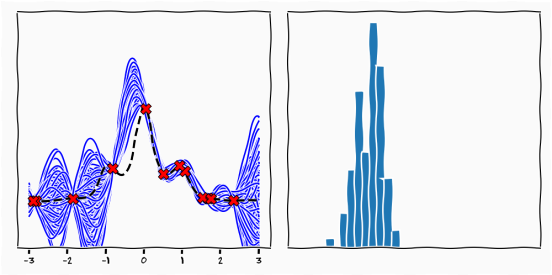

Integrand variance $\alpha(x) = k(x, x)$

Integral Variance Reduction $\alpha(x) = k_F(X, X) - k_F(X, x)$

[3]sometimes called a "Design Rule"

Choice of Covariance

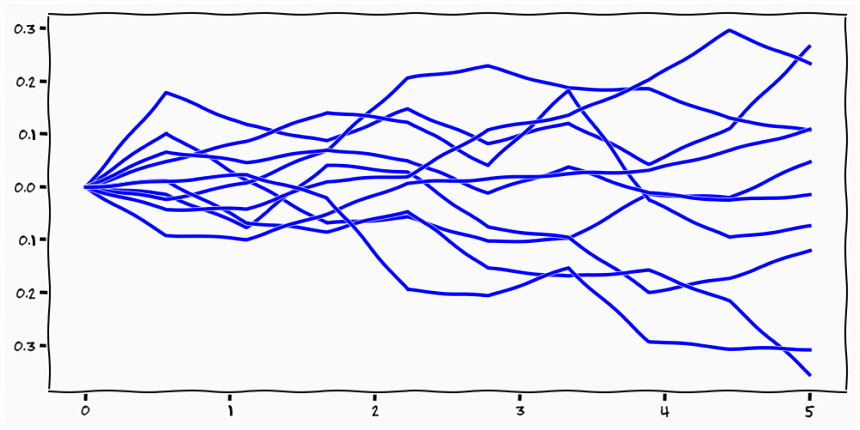

$$p(f) = \mathcal{GP}\left(\mathbf{0}, \theta^2(\min(x, x') - \kappa)\right)$$

Quadrature Rule

$$\mathbb{E}[F] = \mathbb{E}_{p(f|Y)}\left[\int f(x)\mathrm{d}x\right] = \sum_{i=1}^{N-1} \frac{x_{i+1} - x_i}{2}(f(x_{i+1}) + f(x_i))$$

Quadrature Rule

$$\mathbb{E}[F] = \mathbb{E}_{p(f|Y)}\left[\int f(x)\mathrm{d}x\right] = \sum_{i=1}^{N-1} \frac{x_{i+1} - x_i}{2}(f(x_{i+1}) + f(x_i))$$

- This is the normal trapezoid rule!!!

Quadrature Rule

$$\mathbb{E}[F] = \mathbb{E}_{p(f|Y)}\left[\int f(x)\mathrm{d}x\right] = \sum_{i=1}^{N-1} \frac{x_{i+1} - x_i}{2}(f(x_{i+1}) + f(x_i))$$

- This is the normal trapezoid rule!!!
- The algorithm is now tied to our belief in the function!!!!

Quadrature Rule

$$\mathbb{E}[F] = \mathbb{E}_{p(f|Y)}\left[\int f(x)\mathrm{d}x\right] = \sum_{i=1}^{N-1} \frac{x_{i+1} - x_i}{2}(f(x_{i+1}) + f(x_i))$$

- This is the normal trapezoid rule!!!
- The algorithm is now tied to our belief in the function!!!!
- We can do inference over where to sample!!!!!!!!

Definition (Trapedzoid Rule)
The trapezoidal rule is the posterior mean estimate for the integral $F = \int_a^b f(x)\mathrm{d}x$ under any centred Wiener process prior $p(f) = \mathcal{GP}(0, k)$ with $k(x, x) = \theta^2(\min(x, x') - \kappa)$


I'M NOT IMPRESSED

The Scientific Principle

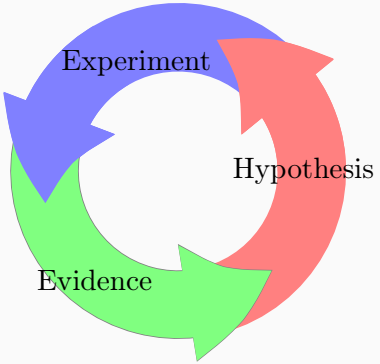

$$\text{Data} + \text{Model} \overset{\textit{Compute}}{\rightarrow} \text{Prediction}$$

Least Squares Regression

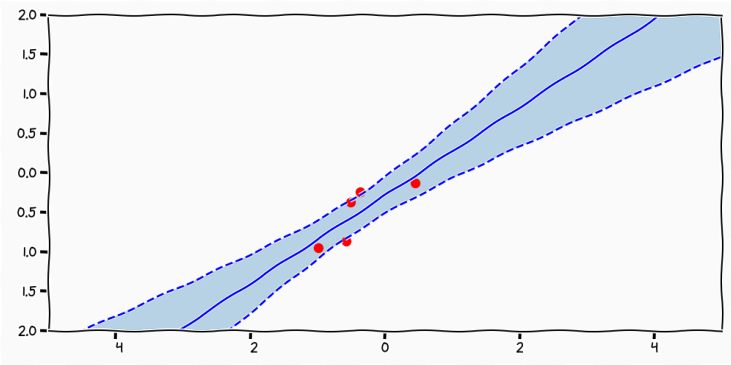

Legendre (1805) algorithm that reduces "error"

Gauss (1809) statistical model assuming i.i.d. Gaussian noise

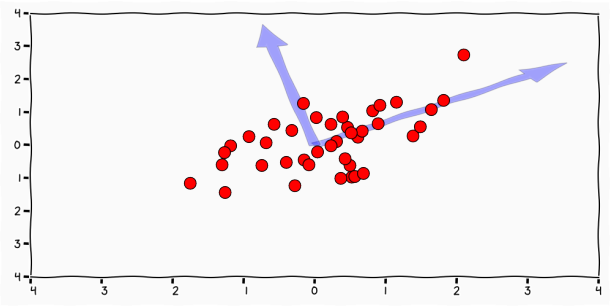

Spearman (1904) proposed an algorithm to extract "factors" from data
Spearman, 1904

Hotelling (1936) concept of factor is clearly defined through a statistical model Hotelling, 1933

Code

```
def minimize(fun, x0, args=(), method=None,
             jac=None, hess=None,
             hessp=None, bounds=None,
             constraints=(), tol=None,
             callback=None, options=None):
```

method Nelder-Mead, Powell, CG, BFGS, Newton-CG, L-BFGS-B, TNC , COBYLA , SLSQP , trust-constr , dogleg , trust-ncg , trust-exact , trust-krylov

Probabilistic Numerics [Hennig et al., 2015]

- There are tons of numerical algorithms for every problem under the sun

Probabilistic Numerics [Hennig et al., 2015]

- There are tons of numerical algorithms for every problem under the sun
- They work really well

Probabilistic Numerics [Hennig et al., 2015]

- There are tons of numerical algorithms for every problem under the sun
- They work really well
- They give different results on the same problem

Probabilistic Numerics [Hennig et al., 2015]

- There are tons of numerical algorithms for every problem under the sun
- They work really well
- They give different results on the same problem
- *what is the prior they implement?*

Neural Networks (Maybe) Evolved to Make Adam The Best Optimizer

In a talk, Olivier Bousquet has described the deep learning community as a giant genetic algorithm: Researchers in this community are exploring the space of all variants of algorithms and architectures in a semi-random way. Things that consistently work in large experiments are kept, the ones not working are discarded. .... the community is evolving only one set of parameters (architectures, initialization strategies, hyperparameters search algorithms, etc.) keeping most of the time the optimizer fixed to Adam.

No Free Lunch[4]

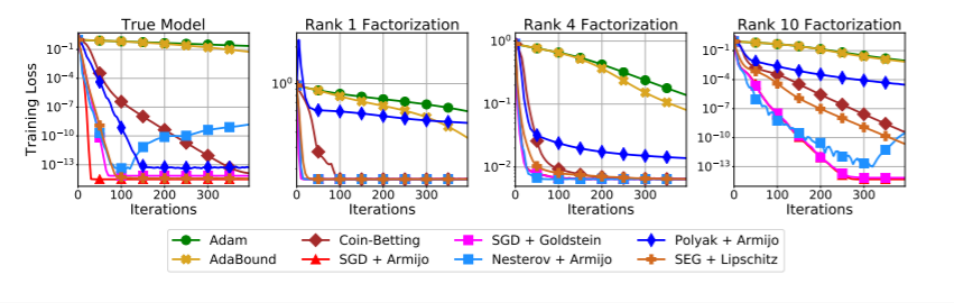

[4]Vaswani et al., 2019.

"The Machine Learning Principle"[5]

"There is a notion of success . . . which I think is novel in the history of science. It interprets success as approximating unanalyzed data."

– Prof. Noam Chomsky

[5] Chomsky et al., 1980

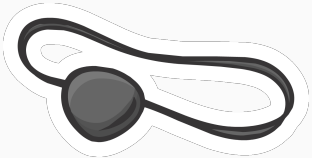

Statistical Learning

$\mathcal{A}_{\mathcal{H}}(\mathcal{S})$

Summary

Summary

- Probabilistic Numerics extends the notion of statistical inference to computation[6]

[6] these thoughts have been around for a long time

Summary

- Probabilistic Numerics extends the notion of statistical inference to computation[6]
- Computation is the process of extracting a latent property, machine learning is the statistical process of updating beliefs about latent properties

[6] these thoughts have been around for a long time

Summary

- Probabilistic Numerics extends the notion of statistical inference to computation[6]
- Computation is the process of extracting a latent property, machine learning is the statistical process of updating beliefs about latent properties
- Computation is often not "truth" therefore we should quantify our ignorance about its results

[6] these thoughts have been around for a long time

Summary

- Probabilistic Numerics extends the notion of statistical inference to computation[6]
- Computation is the process of extracting a latent property, machine learning is the statistical process of updating beliefs about latent properties
- Computation is often not "truth" therefore we should quantify our ignorance about its results
- Why?

[6] these thoughts have been around for a long time

Summary

- Probabilistic Numerics extends the notion of statistical inference to computation[6]
- Computation is the process of extracting a latent property, machine learning is the statistical process of updating beliefs about latent properties
- Computation is often not "truth" therefore we should quantify our ignorance about its results
- Why?
  - efficiency

[6] these thoughts have been around for a long time

Summary

- Probabilistic Numerics extends the notion of statistical inference to computation[6]
- Computation is the process of extracting a latent property, machine learning is the statistical process of updating beliefs about latent properties
- Computation is often not "truth" therefore we should quantify our ignorance about its results
- Why?
  - efficiency
  - down-stream tasks, uncertainty in computation should be part of decision

[6] these thoughts have been around for a long time

Summary

- Probabilistic Numerics extends the notion of statistical inference to computation[6]
- Computation is the process of extracting a latent property, machine learning is the statistical process of updating beliefs about latent properties
- Computation is often not "truth" therefore we should quantify our ignorance about its results
- Why?
  - efficiency
  - down-stream tasks, uncertainty in computation should be part of decision
  - learning/understanding algorithms in relation to problems/data

[6] these thoughts have been around for a long time

Is BO PN?

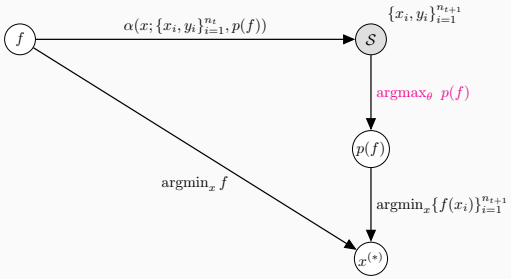

Yes it uses a probablistic model as a proxy for decision loop

Is BO PN?

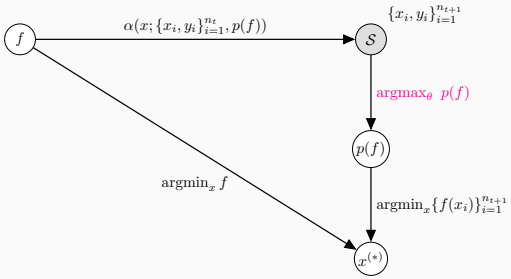

Yes it uses a probablistic model as a proxy for decision loop

No the probabilistic model is not over the quantity of interest

Formalisation

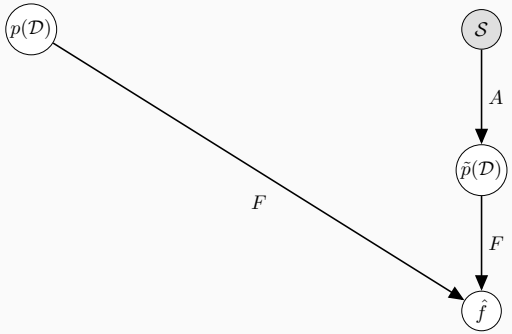

$$A \circ \mathcal{S} = \tilde{p}(\mathcal{D}) \approx p(\mathcal{D})$$

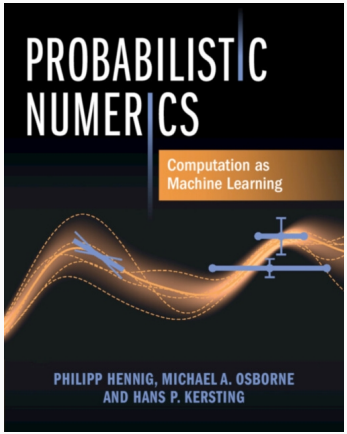

http://probnumschool.org

eof

References

- Chomsky, Noam A and Jerry A Fodor (1980). "The inductivist fallacy." In: *Language and Learning: The Debate between Jean Piaget and Noam Chomsky*.
- Cockayne, Jon, Chris Oates, Tim Sullivan, and Mark Girolami (2017). "Bayesian Probabilistic Numerical Methods." In: *CoRR*.
- Hennig, Philipp, Michael A Osborne, and Mark Girolami (July 2015). "Probabilistic numerics and uncertainty in computations." In: *Proc. R. Soc. A* 471.2179, p. 20150142.

- Hotelling, H (Sept. 1933). "Analysis of a complex of statistical variables into principal components.." In: *Journal of Educational Psychology* 24.6, pp. 417–441.
- Lawrence, Neil D (2005). "Probabilistic non-linear principal component analysis with Gaussian process latent variable models." In: *Journal of Machine Learning Research* 6, pp. 1783–1816.
- Mackay, David J C (Dec. 1991). "Bayesian methods for adaptive models." PhD thesis. California Institute of Technology: California Institute of Technology.
- Neumann, John von and H. H. Goldstine (1947). "Numerical Inverting of Matrices of High Order." In: *Bulletin of the American Mathematical Society* 53.11, pp. 1021–1100.

- O'Hagan, A. (Nov. 1991). "Bayes-Hermite quadrature." In: *Journal of Statistical Planning and Inference* 29.3, pp. 245–260.
- Spearman, Charles (1904). "" General Intelligence," Objectively Determined and Measured." In: *The American Journal of Psychology* 15.2, pp. 201–292.
- Vaswani, Sharan et al. (2019). "Painless Stochastic Gradient: Interpolation, Line-Search, and Convergence Rates." In: *CoRR*.

Learning Hyper-parameters

$$\{\hat{\beta}, \hat{\ell}, \hat{\sigma}\} = \underset{\beta,\ell,\sigma}{\operatorname{argmax}} \quad \log \int p(y \mid f) p(f) \mathrm{d}f$$

- Update the hyper-parameters of the GP
- We can do this by gradient based optimisation
  - *or with a Bayes Opt loop!*

[7] *. . . .and here I was thinking that you guys had some principles in life*

Learning Hyper-parameters

$$\{\hat{\beta}, \hat{\ell}, \hat{\sigma}\} = \underset{\beta,\ell,\sigma}{\operatorname{argmax}} \quad \log \int p(y \mid f) p(f) \mathrm{d}f$$

- Update the hyper-parameters of the GP
- We can do this by gradient based optimisation
  - *or with a Bayes Opt loop!*
- *Are we doing MLE*[7]*?*

[7] *. . . .and here I was thinking that you guys had some principles in life*

Marginal Likelihood

$$
\begin{aligned}
\{\hat{\beta}, \hat{\ell}, \hat{\sigma}\} &= \underset{\beta,\ell,\sigma}{\operatorname{argmax}} \quad \log \int p(y \mid f) p(f) \mathrm{d}f \\
&= \underset{\beta,\ell,\sigma}{\operatorname{argmin}} - \log \quad p(y) \\
&= \underset{\beta,\ell,\sigma}{\operatorname{argmin}} \frac{1}{2} \operatorname{trace}\left(\mathbf{Y}\mathbf{K}^{-1}\mathbf{Y}^{\mathrm{T}}\right) + \frac{1}{2} \log|K| + \frac{N}{2} \log(2\pi)
\end{aligned}
$$

- $Data - fit$ *how well does the observations fit the model*
- $"Complexity"$ *how "smooth" is the functions*

Determinant

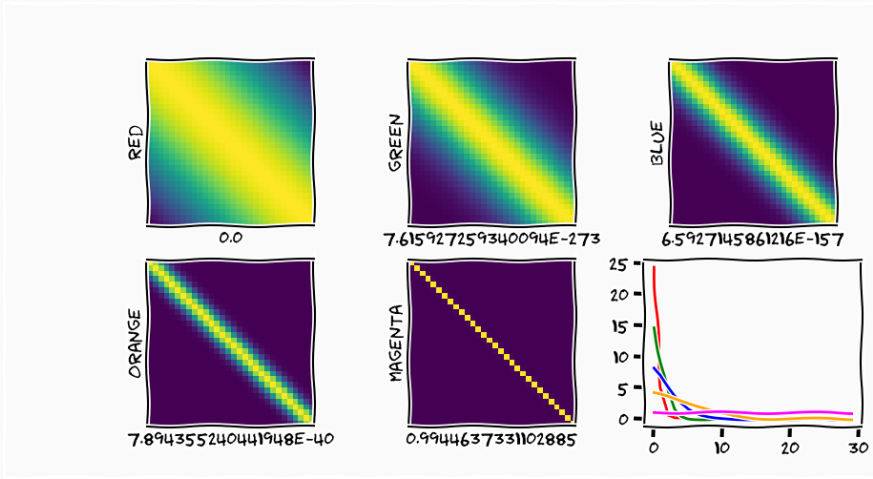

$$\frac{N}{}$$

How to implement

Semantics

$$p(\theta \mid y) = \frac{p(y \mid \theta)p(\theta)}{\int p(y \mid \theta)p(\theta)\mathrm{d}\theta}$$

Likelihood How much evidence is there in the data for a specific hypothesis

Prior What are my beliefs about different hypothesis

Posterior What is my updated belief after having seen data

Evidence What is my belief about the data

Regression Model

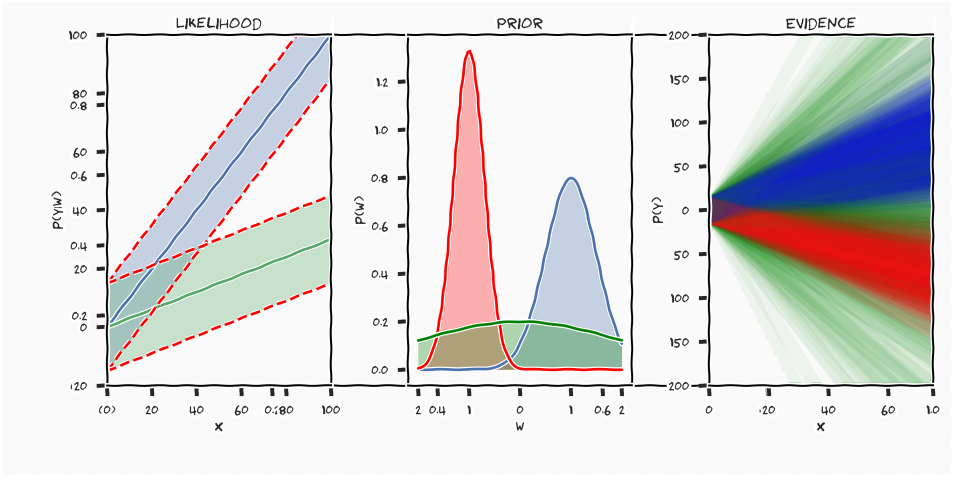

Marginalisation

Next time you want to give your friends a compliment, tell them that you have completely marginalised them from your life

Linear Linear Model

$$p(y_i|x_i, \mathbf{w}) = \mathcal{N}(w_0 + w_1 \cdot x_i, \beta^{-1})$$

Basis function

$$p(y_i|x_i, \mathbf{w}) = \mathcal{N}(\sum_{i=1}^{6} w_i \phi(x_i), \beta^{-1})$$

Linear Linear Regression

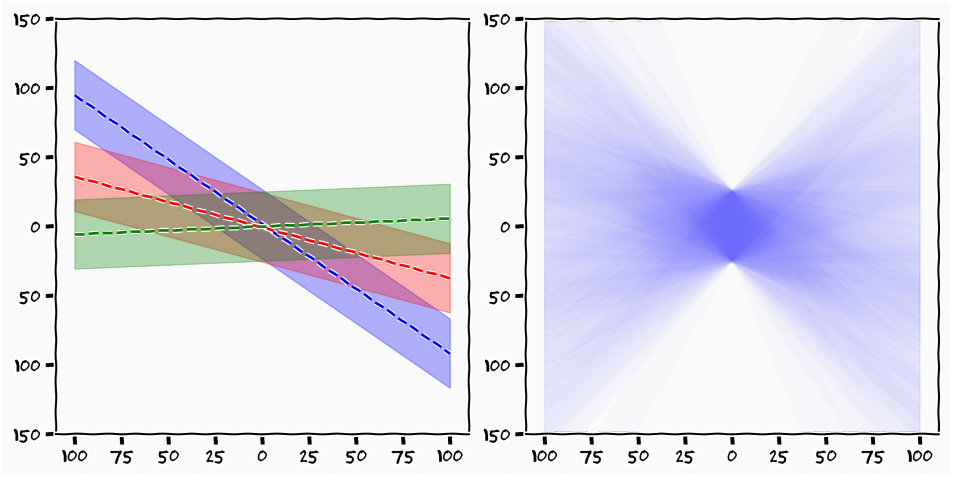

Linear Regression

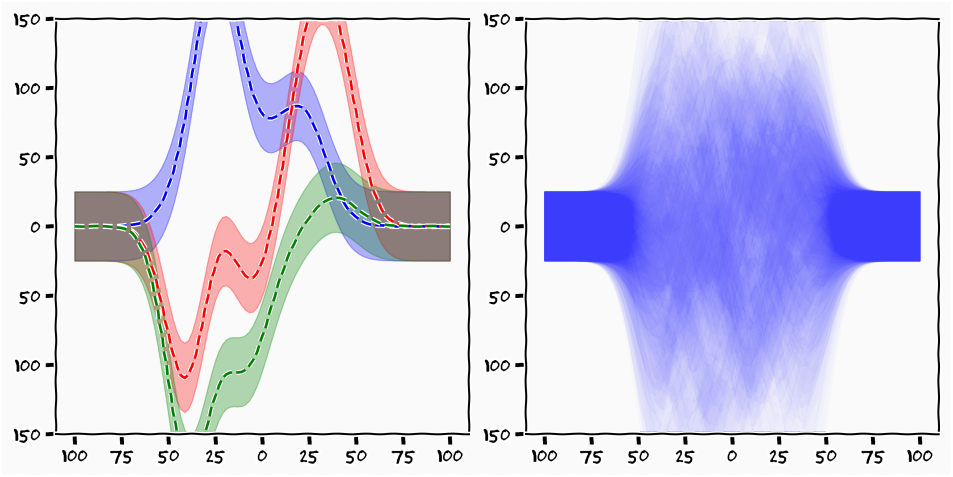

Evidence

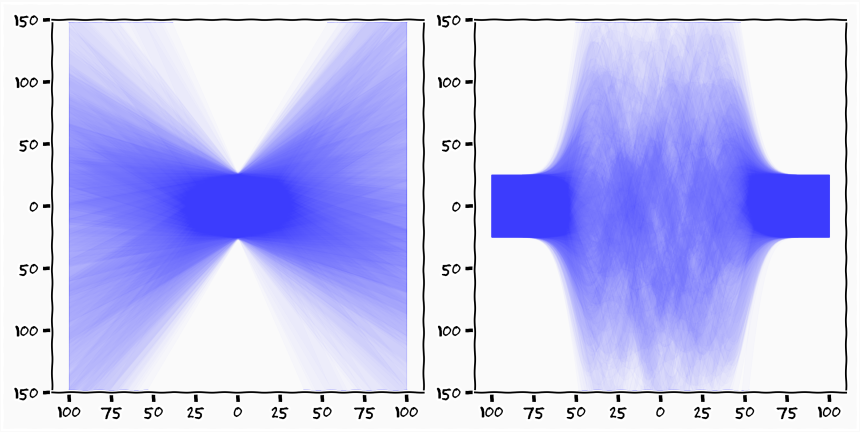

$$p(\mathcal{Y}) = \int p(\mathcal{Y}|\mathbf{W})p(\mathbf{W})\mathrm{d}\mathbf{W}$$

Probabilities are a zero-sum game

The MacKay Plot Mackay, 1991

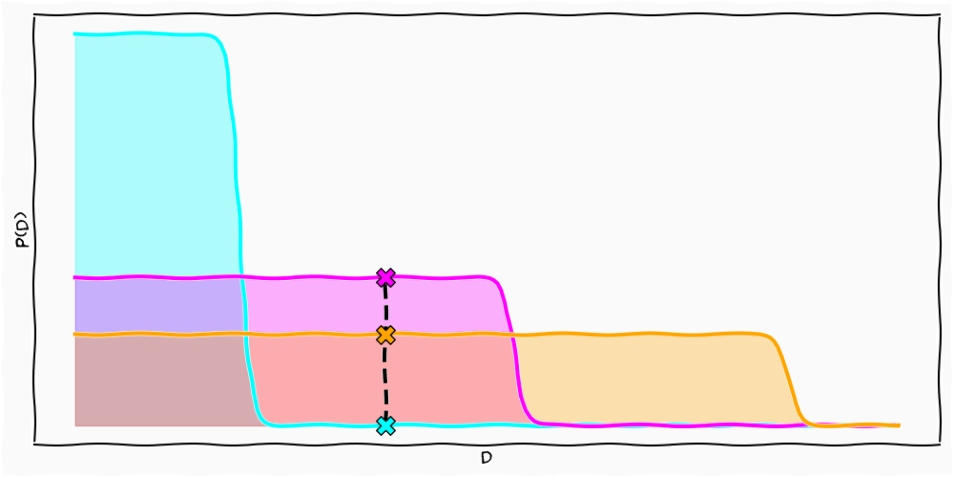

Occams Razor

Definition (Occams Razor)
"All things being equal, the simplest solution tends to be the best one"

– William of Ockham

What is Simple?[8]

[8] https://www.imdb.com/title/tt8132700/

The MacKay Plot Mackay, 1991

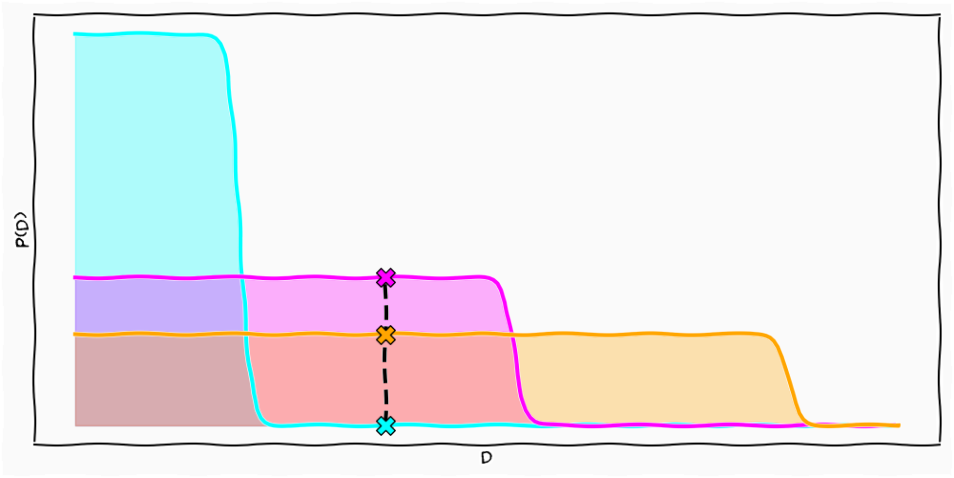

Unsupervised Learning [Lawrence, 2005]

- Regression

$$p(y \mid x) = \int p(y \mid f)p(f \mid x)\mathrm{d}f$$

- "Unsupervised" Learning

$$p(y) = \int p(y \mid f)p(f \mid x)p(x)\mathrm{d}f\mathrm{d}x$$

Unsupervised Learning

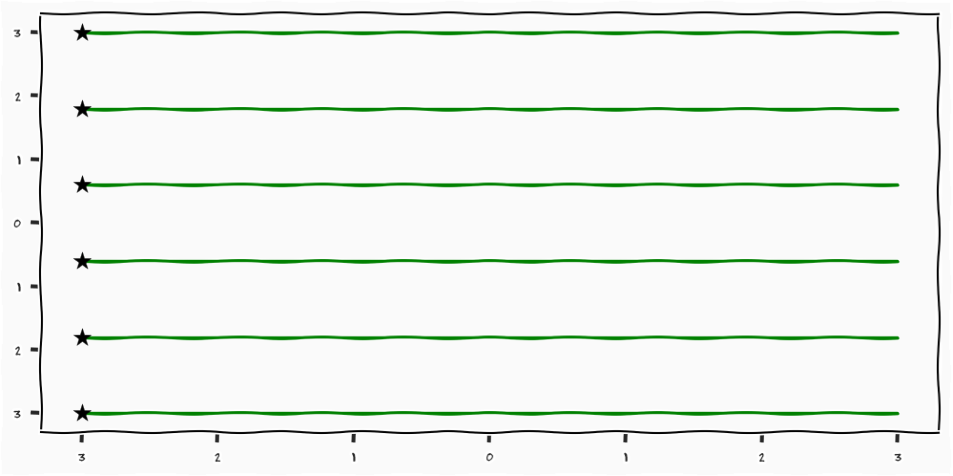

Unsupervised Learning

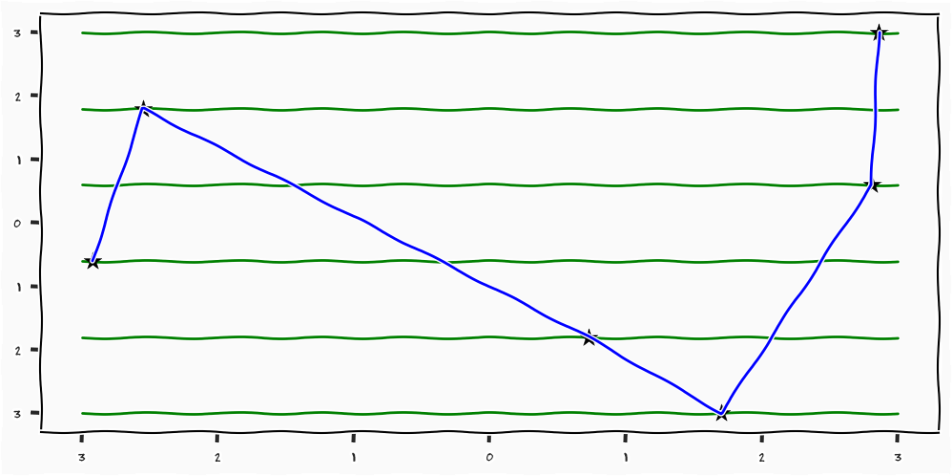

Unsupervised Learning

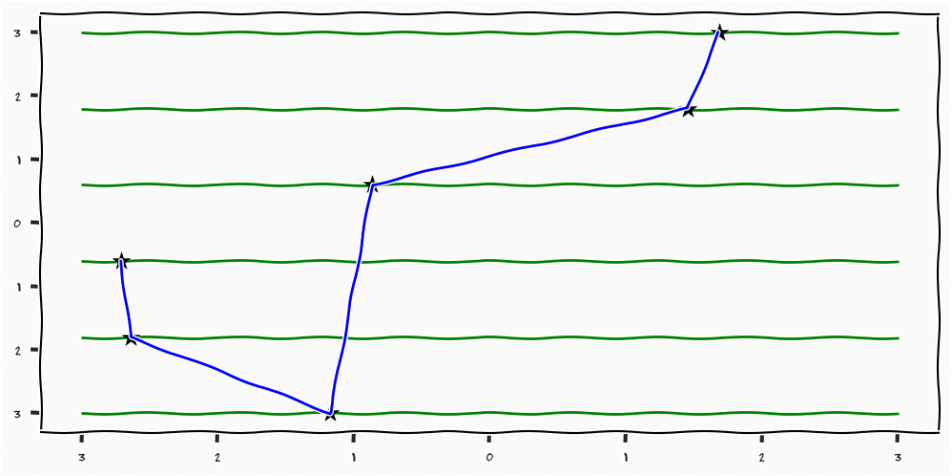

Gaussian Process Latent Variable Model [Lawrence, 2005]

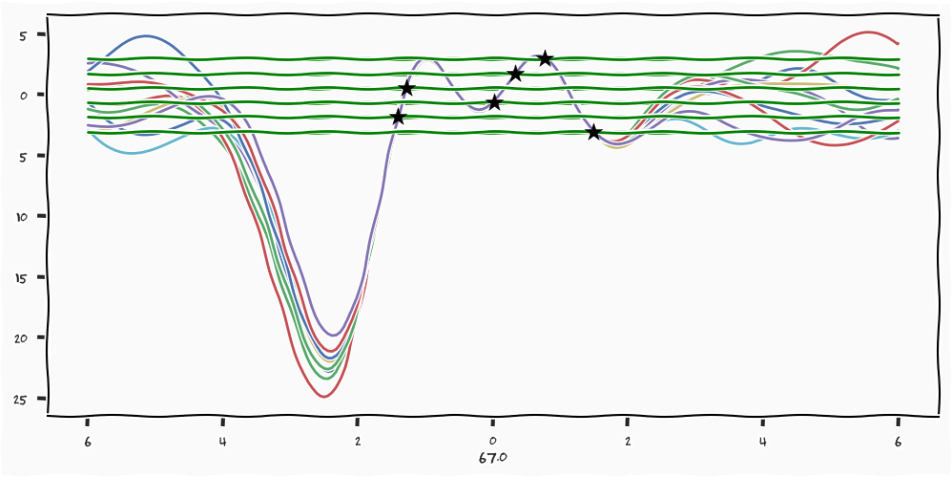

Gaussian Process Latent Variable Model [Lawrence, 2005]

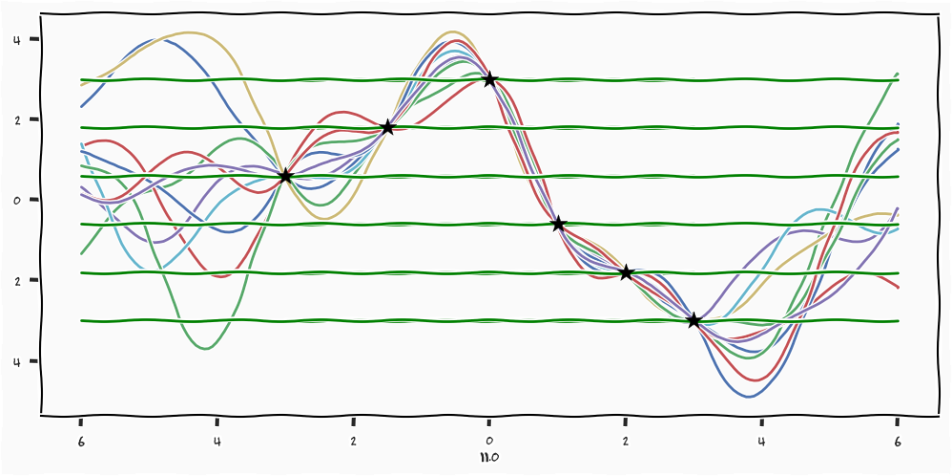

Gaussian Process Latent Variable Model [Lawrence, 2005]

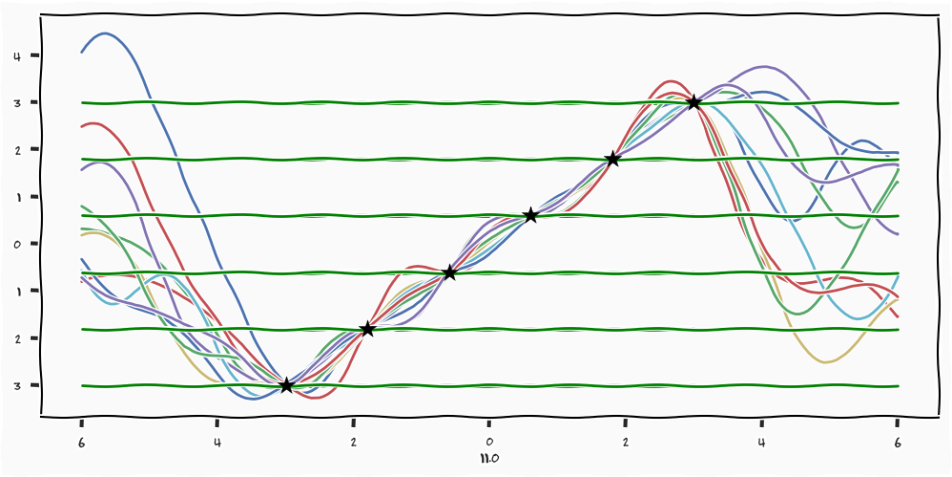

Gaussian Process Latent Variable Model [Lawrence, 2005]